

Campus news

The largest study of AI use by undergrads is in, revealing disparities in access — and in cheating

Researchers found that how students use (and misuse) AI varies by subject and by socioeconomics.

By [Maya L. Kapoor](#)

[↑ Back to Top](#)

UC Berkeley

Yan KruKrau via Pexels

May 21, 2026

In a world where AI can generate research papers, solve equations or create art, educators worry about how college students may be using it, misusing it or missing out on it. Yet there have been few comprehensive studies of college students and their AI use.

Now, [Igor Chirikov](#), a senior researcher at UC Berkeley's Center for Studies in Higher Education, has published [the largest study of generative AI use by undergraduates](#), in collaboration with researchers from the University of Technology Sydney and Cornell University. More than 95,000 students at 20 research-intensive public universities responded to questions about how they use AI, including whether they use it to cheat. The findings were published on May 21 in *Science*.

[↑ Back to Top](#)

UC Berkeley



Igor Chirikov, a senior researcher at UC Berkeley's Center for Studies in Higher Education.

Bora Reed/UC Berkeley

“The arrival of artificial intelligence technologies and GenAI tools like ChatGPT was a big shock to higher education and to faculty, to students — to everybody involved,” Chirikov said. “We didn’t know much about how students were using it and misusing it.”

The study, conducted in the spring of 2024, used data collected by Berkeley’s [Student Experience in the Research University \(SERU\) Consortium](#), a group of research universities that collaborate on surveying students to improve higher education. About two-thirds of respondents said they used GenAI, and almost 40% used it monthly or even more frequently. What’s more, at least 9% of students who used AI reported using it to cheat. That number varied significantly by academic discipline, with more non-STEM students cheating with AI than STEM students. But the researchers caution that banning GenAI won’t stop ch [↑ Back to Top](#) may even harm students when they look for work in industries that expect AI proficiency.

UC Berkeley

The study also uncovered a worrying disparity in use of AI by different demographic groups, with low-income, racially underrepresented and female students using AI less. These students may fall behind in college and eventually the workplace because of unequal access to or practice using AI.

Generative AI is evolving so rapidly, Chirikov said, that this survey feels like it's "from a past life already." Still, he says his findings matter because they can help universities and students alike think about what uses of AI will best serve them.

UC Berkeley News recently spoke with Chirikov about the rising number of students who are using AI, what universities can do to test student learning and how to make sure that students with fewer resources can still develop the AI proficiency they may need in their careers.

Your study found that it's not always clear to students what is cheating versus what is an acceptable use of generative AI, and that there can be a slippery slope to cheating — 26% of daily AI users said they used it to cheat, while only 7% of monthly users did. Could you talk about that a little bit?

AI policies are very different across courses, from faculty allowing AI throughout, including on exams, to completely banning it.

Right now, when you Google something, you use AI, because there's usually a quick summary that pops up that's AI-generated. Or when you use a grammar and editing tool, sometimes there are also AI integrations, and you're just one click away from AI rewriting the whole piece. The level of integration of those tools is incredible, and they are very tempting to use. For students, it's very hard to self-regulate and to navigate this complex environment in terms of classroom AI policies.

We don't know whether frequency of AI use causes students to cheat more, or just students who are more likely to cheat in general tend to use those instruments more often. P

[↑ Back to Top](#)



What do you think motivated the students who knew they were cheating, or thought they were, and did it anyway?

The survey is fielded at research universities that are selective, where a lot of bright students are in this competitive environment where grades matter so much, and it's important to have a perfect GPA to get an internship or get into graduate school. And because higher education is expensive for students and families, there is also pressure to finish as quickly as possible without falling behind.

On the other hand, there are all these AI tools that make their use very easy. Instead of spending an all-nighter on an assignment, you can generate something in under 30 minutes and submit it.

I think those two forces create a perfect storm for AI-assisted cheating. And learning is a casualty of that perfect storm to some extent.

One thing that the study found was that generative AI cheating is not actually as widespread as has been previously reported. Why do you think that is?

There are several caveats to that statement. First of all, this is data from two years ago, relatively early in the cycle of AI adoption. In terms of GenAI capabilities and usage, it was definitely lower than what we're seeing now.


Second, our numbers are conservative. We used an indirect survey method to make it easier for students to answer honestly about a sensitive behavior. But students still had to recognize when their use was not allowed, and that is challenging.

But even with those restrictions, it's still a significant number of students. I show in a [recent paper on grade inflation](#) that when students use AI on assignments, an entire course's grades may become inflated compared to what the students actually know.

[↑ Back to Top](#)

UC Berkeley

Your study suggests that different disciplines should develop different policies about AI use, and different ways of assessing students. Why is that?



When students don't have equal access to AI tools, it may become a disadvantage for their careers. But students also need to be careful about relying too much on AI instead of developing their own skills, Chirikov warns.

Vitaly Gariev via Unsplash

From our findings, we saw that students in different disciplines used AI differently. Any solution should be discipline-specific, or maybe even course-specific.

One response that is gaining traction right now is to move all assessments into controlled environments. For example, proctored oral exams or hand-written in-class exams.

The problem is, those types of assessments only cover a narrow group of skills that can be tested in a time-controlled environment. Universities, and especially research universities, teach students a much broader set of skills. And some of those skills require prolonged engagement with material.

Going back and forth, writing or coding or just struggling intellectually — that's part of how you learn.

[↑ Back to Top](#)



You conclude that the solution to cheating with AI is not blanket policies across universities. Departments need to develop their own policies. Why did you come to that conclusion?

One of the issues that a lot of faculty are running into is that it's not easy to detect GenAI. Sometimes you think student work is by AI, but it may not be. And even if you detect AI use, you may spend a lot of time trying to prove that compared to plagiarism cases, where there's evidence that is easier to collect.

AI detection software is evolving, so there are newer detectors that are able to find AI-generated text better, but it's a cat-and-mouse game, because at the same time, there are AI humanizers, services that allow you to make your text or code look human-written. So it will be a never-ending battle.

Another solution is to ban AI throughout. This is not a productive solution. Students will continue using AI. Some use it in their courses to learn better, to explain material and to ask questions that they would not be comfortable asking their faculty.

So it's challenging for universities to stay away from that wave of AI adoption, and they need to teach students to use AI. But what responsible AI use looks like is different in different fields — writing, coding, problem-solving, lab work and creative work all raise different questions. For that reason, a blanket ban probably wouldn't work.

The study found disparities in the use of generative AI by different members of demographic groups, with low-income, racially underrepresented and female students less likely to use it. Why is that concerning?

This is even more important than the cheating part. We have not had much systematic evidence on disparities in students' use of AI tools.

[↑ Back to Top](#)

UC Berkeley

My concern is that students from wealthier families can access advanced AI tools with stronger capabilities and fewer usage limits. But students who don't have resources may only be able to use free AI tools that are clunkier and have limits.

A lot of employers are interested in graduates who have experience working with AI tools. Students from higher socioeconomic backgrounds may get an advantage that's not necessarily about their skills, but rather about the ability to pay for those tools. That's a very important component of this study.

Why do you think that these findings should matter to students?

Using AI in your learning can do a trick on you. You may produce something polished for class, and may even get a good grade, but not develop the skill the assignment was meant to build.

There are some experimental studies that show that people learn much worse with AI and don't develop durable skills compared to learning without AI. A lot of students may just be very misinformed in terms of how good or bad they are in certain areas, and they may underinvest in some foundational skills.

And we don't know the future of AI in the workplace. In that kind of uncertain world, it's important that you understand how AI is impacting your education.

I encourage students to pause and ask themselves when using AI: "Could I explain this without the tool? Could I do a similar task on my own tomorrow? Did AI help me understand the material better, or did it mainly help me finish faster?" Those simple questions can help students track whether AI is supporting their learning or replacing it.

But it's challenging. I really feel for students.

Why do your study's findings matter for universities?

There was already a crisis of trust in higher education long before AI arrived. But AI creates another point of critique of universities — are they up to the task of teaching and assessing student skills during the Age of AI?

[↑ Back to Top](#)

UC Berkeley

There are already a lot of important efforts under way by universities to navigate the impacts of AI, including at UC Berkeley. But the evidence from our paper shows that these initiatives need more resources and higher prioritization.

This interview has been edited for length and clarity.

UC Berkeley

About

About Berkeley News

Find an expert

Media training

Contact Berkeley News

Browse News

Changemakers

Climate Change

Democracy in Distress

Connect

Berkeleyan newsletter

Media inquiries

Campus resources

Libraries

Equity, Inclusion & Diversity

Communications & Public Affairs

Berkeley Brand

UC Berkeley

Visit

[↑ Back to Top](#)

UC Berkeley

[Accreditation](#)

[HEERF Report](#)

[Contact Us](#)

[Websites A-Z](#)

[Directory](#)

[Map](#)

[Accessibility](#)

[Nondiscrimination](#)

[Privacy](#)



Copyright © 2026 UC Regents; all rights reserved

[↑ Back to Top](#)