

Gopnik AI

Tran Le, tnl22

2026-06-06

Read Survey Responses

```
library(ggplot2)
library(gridExtra)
```

```
## Warning: package 'gridExtra' was built under R version 4.4.3
```

```
survey_path <- "227_participants.csv"

survey_raw <- read.csv(
  survey_path,
  stringsAsFactors = FALSE,
  na.strings = c("", "NA"),
  fileEncoding = "UTF-8"
)

survey <- survey_raw
```

Clean Variable Types

```
# Convert true/false fields.
survey$ai_used <- tolower(survey$ai_used) == "true"
survey$assigned_was_top_choice <- tolower(survey$assigned_was_top_choice) == "true"
survey$low_domain_frequency <- tolower(survey$low_domain_frequency) == "true"

# Convert numeric survey and task variables.
numeric_cols <- c(
  "total_duration_ms",
  "consent_page_duration_ms",
  "first_impressions_duration_ms",
  "what_kind_of_thing_is_ai_duration_ms",
  "how_you_use_ai_duration_ms",
  "ai_use_frequency_duration_ms",
  "your_experience_of_ai_duration_ms",
  "about_you_duration_ms",
  "a_short_task_duration_ms",
  "freq_information_guidance",
```

```

"freq_writing_communication",
"freq_coding_math_technical",
"freq_creative_work",
"freq_learning_understanding",
"freq_advice_decisions",
"freq_summarizing_opinions",
"freq_automation_repetitive",
"hours_per_week",
"exp_ai_does_work_vs_user_decides",
"exp_independent_vs_reliant",
"exp_existing_info_vs_creating_new",
"exp_instrument_vs_point_of_view",
"age",
"task_selection_top_frequency",
"self_reported_freq_in_assigned_category",
"confidence_level",
"epistemic_ownership",
"time_before_ai_ms",
"ai_visible_duration_ms",
"response_started_ms",
"response_submitted_ms",
"task_duration_ms",
"jaccard_similarity_to_ai"
)

for (col in intersect(numeric_cols, names(survey))) {
  survey[[col]] <- suppressWarnings(as.numeric(survey[[col]]))
}

# Make key categorical variables easier to work with.
survey$ai_kind_choice <- factor(survey$ai_kind_choice)
survey$assigned_category <- factor(survey$assigned_category)
survey$highest_education <- factor(survey$highest_education)
survey$ai_understanding <- factor(survey$ai_understanding)

survey$ai_use_tenure <- factor(
  survey$ai_use_tenure,
  levels = c(
    "I do not use them, or am just starting out.",
    "Less than 6 months",
    "6 to 12 months",
    "1 to 2 years",
    "More than 2 years"
  ),
  ordered = TRUE
)

survey$ai_reading_depth <- factor(
  survey$ai_reading_depth,
  levels = c("not_read", "glanced", "skimmed", "read_carefully"),
  ordered = TRUE
)

```

Mixed Regression Models

The main research question is whether participants' pre-task beliefs about AI and independence predict cognitive offloading during the task.

Here, “cognitive offloading” is measured in three related ways:

1. **AI access:** whether the participant used the AI assistant at all (`ai_used`).
2. **AI uptake:** how similar the participant's final answer was to the pre-written AI response. I report both lexical overlap (`jaccard_similarity_to_ai`) and, after the embedding section is computed, semantic similarity (`embedding_cosine_similarity_to_ai`).
3. **Epistemic ownership:** whether the participant said the final response reflected mostly AI thinking or mostly their own thinking (`epistemic_ownership`).

The mixed part of the model is (`1 | task_id`). This means the model allows each task prompt/pre-written AI response to have its own baseline. That matters because some tasks may naturally make people more likely to use or copy AI than others.

```
library(lme4)
```

```
## Warning: package 'lme4' was built under R version 4.4.3
```

```
## Loading required package: Matrix
```

```
# lmerTest adds p-values to the linear mixed models.
```

```
library(lmerTest)
```

```
## Warning: package 'lmerTest' was built under R version 4.4.3
```

```
##
```

```
## Attaching package: 'lmerTest'
```

```
## The following object is masked from 'package:lme4':
```

```
##
```

```
## lmer
```

```
## The following object is masked from 'package:stats':
```

```
##
```

```
## step
```

```
z <- function(x) as.numeric(scale(x))
```

```
# z = (person's score - sample mean) / sample standard deviation
```

```
model_data <- survey
```

```
model_data$task_id <- factor(model_data$task_id)
```

```
# Two pre-task independence sliders, kept as separate predictors (r = .20,  
# too low to justify combining into a composite).
```

```
# Higher values mean the participant described themselves as doing the work /  
# staying independent.
```

```
model_data$z_user_decides <- z(model_data$exp_ai_does_work_vs_user_decides)
```

```

model_data$z_independent <- z(-model_data$exp_independent_vs_reliant) # reverse

# Low inter-item correlation; treat as separate predictors rather than a composite.
cor(model_data$z_user_decides, model_data$z_independent, use = "complete.obs")

## [1] 0.2100162

# Control variables.
model_data$z_hours_per_week <- z(model_data$hours_per_week)
model_data$z_instrument_vs_point_of_view <- z(model_data$exp_instrument_vs_point_of_view)
model_data$z_existing_info_vs_creating_new <- z(model_data$exp_existing_info_vs_creating_new)
model_data$z_jaccard_similarity_to_ai <- z(model_data$jaccard_similarity_to_ai)

# Binary recode of epistemic ownership.
# 1-3 = primarily AI's thinking (0); 5-7 = primarily own thinking (1).
# Six participants at the midpoint (4) are excluded from binary ownership models.
model_data$epistemic_ownership_binary <- ifelse(
  model_data$epistemic_ownership <= 3, 0L,
  ifelse(model_data$epistemic_ownership >= 5, 1L, NA_integer_)
)

# Binary reading depth: read_carefully (1) vs. anything less (0).
# Used in place of the ordered 1-4 numeric scale because ~57% of AI users
# reported reading_carefully, leaving few observations in the lower categories.
model_data$read_carefully <- as.integer(
  model_data$ai_reading_depth == "read_carefully"
)

# Use the original epistemic ownership scale.
# Original scale: 1 = mostly AI's thinking, 7 = mostly my thinking.
# Higher values mean more of the participant's own thinking.
ai_users <- subset(model_data, ai_used == TRUE)

```

Model 1: Do Pre-Task Independence Beliefs Predict Whether People Use AI?

This model asks whether participants who described themselves as more independent on either pre-task slider were less likely to consult the AI assistant during the task. The two independence items are entered as separate predictors: **z_user_decides** (higher = participant prefers to do the work themselves) and **z_independent** (higher = participant expects to stay more independent).

Outcome: **ai_used**

Model type: mixed logistic regression

Random intercept: **task_id**

```

model_ai_used <- glmer(
  ai_used ~ z_user_decides +
    z_independent +
    z_hours_per_week +
    z_instrument_vs_point_of_view +
    (1 | task_id),
  data = model_data,
  family = binomial,

```

```
control = glmerControl(optimizer = "bobyqa")
)
```

```
## boundary (singular) fit: see help('isSingular')
```

```
summary(model_ai_used)
```

```
## Generalized linear mixed model fit by maximum likelihood (Laplace
## Approximation) [glmerMod]
## Family: binomial ( logit )
## Formula: ai_used ~ z_user_decides + z_independent + z_hours_per_week +
##           z_instrument_vs_point_of_view + (1 | task_id)
## Data: model_data
## Control: glmerControl(optimizer = "bobyqa")
##
##           AIC          BIC      logLik -2*log(L)  df.resid
##          262.0          282.6      -125.0      250.0       221
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -3.7115 -0.9226  0.4116  0.7047  1.3366
##
## Random effects:
##   Groups Name            Variance Std.Dev.
##   task_id (Intercept) 0          0
## Number of obs: 227, groups: task_id, 18
##
## Fixed effects:
##
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)      0.9816    0.1660   5.912 3.38e-09 ***
## z_user_decides    -0.7946    0.1989  -3.996 6.45e-05 ***
## z_independent     -0.1246    0.1600  -0.779  0.4361
## z_hours_per_week   0.3117    0.1807   1.725  0.0846 .
## z_instrument_vs_point_of_view 0.1682    0.1591   1.057  0.2905
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Correlation of Fixed Effects:
##              (Intr) z_sr_d z_ndpn z_hr__
## z_user_dcds -0.350
## z_indepndnt -0.025 -0.166
## z_hrs_pr_wk  0.119 -0.053  0.032
## z_nstrm_____ 0.014 -0.020  0.133 -0.224
## optimizer (bobyqa) convergence code: 0 (OK)
## boundary (singular) fit: see help('isSingular')
```

```
fixed_effects_table(model_ai_used, exponentiate = TRUE)
```

```
##              term      estimate std_error      p_value odds_ratio
## 1              (Intercept)  0.9815628 0.1660341 3.383608e-09 2.6686235
## 2      z_user_decides -0.7946191 0.1988651 6.448468e-05 0.4517533
## 3      z_independent -0.1245733 0.1599696 4.361383e-01 0.8828736
```

```
## 4                z_hours_per_week  0.3116570 0.1807144 8.460266e-02 1.3656862
## 5 z_instrument_vs_point_of_view  0.1682300 0.1591488 2.904837e-01 1.1832087
```

The **user_decides** independence slider significantly predicted lower odds of using the AI assistant, **OR = 0.45, $p < .001$** . Participants who more strongly said the user does the work / makes the decisions were less likely to consult AI. The separate **z_independent** slider was not significant, **OR = 0.88, $p = .436$** . Hours per week was marginal but not significant, **OR = 1.37, $p = .085$** , and instrument-vs-point-of-view beliefs were not significant, **OR = 1.18, $p = .290$** .

The mixed model produced a singular fit (task_id random intercept variance = 0). As a robustness check, a fixed-effects logistic regression was fit without the random intercept.

```
model_ai_used_glm <- glm(
  ai_used ~ z_user_decides +
    z_independent +
    z_hours_per_week +
    z_instrument_vs_point_of_view,
  data = model_data,
  family = binomial
)

summary(model_ai_used_glm)
```

```
##
## Call:
## glm(formula = ai_used ~ z_user_decides + z_independent + z_hours_per_week +
##      z_instrument_vs_point_of_view, family = binomial, data = model_data)
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)      0.9816    0.1660   5.912 3.38e-09 ***
## z_user_decides    -0.7946    0.1989  -3.996 6.44e-05 ***
## z_independent     -0.1246    0.1600  -0.779  0.4361
## z_hours_per_week   0.3117    0.1807   1.725  0.0846 .
## z_instrument_vs_point_of_view  0.1682    0.1591   1.057  0.2905
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##    Null deviance: 278.84  on 226  degrees of freedom
## Residual deviance: 250.02  on 222  degrees of freedom
## AIC: 260.02
##
## Number of Fisher Scoring iterations: 4
```

```
fixed_effects_table(model_ai_used_glm, exponentiate = TRUE)
```

	term	estimate	std_error	p_value	odds_ratio
## 1	(Intercept)	0.9815628	0.1660304	3.380870e-09	2.6686235
## 2	z_user_decides	-0.7946191	0.1988579	6.444508e-05	0.4517533
## 3	z_independent	-0.1245733	0.1599683	4.361344e-01	0.8828736
## 4	z_hours_per_week	0.3116570	0.1807129	8.460004e-02	1.3656862
## 5	z_instrument_vs_point_of_view	0.1682300	0.1591477	2.904805e-01	1.1832087

Model 2: Do Pre-Task Independence Beliefs Predict Text Similarity To AI?

This model only uses participants who consulted the AI assistant. It asks whether people who described themselves as more independent on either slider nevertheless produced answers that were more similar to the AI response.

Note on distribution: Among AI users ($n = 158$), `jaccard_similarity_to_ai` is severely bimodal. Approximately 35% of AI users ($n = 55$) had a Jaccard score of exactly 1.0, meaning their response was an exact copy of the pre-written AI response. The remaining participants show a spread of lower values (median approx. 0.31). This bimodality should be kept in mind when interpreting linear models on this outcome. Models are retained as-is, but effects should be interpreted as reflecting this two-group structure.

Outcome: `jaccard_similarity_to_ai`

Model type: linear mixed regression

Random intercept: `task_id`

```
model_similarity <- lmer(
  jaccard_similarity_to_ai ~ z_user_decides +
    z_independent +
    z_hours_per_week +
    z_instrument_vs_point_of_view +
    (1 | task_id),
  data = ai_users
)

summary(model_similarity)
```

```
## Linear mixed model fit by REML. t-tests use Satterthwaite's method [
## lmerModLmerTest]
## Formula:
## jaccard_similarity_to_ai ~ z_user_decides + z_independent + z_hours_per_week +
##      z_instrument_vs_point_of_view + (1 | task_id)
##      Data: ai_users
##
## REML criterion at convergence: 191.1
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -1.6370 -0.8175 -0.2433  0.9534  1.5981
##
## Random effects:
##  Groups   Name                Variance Std.Dev.
## task_id  (Intercept)  0.02084   0.1444
## Residual                    0.15934   0.3992
## Number of obs: 158, groups: task_id, 18
##
## Fixed effects:
##
##              Estimate Std. Error    df t value Pr(>|t|)
## (Intercept)    0.50358    0.04748 16.07656  10.607 1.15e-08
## z_user_decides -0.04194    0.03304 151.16617  -1.269   0.206
## z_independent  -0.04106    0.03465 150.52019  -1.185   0.238
## z_hours_per_week  0.01427    0.03207 146.94145   0.445   0.657
## z_instrument_vs_point_of_view  0.01494    0.03418 150.61159   0.437   0.663
##
## (Intercept)          ***
```

```
## z_user_decides
## z_independent
## z_hours_per_week
## z_instrument_vs_point_of_view
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Correlation of Fixed Effects:
##              (Intr) z_sr_d z_ndpn z_hr__
## z_user_dcds   0.126
## z_indepndnt   0.041 -0.213
## z_hrs_pr_wk  -0.079 -0.087 -0.067
## z_nstrm_____ -0.058 -0.177  0.203 -0.124
```

```
fixed_effects_table(model_similarity)
```

```
##              term      estimate std_error    p_value
## 1      (Intercept)  0.50357786 0.04747811 1.146702e-08
## 2      z_user_decides -0.04194347 0.03304321 2.062663e-01
## 3      z_independent -0.04106184 0.03464951 2.378603e-01
## 4      z_hours_per_week  0.01426732 0.03207427 6.571034e-01
## 5 z_instrument_vs_point_of_view  0.01493647 0.03417693 6.627129e-01
```

Among participants who used the AI assistant, neither pre-task independence predictor significantly predicted Jaccard similarity to the AI response: **z_user_decides**, $b = -0.042$, $p = .206$; **z_independent**, $b = -0.041$, $p = .238$. Hours per week, $b = 0.014$, $p = .657$, and instrument-vs-point-of-view beliefs, $b = 0.015$, $p = .663$, were also not significant predictors. This AI-user-only result is weaker than the all-participant correlations reported in the sample-scope check below.

Model 3: Does Objective AI Uptake Predict Epistemic Ownership?

This model asks whether participants whose answers were more similar to the AI response were less likely to feel the response reflected their own thinking. Because **epistemic_ownership** is U-shaped (many 1s, 2s, 6s, 7s; few responses near the midpoint), it is recoded as binary: 0 = primarily AI's thinking (scores 1-3) vs. 1 = primarily own thinking (scores 5-7). The 14 AI users at the midpoint (4) are excluded.

Outcome: **epistemic_ownership_binary** (0 = AI's thinking, 1 = own thinking)

Model type: mixed logistic regression

Random intercept: **task_id**

```
model_epistemic_ownership <- glmer(
  epistemic_ownership_binary ~ z_jaccard_similarity_to_ai +
    z_user_decides +
    z_independent +
    z_hours_per_week +
    (1 | task_id),
  data = ai_users,
  family = binomial,
  control = glmerControl(optimizer = "bobyqa")
)

summary(model_epistemic_ownership)
```



```

## Generalized linear mixed model fit by maximum likelihood (Laplace
## Approximation) [glmerMod]
## Family: binomial ( logit )
## Formula:
## epistemic_ownership_binary ~ z_jaccard_similarity_to_ai + z_user_decides +
## z_independent + z_hours_per_week + (1 | task_id)
## Data: ai_users
## Control: glmerControl(optimizer = "bobyqa")
##
##          AIC          BIC      logLik -2*log(L)  df.resid
##        160.8        178.6       -74.4     148.8      138
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -1.9287 -0.6263 -0.3268  0.7832  3.3727
##
## Random effects:
## Groups Name          Variance Std.Dev.
## task_id (Intercept) 0.03836  0.1959
## Number of obs: 144, groups: task_id, 18
##
## Fixed effects:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    -0.47898    0.21415  -2.237   0.0253 *
## z_jaccard_similarity_to_ai -1.12868    0.24547  -4.598 4.27e-06 ***
## z_user_decides      0.20344    0.20164   1.009   0.3130
## z_independent     -0.07176    0.20196  -0.355   0.7224
## z_hours_per_week    0.13194    0.19967   0.661   0.5088
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Correlation of Fixed Effects:
##              (Intr) z_j___ z_sr_d z_ndpn
## z_jccrd_s__ -0.041
## z_user_dcds  0.151 -0.105
## z_indepndnt  0.044  0.054 -0.164
## z_hrs_pr_wk -0.154 -0.020 -0.152 -0.038

```

```
fixed_effects_table(model_epistemic_ownership, exponentiate = TRUE)
```

```

##              term      estimate std_error      p_value odds_ratio
## 1              (Intercept) -0.4789847 0.2141522 2.530887e-02 0.6194119
## 2 z_jaccard_similarity_to_ai -1.1286803 0.2454708 4.265212e-06 0.3234598
## 3           z_user_decides  0.2034378 0.2016415 3.130186e-01 1.2256089
## 4           z_independent -0.0717574 0.2019631 7.223652e-01 0.9307567
## 5           z_hours_per_week  0.1319359 0.1996741 5.087686e-01 1.1410352

```

Greater Jaccard similarity to the AI response significantly predicted lower odds of reporting primarily own-thinking ownership, OR = 0.32, $p < .001$. Neither independence predictor nor hours per week was significant ($z_{\text{user_decides}}$, OR = 1.23, $p = .313$; $z_{\text{independent}}$, OR = 0.93, $p = .722$; hours per week, OR = 1.14, $p = .509$).

Model 4: Do Pre-Task Independence Beliefs Moderate How People Interpret AI Influence?

For example, imagine two participants both have high similarity to the AI response:

Participant A rated themselves as not very independent. Participant B rated themselves as very independent.

So the interaction asks: does Participant B, the independent person, give a higher epistemic ownership rating than Participant A, although both produced AI-like answers?

Because the two independence sliders are separate predictors, there are two interaction terms:

- `z_user_decides:z_jaccard_similarity_to_ai`
- `z_independent:z_jaccard_similarity_to_ai`

Although sample size has increased to $n = 158$ AI users, interaction effects require substantial power and should still be interpreted cautiously.

Outcome: `epistemic_ownership_binary` (0 = AI's thinking, 1 = own thinking)

Model type: mixed logistic regression with interactions

Random intercept: `task_id`

```
model_contrast <- glmer(
  epistemic_ownership_binary ~
    z_user_decides * z_jaccard_similarity_to_ai +
    z_independent * z_jaccard_similarity_to_ai +
    z_hours_per_week +
    (1 | task_id),
  data = ai_users,
  family = binomial,
  control = glmerControl(optimizer = "bobyqa")
)

summary(model_contrast)

## Generalized linear mixed model fit by maximum likelihood (Laplace
##   Approximation) [glmerMod]
## Family: binomial ( logit )
## Formula:
## epistemic_ownership_binary ~ z_user_decides * z_jaccard_similarity_to_ai +
##   z_independent * z_jaccard_similarity_to_ai + z_hours_per_week +
##   (1 | task_id)
## Data: ai_users
## Control: glmerControl(optimizer = "bobyqa")
##
##           AIC          BIC      logLik -2*log(L)  df.resid
##        163.9        187.7       -74.0     147.9      136
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -1.6938 -0.6330 -0.3097  0.7949  3.9263
##
## Random effects:
## Groups Name          Variance Std.Dev.
## task_id (Intercept) 0.03176  0.1782
```

```

## Number of obs: 144, groups: task_id, 18
##
## Fixed effects:
##
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)      -0.5221    0.2213  -2.359  0.0183
## z_user_decides       0.2201    0.2115   1.041  0.2981
## z_jaccard_similarity_to_ai -1.1775    0.2688  -4.380 1.18e-05
## z_independent       -0.1264    0.2176  -0.581  0.5613
## z_hours_per_week     0.1278    0.2001   0.639  0.5229
## z_user_decides:z_jaccard_similarity_to_ai  0.1121    0.2448   0.458  0.6471
## z_jaccard_similarity_to_ai:z_independent -0.2268    0.2669  -0.850  0.3954
##
## (Intercept)          *
## z_user_decides
## z_jaccard_similarity_to_ai      ***
## z_independent
## z_hours_per_week
## z_user_decides:z_jaccard_similarity_to_ai
## z_jaccard_similarity_to_ai:z_independent
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Correlation of Fixed Effects:
##              (Intr) z_sr_d z_j___ z_ndpn z_hr__ z__:_
## z_user_dcds  0.088
## z_jccrd_s__  0.061 -0.116
## z_indepndnt  0.179 -0.163  0.215
## z_hrs_pr_wk -0.141 -0.170 -0.014 -0.023
## z_sr_d:_____ -0.108  0.269  0.020 -0.056 -0.068
## z_jccr___:_  0.254 -0.041  0.353  0.337  0.018 -0.206

```

```
fixed_effects_table(model_contrast, exponentiate = TRUE)
```

```

##              term      estimate std_error      p_value
## 1              (Intercept) -0.5221174 0.2212833 1.829971e-02
## 2              z_user_decides  0.2200928 0.2115026 2.980542e-01
## 3              z_jaccard_similarity_to_ai -1.1774767 0.2688091 1.184917e-05
## 4              z_independent -0.1264309 0.2176280 5.612745e-01
## 5              z_hours_per_week  0.1278081 0.2000686 5.229391e-01
## 6 z_user_decides:z_jaccard_similarity_to_ai  0.1120546 0.2447612 6.470875e-01
## 7 z_jaccard_similarity_to_ai:z_independent -0.2267967 0.2668848 3.954403e-01
## odds_ratio
## 1  0.5932630
## 2  1.2461923
## 3  0.3080551
## 4  0.8812351
## 5  1.1363350
## 6  1.1185739
## 7  0.7970828

```

Greater AI similarity again significantly predicted lower odds of reporting primarily own-thinking ownership, OR = 0.31, $p < .001$. Neither interaction with independence beliefs was significant (z_user_decides x similarity: OR = 1.12, $p = .647$; z_independent x similarity: OR = 0.80, $p = .395$), indicating that the uptake-to-ownership link did not differ reliably by pre-task independence beliefs.

Altogether: Pre-task independence beliefs mattered most for whether participants opened AI, not how much they copied after opening it. **The user_decides slider predicted lower AI use, OR = 0.45, $p < .001$,** but neither independence slider significantly predicted Jaccard similarity among AI users. **Participants' post-task ownership ratings tracked behavioral uptake: higher similarity predicted lower odds of primarily own-thinking ownership in both ownership models, ORs = 0.32 and 0.31, $ps < .001$.** This suggests that participants were at least partly aware of the AI's contribution. Neither independence predictor moderated this uptake-to-ownership link.

Sample-Scope Check: All Participants vs. AI Users

Several apparent contrasts depend on whether the analysis includes all participants or only people who actually opened the AI assistant. This matters because non-AI users have low or zero AI uptake by design, so all-participant associations partly capture the decision to use AI at all.

Across all participants ($N = 227$), the composite independence score was negatively related to Jaccard similarity, **Pearson $r = -.23$, $p < .001$** ; Spearman $\rho = -.22$, $p < .001$. This matches the broader direction found by the mentor: participants who described themselves as more independent tended to produce less AI-similar final responses overall. However, among AI users only, the individual independence sliders did not significantly predict Jaccard similarity in Model 2. This suggests independence beliefs mainly predicted whether people opened AI, rather than how much they copied after opening it.

Domain-specific AI-use frequency shows the same sample-scope issue. Across all participants, domain frequency was weakly positively related to Jaccard similarity using Pearson correlation, $r = .13$, $p = .043$, but not using Spearman correlation, $\rho = .10$, $p = .116$. In the AI-user-only regression in Analysis 7, domain frequency did not predict Jaccard similarity, $b = -0.013$, $p = .757$. **The clearest domain-frequency result is AI access itself: people who reported using AI more often in the assigned task domain were more likely to open AI, OR = 1.82, $p < .001$.**

The timing data also show that AI use usually happened early. Among AI users, 78 of 158 participants (49%) opened AI within 50 seconds, and 153 of the 157 participants with both timing values started writing after opening AI. This supports the interpretation that most AI users consulted the assistant before drafting their own response.

Note on exploratory analyses (Analyses 1-8): All analyses below are exploratory. They were motivated by theory and by inspection of the data but were not pre-registered. No correction for multiple comparisons is applied across analyses. Within each analysis, Tukey adjustments are used for pairwise comparisons. Results should be interpreted as hypothesis-generating rather than confirmatory, and any significant findings should be treated with appropriate caution pending replication.

Analysis 1: Does Conceptualizing AI Differently Predict Offloading?

This analysis asks whether participants' answer to "What kind of thing is AI?" predicts cognitive offloading during the short task.

This connects most directly to the paper *Large AI models are cultural and social technologies*. That paper argues that large AI models should not only be understood as agents, but also as cultural/social technologies similar to writing, print, libraries, markets, bureaucracies, and other systems for organizing human knowledge.

Because this is exploratory, I do pairwise comparisons. This means I compare every AI conceptualization category with every other category. The models use sum/effects coding for `ai_kind_choice`, so no conceptualization group is treated as the reference group; the pairwise `emmeans()` results are the main comparisons to interpret.

Only 8 participants selected “Writing or print” overall, and only 6 of those participants used the AI assistant. Because that group is too small for stable comparisons, I exclude it from Analysis 1. This means the pairwise comparisons below include the remaining four categories.

```
library(emmeans)
```

```
## Warning: package 'emmeans' was built under R version 4.4.3
```

```
## Welcome to emmeans.
```

```
## Caution: You lose important information if you filter this package's results.
```

```
## See '? untidy'
```

```
concept_data <- subset(model_data, ai_kind_choice != "Writing or print")
concept_data$ai_kind_choice <- droplevels(concept_data$ai_kind_choice)
```

```
# Use sum/effects coding so no AI conceptualization category is treated as
# the reference category in the model parameterization.
```

```
contrasts(concept_data$ai_kind_choice) <- contr.sum(
  nlevels(concept_data$ai_kind_choice)
)
```

```
concept_ai_users <- subset(concept_data, ai_used == TRUE)
concept_ai_users$ai_kind_choice <- droplevels(concept_ai_users$ai_kind_choice)
contrasts(concept_ai_users$ai_kind_choice) <- contr.sum(
  nlevels(concept_ai_users$ai_kind_choice)
)
```

```
table(concept_data$ai_kind_choice, useNA = "ifany")
```

```
##
## A calculator or specialized tool      A colleague or advisor
##                                30                                71
##      A library or search engine      A market or aggregator
##                                104                               14
```

```
table(concept_ai_users$ai_kind_choice, useNA = "ifany")
```

```
##
## A calculator or specialized tool      A colleague or advisor
##                                20                                56
##      A library or search engine      A market or aggregator
##                                67                                9
```

Analysis 1A: AI Conceptualization Predicting Whether Participants Used AI

Outcome: `ai_used`

Model type: mixed logistic regression

Predictor: ai_kind_choice
Random intercept: task_id

This model asks whether people who conceptualized AI differently were more or less likely to consult the optional AI assistant during the task. The model uses sum/effects coding so no category is treated as the reference category. The `drop1()` result tests whether ai_kind_choice as a whole improves the model. The `emmeans()` results show estimated AI-use probabilities for each conceptualization group, and the pairwise comparisons show Tukey-adjusted differences between groups.

```
model_ai_kind_ai_used <- glmer(  
  ai_used ~ ai_kind_choice + (1 | task_id),  
  data = concept_data,  
  family = binomial,  
  control = glmerControl(optimizer = "bobyqa")  
)
```

```
summary(model_ai_kind_ai_used)
```

```
## Generalized linear mixed model fit by maximum likelihood (Laplace  
## Approximation) [glmerMod]  
## Family: binomial ( logit )  
## Formula: ai_used ~ ai_kind_choice + (1 | task_id)  
## Data: concept_data  
## Control: glmerControl(optimizer = "bobyqa")  
##  
##           AIC          BIC      logLik -2*log(L)  df.resid  
##      275.1       292.0     -132.5     265.1      214  
##  
## Scaled residuals:  
##      Min       1Q   Median       3Q      Max  
## -1.9513 -1.3413  0.5183  0.7401  0.7495  
##  
## Random effects:  
## Groups Name          Variance Std.Dev.  
## task_id (Intercept) 0.006044 0.07774  
## Number of obs: 219, groups: task_id, 18  
##  
## Fixed effects:  
##              Estimate Std. Error z value Pr(>|z|)  
## (Intercept)      0.8005     0.2002   3.999 6.37e-05 ***  
## ai_kind_choice1 -0.1061     0.3356  -0.316  0.7520  
## ai_kind_choice2  0.5198     0.2815   1.846  0.0649 .  
## ai_kind_choice3 -0.2066     0.2462  -0.839  0.4013  
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
##  
## Correlation of Fixed Effects:  
##              (Intr) a_kn_1 a_kn_2  
## ai_knd_chc1 -0.009  
## ai_knd_chc2 -0.268 -0.235  
## ai_knd_chc3 -0.591 -0.120  0.066
```

```
tests_ai_kind_ai_used <- drop1(model_ai_kind_ai_used, test = "Chisq")
```

```
## boundary (singular) fit: see help('isSingular')
```

```
tests_ai_kind_ai_used
```

```
## Single term deletions
##
## Model:
## ai_used ~ ai_kind_choice + (1 | task_id)
##               npar      AIC      LRT Pr(Chi)
## <none>                275.05
## ai_kind_choice      3 273.73 4.6708 0.1976
```

```
emm_ai_kind_ai_used <- emmeans(
  model_ai_kind_ai_used,
  ~ ai_kind_choice,
  type = "response"
)
```

```
emm_ai_kind_ai_used
```

```
## ai_kind_choice      prob      SE  df asymp.LCL asymp.UCL
## A calculator or specialized tool 0.667 0.0865 Inf      0.483      0.811
## A colleague or advisor          0.789 0.0497 Inf      0.676      0.870
## A library or search engine       0.644 0.0472 Inf      0.547      0.731
## A market or aggregator          0.644 0.1310 Inf      0.371      0.847
##
## Confidence level used: 0.95
## Intervals are back-transformed from the logit scale
```

```
pairs_ai_kind_ai_used <- pairs(
  emm_ai_kind_ai_used,
  adjust = "tukey"
)
```

```
summary(pairs_ai_kind_ai_used, type = "response")
```

```
## contrast                                odds.ratio      SE
## A calculator or specialized tool / A colleague or advisor      0.535 0.260
## A calculator or specialized tool / A library or search engine   1.106 0.486
## A calculator or specialized tool / A market or aggregator      1.106 0.759
## A colleague or advisor / A library or search engine            2.068 0.747
## A colleague or advisor / A market or aggregator                2.068 1.310
## A library or search engine / A market or aggregator            1.000 0.607
## df null z.ratio p.value
## Inf  1 -1.287 0.5711
## Inf  1  0.229 0.9958
## Inf  1  0.147 0.9989
## Inf  1  2.009 0.1844
## Inf  1  1.149 0.6588
```

```
## Inf      1    0.001  1.0000
##
## P value adjustment: tukey method for comparing a family of 4 estimates
## Tests are performed on the log odds ratio scale
```

The overall effect of AI-kind was not significant, $\chi^2(3) = 4.67$, $p = .198$. The colleague/advisor category had the highest estimated AI-use probability (.79), followed by calculator/tool (.67), library/search engine (.64), and market/aggregator (.64), but Tukey-adjusted pairwise comparisons were not significant. Overall, AI conceptualization did not reliably predict AI use, though descriptively AI use was highest when AI was seen as a colleague/advisor.

The mixed model produced a singular-fit warning, with the random intercept variance for task_id estimated at zero. This suggests that the model found little evidence of residual variation between task prompts for this outcome, or that the random effect was too weak to estimate reliably. As a robustness check, I fit a simpler logistic regression without (1 | task_id) to test whether the substantive conclusion changed.

```
model_ai_kind_ai_used_glm <- glm(
  ai_used ~ ai_kind_choice,
  data = concept_data,
  family = binomial
)

summary(model_ai_kind_ai_used_glm)

##
## Call:
## glm(formula = ai_used ~ ai_kind_choice, family = binomial, data = concept_data)
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    0.7980     0.1916   4.164 3.12e-05 ***
## ai_kind_choice1 -0.1049     0.3343  -0.314  0.7537
## ai_kind_choice2  0.5193     0.2810   1.848  0.0646 .
## ai_kind_choice3 -0.2042     0.2402  -0.850  0.3952
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 269.72  on 218  degrees of freedom
## Residual deviance: 265.06  on 215  degrees of freedom
## AIC: 273.06
##
## Number of Fisher Scoring iterations: 4
```

```
tests_ai_kind_ai_used_glm <- drop1(
  model_ai_kind_ai_used_glm,
  test = "Chisq"
)

tests_ai_kind_ai_used_glm
```

```
## Single term deletions
```



```
##
## Model:
## ai_used ~ ai_kind_choice
##           Df Deviance    AIC    LRT Pr(>Chi)
## <none>           265.06 273.06
## ai_kind_choice  3   269.73 271.73 4.6687  0.1977
```

```
emm_ai_kind_ai_used_glm <- emmeans(
  model_ai_kind_ai_used_glm,
  ~ ai_kind_choice,
  type = "response"
)
```

```
emm_ai_kind_ai_used_glm
```

```
## ai_kind_choice          prob      SE  df asymp.LCL asymp.UCL
## A calculator or specialized tool 0.667 0.0861 Inf    0.484    0.810
## A colleague or advisor           0.789 0.0484 Inf    0.679    0.868
## A library or search engine        0.644 0.0469 Inf    0.548    0.730
## A market or aggregator           0.643 0.1280 Inf    0.376    0.843
##
## Confidence level used: 0.95
## Intervals are back-transformed from the logit scale
```

```
pairs_ai_kind_ai_used_glm <- pairs(
  emm_ai_kind_ai_used_glm,
  adjust = "tukey"
)
```

```
summary(pairs_ai_kind_ai_used_glm, type = "response")
```

```
## contrast                                odds.ratio    SE
## A calculator or specialized tool / A colleague or advisor      0.536 0.259
## A calculator or specialized tool / A library or search engine  1.104 0.484
## A calculator or specialized tool / A market or aggregator      1.111 0.755
## A colleague or advisor / A library or search engine            2.062 0.733
## A colleague or advisor / A market or aggregator               2.074 1.300
## A library or search engine / A market or aggregator           1.006 0.598
## df null z.ratio p.value
## Inf    1  -1.289 0.5700
## Inf    1   0.227 0.9959
## Inf    1   0.155 0.9987
## Inf    1   2.034 0.1753
## Inf    1   1.160 0.6523
## Inf    1   0.010 1.0000
##
## P value adjustment: tukey method for comparing a family of 4 estimates
## Tests are performed on the log odds ratio scale
```

AI conceptualization did not significantly predict optional assistant use, $\chi^2(3) = 4.67$, $p = .198$. Participants who viewed AI as a colleague/advisor showed the highest estimated AI-use probability (.79), followed by calculator/tool (.67), library/search engine (.64), and market/aggregator (.64), but pairwise differences were not statistically reliable after Tukey adjustment.

```

# AI use rate by conceptualization (all participants including Writing or print)
ai_kind_all <- subset(model_data, !is.na(ai_kind_choice))
ai_kind_use_df <- do.call(rbind, lapply(
  split(ai_kind_all, ai_kind_all$ai_kind_choice),
  function(sub) {
    data.frame(
      ai_kind_choice = as.character(sub$ai_kind_choice[1]),
      pct_used       = mean(sub$ai_used, na.rm = TRUE),
      n              = nrow(sub),
      stringsAsFactors = FALSE
    )
  }
))


kind_lookup2 <- c(
  "A calculator or specialized tool" = "Calculator /\nspecialized tool",
  "A colleague or advisor"          = "Colleague /\nadvisor",
  "A library or search engine"       = "Library /\nsearch engine",
  "A market or aggregator"           = "Market /\naggregator",
  "Writing or print"                 = "Writing or\nprint"
)

ai_kind_use_df$kind_short <- kind_lookup2[ai_kind_use_df$ai_kind_choice]
ai_kind_use_df$kind_short <- factor(
  ai_kind_use_df$kind_short,
  levels = ai_kind_use_df$kind_short[order(ai_kind_use_df$pct_used)]
)

ggplot(ai_kind_use_df, aes(x = kind_short, y = pct_used)) +
  geom_col(fill = "#4472C4", width = 0.65) +
  geom_text(
    aes(label = paste0(round(pct_used * 100), "%")),
    hjust = -0.12, size = 3.5, fontface = "bold"
  ) +
  geom_text(
    aes(label = paste0("n=", n), y = 0.01),
    hjust = 0, size = 3, color = "white"
  ) +
  coord_flip() +
  scale_y_continuous(
    limits = c(0, 1.05), expand = c(0, 0),
    labels = function(x) paste0(round(x * 100), "%")
  ) +
  labs(
    title     = "AI Use Rate by AI Conceptualization",
    subtitle  = "% of participants who consulted the optional AI assistant",
    x         = NULL,
    y         = "% who used AI assistant"
  ) +
  theme_classic(base_size = 12) +
  theme(
    plot.title     = element_text(face = "bold", size = 12),
    plot.subtitle  = element_text(size = 10),
    axis.text.y    = element_text(size = 10.5)
  )

```

)



Analysis 1B: AI Conceptualization Predicting Similarity To The AI Response

Outcome: `jaccard_similarity_to_ai`
 Model type: linear mixed regression
 Predictor: `ai_kind_choice`
 Random intercept: `task_id`
 Sample: only participants who used the AI assistant

This model asks whether people's conceptualization of AI predicted how much their final response resembled the pre-written AI response. The `emmeans()` results show the adjusted average similarity score for each conceptualization group, and the pairwise comparisons test whether any groups differ from one another.

```
model_ai_kind_similarity <- lmer(
  jaccard_similarity_to_ai ~ ai_kind_choice + (1 | task_id),
  data = concept_ai_users
)
```

```
summary(model_ai_kind_similarity)
```

```
## Linear mixed model fit by REML. t-tests use Satterthwaite's method [
## lmerModLmerTest]
## Formula: jaccard_similarity_to_ai ~ ai_kind_choice + (1 | task_id)
## Data: concept_ai_users
##
## REML criterion at convergence: 179.3
##
## Scaled residuals:
##    Min      1Q  Median      3Q     Max
## -1.6964 -0.8675 -0.3148  0.9684  1.5818
##
```

```
## Random effects:
## Groups Name Variance Std.Dev.
## task_id (Intercept) 0.01529 0.1236
## Residual 0.16519 0.4064
## Number of obs: 152, groups: task_id, 18
##
## Fixed effects:
## Estimate Std. Error df t value Pr(>|t|)
## (Intercept) 0.52870 0.05390 33.22106 9.809 2.44e-11 ***
## ai_kind_choice1 0.05266 0.08015 144.22774 0.657 0.512
## ai_kind_choice2 0.02383 0.06073 146.51328 0.392 0.695
## ai_kind_choice3 -0.06049 0.05827 144.76230 -1.038 0.301
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Correlation of Fixed Effects:
## (Intr) a_kn_1 a_kn_2
## ai_knd_chc1 0.031
## ai_knd_chc2 -0.405 -0.175
## ai_knd_chc3 -0.455 -0.163 0.179
```

```
anova(model_ai_kind_similarity)
```

```
## Type III Analysis of Variance Table with Satterthwaite's method
## Sum Sq Mean Sq NumDF DenDF F value Pr(>F)
## ai_kind_choice 0.29269 0.097562 3 144.87 0.5906 0.6221
```

```
emm_ai_kind_similarity <- emmeans(
  model_ai_kind_similarity,
  ~ ai_kind_choice
)
```

```
## Cannot use mode = "kenward-roger" because *pbkrtest* package is not installed
```

```
emm_ai_kind_similarity
```

```
## ai_kind_choice emmean SE df lower.CL upper.CL
## A calculator or specialized tool 0.581 0.0979 117.7 0.387 0.775
## A colleague or advisor 0.553 0.0628 50.9 0.427 0.679
## A library or search engine 0.468 0.0587 41.4 0.350 0.587
## A market or aggregator 0.513 0.1400 147.3 0.235 0.790
##
## Degrees-of-freedom method: satterthwaite
## Confidence level used: 0.95
```

```
pairs_ai_kind_similarity <- pairs(
  emm_ai_kind_similarity,
  adjust = "tukey"
)
```

```
pairs_ai_kind_similarity
```

```
## contrast
## A calculator or specialized tool - A colleague or advisor      estimate    SE
## A calculator or specialized tool - A library or search engine  0.1131 0.1060
## A calculator or specialized tool - A market or aggregator      0.0687 0.1660
## A colleague or advisor - A library or search engine            0.0843 0.0763
## A colleague or advisor - A market or aggregator               0.0398 0.1490
## A library or search engine - A market or aggregator           -0.0445 0.1470
## df t.ratio p.value
## 145  0.265  0.9934
## 146  1.063  0.7127
## 142  0.413  0.9761
## 148  1.106  0.6866
## 143  0.267  0.9933
## 141 -0.304  0.9902
##
## Degrees-of-freedom method: satterthwaite
## P value adjustment: tukey method for comparing a family of 4 estimates
```

AI conceptualization did not significantly predict Jaccard similarity to the AI response (see ANOVA output above). Descriptively, mean Jaccard similarity ranged from .45 (Library/search engine) to .57 (Calculator/tool) across the four main groups, but differences were small and unreliable. Tukey-adjusted pairwise comparisons showed no significant differences between conceptualization groups.

```
# Include all AI users (including Writing or print) for visualization only.
jaccard_viz_data <- subset(
  ai_users,
  !is.na(ai_kind_choice) & !is.na(jaccard_similarity_to_ai)
)
jaccard_viz_data$ai_kind_choice <- droplevels(jaccard_viz_data$ai_kind_choice)

kind_lookup <- c(
  "A calculator or specialized tool" = "Calculator /\nspecialized tool",
  "A colleague or advisor"          = "Colleague /\nadvisor",
  "A library or search engine"       = "Library /\nsearch engine",
  "A market or aggregator"          = "Market /\naggregator",
  "Writing or print"                = "Writing or\nprint"
)


library(ggplot2)
# Compute group means and n for bar chart
jac_kind_df <- do.call(rbind, lapply(
  split(jaccard_viz_data, jaccard_viz_data$ai_kind_choice),
  function(sub) {
    data.frame(
      ai_kind_choice = as.character(sub$ai_kind_choice[1]),
      mean_jac       = mean(sub$jaccard_similarity_to_ai, na.rm = TRUE),
      n              = sum(!is.na(sub$jaccard_similarity_to_ai)),
      stringsAsFactors = FALSE
    )
  }
))
jac_kind_df$kind_short <- kind_lookup[jac_kind_df$ai_kind_choice]
jac_kind_df$kind_short <- factor(
  jac_kind_df$kind_short,
```

```

levels = jac_kind_df$kind_short[order(jac_kind_df$mean_jac)]
)

ggplot(jac_kind_df, aes(x = kind_short, y = mean_jac)) +
  geom_col(fill = "#4472C4", width = 0.65) +
  geom_text(
    aes(label = sprintf("%.3f", mean_jac)),
    hjust = -0.12, size = 3.5, fontface = "bold"
  ) +
  geom_text(
    aes(label = paste0("n=", n), y = 0.012),
    hjust = 0, size = 3, color = "white"
  ) +
  coord_flip() +
  scale_y_continuous(limits = c(0, 0.74), expand = c(0, 0)) +
  labs(
    title = "Mean Jaccard Similarity to AI Response by AI Conceptualization",
    x = NULL,
    y = "Mean Jaccard Similarity"
  ) +
  theme_classic(base_size = 12) +
  theme(
    plot.title = element_text(face = "bold", size = 12),
    plot.subtitle = element_text(size = 10),
    axis.text.y = element_text(size = 10.5)
  )
)

```



Analysis 1C: AI Conceptualization Predicting Epistemic Ownership

Outcome: `epistemic_ownership_binary` (0 = AI's thinking, 1 = own thinking)

Model type: mixed logistic regression

Predictor: `ai_kind_choice`

Random intercept: task_id

Sample: only participants who used the AI assistant

This model asks whether people's conceptualization of AI predicted whether they rated the final response as closer to AI thinking or their own thinking. Epistemic ownership is recoded as binary (see Model 3).

```
model_ai_kind_ownership <- glmer(  
  epistemic_ownership_binary ~ ai_kind_choice + (1 | task_id),  
  data = concept_ai_users,  
  family = binomial,  
  control = glmerControl(optimizer = "bobyqa")  
)
```

```
## boundary (singular) fit: see help('isSingular')
```

```
summary(model_ai_kind_ownership)
```

```
## Generalized linear mixed model fit by maximum likelihood (Laplace  
## Approximation) [glmerMod]  
## Family: binomial ( logit )  
## Formula: epistemic_ownership_binary ~ ai_kind_choice + (1 | task_id)  
## Data: concept_ai_users  
## Control: glmerControl(optimizer = "bobyqa")  
##  
##           AIC          BIC      logLik -2*log(L)  df.resid  
##      181.8       196.5      -85.9     171.8      133  
##  
## Scaled residuals:  
##      Min       1Q   Median       3Q      Max  
## -0.7746 -0.7171 -0.6172   1.3093   1.6202  
##  
## Random effects:  
## Groups Name          Variance Std.Dev.  
## task_id (Intercept) 0          0  
## Number of obs: 138, groups: task_id, 18  
##  
## Fixed effects:  
##              Estimate Std. Error z value Pr(>|z|)  
## (Intercept)   -0.669970  0.241086  -2.779  0.00545 **  
## ai_kind_choice1  0.130973  0.413785   0.317  0.75160  
## ai_kind_choice2  0.004994  0.316522   0.016  0.98741  
## ai_kind_choice3 -0.295111  0.318241  -0.927  0.35376  
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
##  
## Correlation of Fixed Effects:  
##              (Intr) a_kn_1 a_kn_2  
## ai_knd_chc1 -0.016  
## ai_knd_chc2 -0.486 -0.149  
## ai_knd_chc3 -0.476 -0.152  0.154  
## optimizer (bobyqa) convergence code: 0 (OK)  
## boundary (singular) fit: see help('isSingular')
```

```
drop1(model_ai_kind_ownership, test = "Chisq")
```

```
## boundary (singular) fit: see help('isSingular')
```

```
## Single term deletions
```

```
##
```

```
## Model:
```

```
## epistemic_ownership_binary ~ ai_kind_choice + (1 | task_id)
```

```
##          npar      AIC      LRT Pr(Chi)
```

```
## <none>          181.84
```

```
## ai_kind_choice    3 176.77 0.93322 0.8174
```

```
emm_ai_kind_ownership <- emmeans(
  model_ai_kind_ownership,
  ~ ai_kind_choice,
  type = "response"
)
```

```
emm_ai_kind_ownership
```

```
## ai_kind_choice      prob      SE  df asymp.LCL asymp.UCL
## A calculator or specialized tool 0.368 0.1110 Inf    0.187    0.597
## A colleague or advisor          0.340 0.0651 Inf    0.226    0.476
## A library or search engine       0.276 0.0587 Inf    0.176    0.404
## A market or aggregator          0.375 0.1710 Inf    0.125    0.715
```

```
##
```

```
## Confidence level used: 0.95
```

```
## Intervals are back-transformed from the logit scale
```

```
pairs_ai_kind_ownership <- pairs(
  emm_ai_kind_ownership,
  adjust = "tukey"
)
```

```
summary(pairs_ai_kind_ownership, type = "response")
```

```
## contrast                                odds.ratio      SE
## A calculator or specialized tool / A colleague or advisor      1.134 0.632
## A calculator or specialized tool / A library or search engine  1.531 0.856
## A calculator or specialized tool / A market or aggregator      0.972 0.847
## A colleague or advisor / A library or search engine            1.350 0.557
## A colleague or advisor / A market or aggregator                0.857 0.674
## A library or search engine / A market or aggregator            0.635 0.500
```

```
## df null z.ratio p.value
```

```
## Inf 1 0.226 0.9959
```

```
## Inf 1 0.762 0.8714
```

```
## Inf 1 -0.032 1.0000
```

```
## Inf 1 0.727 0.8864
```

```
## Inf 1 -0.196 0.9973
```

```
## Inf 1 -0.577 0.9390
```

```
##
```

```
## P value adjustment: tukey method for comparing a family of 4 estimates
```

```
## Tests are performed on the log odds ratio scale
```


AI conceptualization did not significantly predict binary epistemic ownership, $\chi^2(3) = 0.93$, $p = .817$. Estimated probabilities of primarily own-thinking ownership ranged from .28 to .38 across conceptualization groups, and Tukey-adjusted pairwise comparisons showed no significant differences.

Altogether: Conceptualizing AI differently did not significantly predict AI use, similarity to the AI response, or epistemic ownership. Across all three outcomes, the colleague/advisor group showed descriptively higher values, but none of these patterns reached statistical significance. This suggests that participants' explicit mental model of AI was less predictive of cognitive offloading behavior than their actual task-level prompting behavior.

Analysis 2: Learning-Curve Analysis

This analysis is inspired by the Anthropic Economic Index learning-curve report. The idea is that more experienced AI users may use AI differently over time: they may consult it differently, copy it differently, or become more aware of when it has shaped their thinking.

Predictors:

- **ai_use_tenure:** how long the participant has been using AI assistants regularly
- **z_hours_per_week:** approximate weekly AI use (z-scored)
- **ai_understanding:** self-reported understanding of how AI assistants work

Random intercept:

- **task_id**, because participants responded to different short-task prompts

```
learning_curve_data <- subset(
  model_data,
  !is.na(ai_used) &
    !is.na(ai_use_tenure) &
    !is.na(hours_per_week) &
    !is.na(ai_understanding) &
    !is.na(task_id)
)

learning_curve_ai_users <- subset(
  learning_curve_data,
  ai_used == TRUE &
    !is.na(jaccard_similarity_to_ai) &
    !is.na(epistemic_ownership)
)

table(learning_curve_data$ai_use_tenure, useNA = "ifany")
```

```
##
## I do not use them, or am just starting out.
##          9
##          Less than 6 months
##          20
##          6 to 12 months
##          49
##          1 to 2 years
##          89
```

```
## More than 2 years
## 60

table(learning_curve_data$ai_understanding, useNA = "ifany")

##
## I have a general sense. It learns from a lot of text and generates responses.
## 83
## I just use it. I do not think much about how it works.
## 14
## I understand it reasonably well. I know about training, prompting, and limitations.
## 118
## I understand it technically. I work with or study AI systems.
## 12
```

Analysis 2A: Do More Experienced Users Consult AI Differently?

Outcome: ai_used

Model type: mixed logistic regression

Predictors: ai_use_tenure, z_hours_per_week, ai_understanding

Random intercept: task_id

This model asks whether more experienced participants were more or less likely to click/use the optional AI assistant during the short task.

```
model_learning_ai_used <- glmer(
  ai_used ~ ai_use_tenure +
    z_hours_per_week +
    ai_understanding +
    (1 | task_id),
  data = learning_curve_data,
  family = binomial,
  control = glmerControl(optimizer = "bobyqa")
)
```

```
## boundary (singular) fit: see help('isSingular')
```

```
summary(model_learning_ai_used)
```

```
## Generalized linear mixed model fit by maximum likelihood (Laplace
## Approximation) [glmerMod]
## Family: binomial ( logit )
## Formula: ai_used ~ ai_use_tenure + z_hours_per_week + ai_understanding +
## (1 | task_id)
## Data: learning_curve_data
## Control: glmerControl(optimizer = "bobyqa")
##
## AIC      BIC    logLik -2*log(L)  df.resid
## 279.8    314.0   -129.9    259.8     217
##
## Scaled residuals:
## Min      1Q  Median      3Q      Max
```

```

## -2.4476 -1.3146 0.5406 0.6718 2.9131
##
## Random effects:
## Groups Name Variance Std.Dev.
## task_id (Intercept) 0 0
## Number of obs: 227, groups: task_id, 18
##
## Fixed effects:
##
## (Intercept)
## ai_use_tenure.L
## ai_use_tenure.Q
## ai_use_tenure.C
## ai_use_tenure^4
## z_hours_per_week
## ai_understandingI just use it. I do not think much about how it works.
## ai_understandingI understand it reasonably well. I know about training, prompting, and limitations.
## ai_understandingI understand it technically. I work with or study AI systems.
##
## (Intercept)
## ai_use_tenure.L
## ai_use_tenure.Q
## ai_use_tenure.C
## ai_use_tenure^4
## z_hours_per_week
## ai_understandingI just use it. I do not think much about how it works.
## ai_understandingI understand it reasonably well. I know about training, prompting, and limitations.
## ai_understandingI understand it technically. I work with or study AI systems.
##
## (Intercept)
## ai_use_tenure.L
## ai_use_tenure.Q
## ai_use_tenure.C
## ai_use_tenure^4
## z_hours_per_week
## ai_understandingI just use it. I do not think much about how it works.
## ai_understandingI understand it reasonably well. I know about training, prompting, and limitations.
## ai_understandingI understand it technically. I work with or study AI systems.
##
## (Intercept)
## ai_use_tenure.L
## ai_use_tenure.Q
## ai_use_tenure.C
## ai_use_tenure^4
## z_hours_per_week
## ai_understandingI just use it. I do not think much about how it works.
## ai_understandingI understand it reasonably well. I know about training, prompting, and limitations.
## ai_understandingI understand it technically. I work with or study AI systems.
##
## (Intercept)
## ai_use_tenure.L
## ai_use_tenure.Q
## ai_use_tenure.C
## ai_use_tenure^4

```

```
## z_hours_per_week
## ai_understandingI just use it. I do not think much about how it works.
## ai_understandingI understand it reasonably well. I know about training, prompting, and limitations.
## ai_understandingI understand it technically. I work with or study AI systems.
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Correlation of Fixed Effects:
##      (Intr) a_s_.L a_s_.Q a_s_.C a_s.^4 z_hr__ ajuiIdntmahiw
## ai_us_tnr.L   -0.601
## ai_us_tnr.Q    0.553 -0.711
## ai_us_tnr.C   -0.322  0.529 -0.661
## ai_us_tnr^4    0.128 -0.180  0.327 -0.523
## z_hrs_pr_wk    0.222 -0.163  0.042 -0.067  0.017
## ajuiIdntmahiw -0.316  0.026 -0.083  0.064  0.028 -0.102
## auirWkatpal   -0.590 -0.018 -0.055  0.095  0.013 -0.125  0.299
## auitIwwosAs   -0.240 -0.038 -0.026  0.032 -0.044 -0.145  0.129
##
##      auirWkatpal
## ai_us_tnr.L
## ai_us_tnr.Q
## ai_us_tnr.C
## ai_us_tnr^4
## z_hrs_pr_wk
## ajuiIdntmahiw
## auirWkatpal
## auitIwwosAs    0.257
## optimizer (bobyqa) convergence code: 0 (OK)
## boundary (singular) fit: see help('isSingular')
```

```
tests_learning_ai_used <- drop1(
  model_learning_ai_used,
  test = "Chisq"
)
```

```
## boundary (singular) fit: see help('isSingular')
## boundary (singular) fit: see help('isSingular')
## boundary (singular) fit: see help('isSingular')
```

```
tests_learning_ai_used
```

```
## Single term deletions
##
## Model:
## ai_used ~ ai_use_tenure + z_hours_per_week + ai_understanding +
##      (1 | task_id)
##      npar      AIC      LRT Pr(Chi)
## <none>          279.80
## ai_use_tenure    4 285.74 13.9376 0.007497 **
## z_hours_per_week 1 278.89  1.0930 0.295815
## ai_understanding 3 274.48  0.6795 0.878007
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
emm_learning_ai_used_tenure <- emmeans(
  model_learning_ai_used,
  ~ ai_use_tenure,
  type = "response"
)
```

```
emm_learning_ai_used_tenure
```

```
## ai_use_tenure          prob    SE  df asymp.LCL
## I do not use them, or am just starting out. 0.145 0.1350 Inf    0.0196
## Less than 6 months          0.741 0.1050 Inf    0.4945
## 6 to 12 months              0.747 0.0743 Inf    0.5774
## 1 to 2 years                 0.708 0.0656 Inf    0.5657
## More than 2 years           0.813 0.0582 Inf    0.6722
## asymp.UCL
##      0.590
##      0.894
##      0.865
##      0.819
##      0.902
##
## Results are averaged over the levels of: ai_understanding
## Confidence level used: 0.95
## Intervals are back-transformed from the logit scale
```

AI-use tenure significantly predicted use of the optional AI assistant, chi-squared(4) = 13.94, $p = .0075$. Participants who reported not using AI or just starting out had a much lower estimated AI-use probability (.15) than all other tenure groups (.71-.81). This should still be interpreted cautiously because the smallest tenure group produces wide uncertainty. Hours per week, $p = .296$, and AI understanding, $p = .878$, were not significant predictors.

```
model_learning_ai_used_glm <- glm(
  ai_used ~ ai_use_tenure +
    z_hours_per_week +
    ai_understanding,
  data = learning_curve_data,
  family = binomial
)
```

```
summary(model_learning_ai_used_glm)
```

```
##
## Call:
## glm(formula = ai_used ~ ai_use_tenure + z_hours_per_week + ai_understanding,
##      family = binomial, data = learning_curve_data)
##
## Coefficients:
##
## (Intercept)
## ai_use_tenure.L
## ai_use_tenure.Q
## ai_use_tenure.C
```

```

## ai_use_tenure^4
## z_hours_per_week
## ai_understandingI just use it. I do not think much about how it works.
## ai_understandingI understand it reasonably well. I know about training, prompting, and limitations.
## ai_understandingI understand it technically. I work with or study AI systems.
##
## (Intercept)
## ai_use_tenure.L
## ai_use_tenure.Q
## ai_use_tenure.C
## ai_use_tenure^4
## z_hours_per_week
## ai_understandingI just use it. I do not think much about how it works.
## ai_understandingI understand it reasonably well. I know about training, prompting, and limitations.
## ai_understandingI understand it technically. I work with or study AI systems.
##
## (Intercept)
## ai_use_tenure.L
## ai_use_tenure.Q
## ai_use_tenure.C
## ai_use_tenure^4
## z_hours_per_week
## ai_understandingI just use it. I do not think much about how it works.
## ai_understandingI understand it reasonably well. I know about training, prompting, and limitations.
## ai_understandingI understand it technically. I work with or study AI systems.
##
## (Intercept)
## ai_use_tenure.L
## ai_use_tenure.Q
## ai_use_tenure.C
## ai_use_tenure^4
## z_hours_per_week
## ai_understandingI just use it. I do not think much about how it works.
## ai_understandingI understand it reasonably well. I know about training, prompting, and limitations.
## ai_understandingI understand it technically. I work with or study AI systems.
##
## (Intercept)
## ai_use_tenure.L
## ai_use_tenure.Q
## ai_use_tenure.C
## ai_use_tenure^4
## z_hours_per_week
## ai_understandingI just use it. I do not think much about how it works.
## ai_understandingI understand it reasonably well. I know about training, prompting, and limitations.
## ai_understandingI understand it technically. I work with or study AI systems.
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 278.84  on 226  degrees of freedom
## Residual deviance: 259.80  on 218  degrees of freedom
## AIC: 277.8
##

```

```
## Number of Fisher Scoring iterations: 4
```

```
tests_learning_ai_used_glm <- drop1(  
  model_learning_ai_used_glm,  
  test = "Chisq"  
)
```

```
tests_learning_ai_used_glm
```

```
## Single term deletions
```

```
##
```

```
## Model:
```

```
## ai_used ~ ai_use_tenure + z_hours_per_week + ai_understanding
```

```
##           Df Deviance   AIC    LRT Pr(>Chi)
```

```
## <none>           259.80 277.80
```

```
## ai_use_tenure     4    273.74 283.74 13.9376 0.007497 **
```

```
## z_hours_per_week  1    260.89 276.89  1.0930 0.295815
```

```
## ai_understanding  3    260.48 272.48  0.6795 0.878007
```

```
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
emm_learning_ai_used_tenure_glm <- emmeans(  
  model_learning_ai_used_glm,  
  ~ ai_use_tenure,  
  type = "response"  
)
```

```
emm_learning_ai_used_tenure_glm
```

```
## ai_use_tenure           prob    SE  df asymp.LCL  
## I do not use them, or am just starting out. 0.145 0.1350 Inf    0.0197  
## Less than 6 months          0.741 0.1050 Inf    0.4945  
## 6 to 12 months              0.747 0.0743 Inf    0.5774  
## 1 to 2 years                 0.708 0.0656 Inf    0.5657  
## More than 2 years           0.813 0.0582 Inf    0.6722  
## asymp.UCL  
##      0.590  
##      0.894  
##      0.865  
##      0.819  
##      0.902
```

```
##
```

```
## Results are averaged over the levels of: ai_understanding
```

```
## Confidence level used: 0.95
```

```
## Intervals are back-transformed from the logit scale
```

Same results but the extreme predicted probabilities and very large standard errors suggest that this result should be interpreted cautiously.

```
tenure_levels <- c(  
  "I do not use them, or am just starting out.",  
  "Less than 6 months",
```


```

    "6 to 12 months",
    "1 to 2 years",
    "More than 2 years"
  )
tenure_labels_short <- c(
  "Not using /\njust starting",
  "< 6 months",
  "6 -- 12 months",
  "1 -- 2 years",
  "> 2 years"
)

tenure_use_df <- do.call(rbind, lapply(
  seq_along(tenure_levels),
  function(i) {
    sub <- subset(learning_curve_data, ai_use_tenure == tenure_levels[i])
    if (nrow(sub) == 0) return(NULL)
    data.frame(
      tenure = tenure_labels_short[i],
      pct_used = mean(sub$ai_used, na.rm = TRUE),
      n = nrow(sub),
      stringsAsFactors = FALSE
    )
  }
))
tenure_use_df$tenure <- factor(tenure_use_df$tenure, levels = tenure_labels_short)

ggplot(tenure_use_df, aes(x = tenure, y = pct_used)) +
  geom_col(fill = "#4472C4", width = 0.65) +
  geom_text(
    aes(label = paste0(round(pct_used * 100), "%")),
    vjust = -0.4, size = 3.5, fontface = "bold"
  ) +
  geom_text(
    aes(label = paste0("n=", n), y = 0.02),
    vjust = 0, size = 3, color = "white"
  ) +
  scale_y_continuous(
    limits = c(0, 1.05), expand = c(0, 0),
    labels = function(x) paste0(round(x * 100), "%")
  ) +
  labs(
    title = "AI Use Rate by AI Experience Tenure",
    subtitle = "% of participants who consulted the optional AI assistant",
    x = "AI Use Tenure",
    y = "% Used AI"
  ) +
  theme_classic(base_size = 12) +
  theme(
    plot.title = element_text(face = "bold", size = 12),
    plot.subtitle = element_text(size = 10),
    axis.text.x = element_text(size = 10)
  )

```

Analysis 2B: Do More Experienced Users Copy AI Differently?

Outcome: `jaccard_similarity_to_ai`

Model type: linear mixed regression

Predictors: `ai_use_tenure`, `z_hours_per_week`, `ai_understanding`

Random intercept: `task_id`

Sample: only participants who used the AI assistant

This model asks whether more experienced participants produced final responses that were more or less similar to the pre-written AI response.

```
model_learning_similarity <- lmer(
  jaccard_similarity_to_ai ~ ai_use_tenure +
    z_hours_per_week +
    ai_understanding +
    (1 | task_id),
  data = learning_curve_ai_users
)
```

```
summary(model_learning_similarity)
```

```
## Linear mixed model fit by REML. t-tests use Satterthwaite's method [
## lmerModLmerTest]
## Formula: jaccard_similarity_to_ai ~ ai_use_tenure + z_hours_per_week +
##          ai_understanding + (1 | task_id)
## Data: learning_curve_ai_users
##
## REML criterion at convergence: 192.5
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -1.7644 -0.8611 -0.2303  0.9309  1.6229
##
```

```

## Random effects:
##   Groups   Name      Variance Std.Dev.
## task_id (Intercept) 0.01908  0.1381
## Residual          0.16268  0.4033
## Number of obs: 158, groups: task_id, 18
##
## Fixed effects:
##
## (Intercept)
## ai_use_tenure.L
## ai_use_tenure.Q
## ai_use_tenure.C
## ai_use_tenure^4
## z_hours_per_week
## ai_understandingI just use it. I do not think much about how it works.
## ai_understandingI understand it reasonably well. I know about training, prompting, and limitations.
## ai_understandingI understand it technically. I work with or study AI systems.
##
## (Intercept)
## ai_use_tenure.L
## ai_use_tenure.Q
## ai_use_tenure.C
## ai_use_tenure^4
## z_hours_per_week
## ai_understandingI just use it. I do not think much about how it works.
## ai_understandingI understand it reasonably well. I know about training, prompting, and limitations.
## ai_understandingI understand it technically. I work with or study AI systems.
##
## (Intercept)
## ai_use_tenure.L
## ai_use_tenure.Q
## ai_use_tenure.C
## ai_use_tenure^4
## z_hours_per_week
## ai_understandingI just use it. I do not think much about how it works.
## ai_understandingI understand it reasonably well. I know about training, prompting, and limitations.
## ai_understandingI understand it technically. I work with or study AI systems.
##
## (Intercept)
## ai_use_tenure.L
## ai_use_tenure.Q
## ai_use_tenure.C
## ai_use_tenure^4
## z_hours_per_week
## ai_understandingI just use it. I do not think much about how it works.
## ai_understandingI understand it reasonably well. I know about training, prompting, and limitations.
## ai_understandingI understand it technically. I work with or study AI systems.
##
## (Intercept)
## ai_use_tenure.L
## ai_use_tenure.Q
## ai_use_tenure.C
## ai_use_tenure^4
## z_hours_per_week

```

```

## ai_understandingI just use it. I do not think much about how it works.
## ai_understandingI understand it reasonably well. I know about training, prompting, and limitations.
## ai_understandingI understand it technically. I work with or study AI systems.
##
## (Intercept)
## ai_use_tenure.L
## ai_use_tenure.Q
## ai_use_tenure.C
## ai_use_tenure^4
## z_hours_per_week
## ai_understandingI just use it. I do not think much about how it works.
## ai_understandingI understand it reasonably well. I know about training, prompting, and limitations.
## ai_understandingI understand it technically. I work with or study AI systems.
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Correlation of Fixed Effects:
##              (Intr) a_s_.L a_s_.Q a_s_.C a_s.^4 z_hr__ ajuiIdntmahiw
## ai_us_tnr.L      -0.796
## ai_us_tnr.Q       0.760 -0.908
## ai_us_tnr.C      -0.629  0.780 -0.854
## ai_us_tnr^4       0.365 -0.463  0.538 -0.654
## z_hrs_pr_wk       0.080 -0.094  0.000 -0.063  0.039
## ajuiIdntmahiw    -0.205  0.003 -0.022  0.012  0.047 -0.175
## auirwIkatpal     -0.463  0.079 -0.118  0.123 -0.019 -0.080  0.306
## auitIwosAs       -0.207  0.005 -0.038  0.051 -0.075 -0.113  0.169
##
##              auirwIkatpal
## ai_us_tnr.L
## ai_us_tnr.Q
## ai_us_tnr.C
## ai_us_tnr^4
## z_hrs_pr_wk
## ajuiIdntmahiw
## auirwIkatpal
## auitIwosAs      0.299

```

```

tests_learning_similarity <- anova(model_learning_similarity)

tests_learning_similarity

```

```

## Type III Analysis of Variance Table with Satterthwaite's method
##              Sum Sq  Mean Sq NumDF  DenDF F value Pr(>F)
## ai_use_tenure    0.68531 0.171328     4 143.31  1.0531 0.3821
## z_hours_per_week 0.00276 0.002759     1 143.19  0.0170 0.8966
## ai_understanding 0.27510 0.091700     3 142.26  0.5637 0.6398

```

```

emm_learning_similarity_tenure <- emmeans(
  model_learning_similarity,
  ~ ai_use_tenure
)

```

```

## Cannot use mode = "kenward-roger" because *pbkrtest* package is not installed

```

```
emm_learning_similarity_tenure
```

```
## ai_use_tenure          emmean      SE    df lower.CL
## I do not use them, or am just starting out.  0.119 0.4220 146.4   -0.714
## Less than 6 months          0.527 0.1230 139.6    0.283
## 6 to 12 months              0.630 0.0863  98.3    0.459
## 1 to 2 years                 0.493 0.0734  70.4    0.347
## More than 2 years           0.600 0.0776  80.8    0.446
## upper.CL
##      0.952
##      0.770
##      0.801
##      0.639
##      0.755
##
## Results are averaged over the levels of: ai_understanding
## Degrees-of-freedom method: satterthwaite
## Confidence level used: 0.95
```

AI-use tenure did not significantly predict Jaccard similarity to the AI response (see ANOVA output above). Descriptively, observed mean Jaccard ranged from .45 (1–2 years) to .57 (6–12 months) across tenure groups, with no consistent trend with experience level. Hours of AI use per week and self-rated AI understanding were also not significant predictors.

Analysis 2C: Do More Experienced Users Recognize AI Influence Differently?

Outcome: `epistemic_ownership_binary` (0 = AI's thinking, 1 = own thinking)

Model type: mixed logistic regression

Predictors: `z_jaccard_similarity_to_ai`, `ai_use_tenure`, `z_hours_per_week`, `ai_understanding`

Interaction: `z_jaccard_similarity_to_ai * ai_use_tenure`

Random intercept: `task_id`

Sample: only participants who used the AI assistant

This model asks whether the link between objective uptake and epistemic ownership differs by experience. Epistemic ownership is recoded as binary (see Model 3). The Jaccard predictor is z-scored for comparability.

```
model_learning_ownership <- glmer(
  epistemic_ownership_binary ~ z_jaccard_similarity_to_ai *
    ai_use_tenure +
    z_hours_per_week +
    ai_understanding +
    (1 | task_id),
  data = learning_curve_ai_users,
  family = binomial,
  control = glmerControl(optimizer = "bobyqa")
)
```

```
## fixed-effect model matrix is rank deficient so dropping 1 column / coefficient
```

```
## Warning in checkConv(attr(opt, "derivs"), opt$par, ctrl = control$checkConv, : Model is nearly unidentifiable:
## - Rescale variables?
```

```
summary(model_learning_ownership)
```

```
## Generalized linear mixed model fit by maximum likelihood (Laplace
## Approximation) [glmerMod]
## Family: binomial ( logit )
## Formula:
## epistemic_ownership_binary ~ z_jaccard_similarity_to_ai * ai_use_tenure +
## z_hours_per_week + ai_understanding + (1 | task_id)
## Data: learning_curve_ai_users
## Control: glmerControl(optimizer = "bobyqa")
##
##          AIC          BIC      logLik -2*log(L)  df.resid
##       173.6       215.2      -72.8     145.6      130
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -1.8062 -0.5751 -0.2845  0.7518  3.5317
##
## Random effects:
##  Groups Name      Variance Std.Dev.
## task_id (Intercept) 0.02015  0.142
## Number of obs: 144, groups: task_id, 18
##
## Fixed effects:
##
## (Intercept)
## z_jaccard_similarity_to_ai
## ai_use_tenure.L
## ai_use_tenure.Q
## ai_use_tenure.C
## ai_use_tenure^4
## z_hours_per_week
## ai_understandingI just use it. I do not think much about how it works.
## ai_understandingI understand it reasonably well. I know about training, prompting, and limitations.
## ai_understandingI understand it technically. I work with or study AI systems.
## z_jaccard_similarity_to_ai:ai_use_tenure.L
## z_jaccard_similarity_to_ai:ai_use_tenure.Q
## z_jaccard_similarity_to_ai:ai_use_tenure.C
##
## (Intercept)
## z_jaccard_similarity_to_ai
## ai_use_tenure.L
## ai_use_tenure.Q
## ai_use_tenure.C
## ai_use_tenure^4
## z_hours_per_week
## ai_understandingI just use it. I do not think much about how it works.
## ai_understandingI understand it reasonably well. I know about training, prompting, and limitations.
## ai_understandingI understand it technically. I work with or study AI systems.
## z_jaccard_similarity_to_ai:ai_use_tenure.L
## z_jaccard_similarity_to_ai:ai_use_tenure.Q
## z_jaccard_similarity_to_ai:ai_use_tenure.C
##
```

```

## (Intercept)
## z_jaccard_similarity_to_ai
## ai_use_tenure.L
## ai_use_tenure.Q
## ai_use_tenure.C
## ai_use_tenure^4
## z_hours_per_week
## ai_understandingI just use it. I do not think much about how it works.
## ai_understandingI understand it reasonably well. I know about training, prompting, and limitations.
## ai_understandingI understand it technically. I work with or study AI systems.
## z_jaccard_similarity_to_ai:ai_use_tenure.L
## z_jaccard_similarity_to_ai:ai_use_tenure.Q
## z_jaccard_similarity_to_ai:ai_use_tenure.C
##
## (Intercept)
## z_jaccard_similarity_to_ai
## ai_use_tenure.L
## ai_use_tenure.Q
## ai_use_tenure.C
## ai_use_tenure^4
## z_hours_per_week
## ai_understandingI just use it. I do not think much about how it works.
## ai_understandingI understand it reasonably well. I know about training, prompting, and limitations.
## ai_understandingI understand it technically. I work with or study AI systems.
## z_jaccard_similarity_to_ai:ai_use_tenure.L
## z_jaccard_similarity_to_ai:ai_use_tenure.Q
## z_jaccard_similarity_to_ai:ai_use_tenure.C

##
## Correlation matrix not shown by default, as p = 13 > 12.
## Use print(value, correlation=TRUE) or
##     vcov(value)         if you need it

## fit warnings:
## fixed-effect model matrix is rank deficient so dropping 1 column / coefficient
## optimizer (bobyqa) convergence code: 0 (OK)
## Model is nearly unidentifiable: large eigenvalue ratio
## - Rescale variables?

tests_learning_ownership <- drop1(model_learning_ownership, test = "Chisq")

## fixed-effect model matrix is rank deficient so dropping 1 column / coefficient

## Warning in checkConv(attr(opt, "derivs"), opt$par, ctrl = control$checkConv, : Model is nearly unidentifiable: large eigenvalue ratio
## - Rescale variables?

## fixed-effect model matrix is rank deficient so dropping 1 column / coefficient

## Warning in checkConv(attr(opt, "derivs"), opt$par, ctrl = control$checkConv, :
## unable to evaluate scaled gradient

## Warning in checkConv(attr(opt, "derivs"), opt$par, ctrl = control$checkConv, : Model failed to converge:
## See ?lme4::convergence and ?lme4::troubleshooting.

```

```
## boundary (singular) fit: see help('isSingular')
```

```
tests_learning_ownership
```

```
## Single term deletions
##
## Model:
## epistemic_ownership_binary ~ z_jaccard_similarity_to_ai * ai_use_tenure +
##   z_hours_per_week + ai_understanding + (1 | task_id)
##               npar      AIC      LRT Pr(Chi)
## <none>                173.59
## z_hours_per_week      1 171.96 0.36570 0.5454
## ai_understanding      3 169.28 1.68867 0.6395
## z_jaccard_similarity_to_ai:ai_use_tenure 3 168.17 0.58377 0.9001
```

```
ownership_similarity_slopes_by_tenure <- emtrends(
  model_learning_ownership,
  ~ ai_use_tenure,
  var = "z_jaccard_similarity_to_ai"
)
```

```
ownership_similarity_slopes_by_tenure
```

```
## ai_use_tenure z_jaccard_similarity_to_ai.trend
## I do not use them, or am just starting out. nonEst
## Less than 6 months -1.39
## 6 to 12 months -1.43
## 1 to 2 years -1.15
## More than 2 years -0.98
## SE df asymp.LCL asymp.UCL
## NA NA NA NA
## 0.852 Inf -3.06 0.281
## 0.593 Inf -2.59 -0.262
## 0.391 Inf -1.92 -0.386
## 0.368 Inf -1.70 -0.258
##
## Results are averaged over the levels of: ai_understanding
## Confidence level used: 0.95
```

```
pairs(
  ownership_similarity_slopes_by_tenure,
  adjust = "tukey"
)
```

```
## contrast estimate
## I do not use them, or am just starting out. - Less than 6 months nonEst
## I do not use them, or am just starting out. - 6 to 12 months nonEst
## I do not use them, or am just starting out. - 1 to 2 years nonEst
## I do not use them, or am just starting out. - More than 2 years nonEst
## Less than 6 months - 6 to 12 months 0.0353
## Less than 6 months - 1 to 2 years -0.2380
## Less than 6 months - More than 2 years -0.4095
```

```
## 6 to 12 months - 1 to 2 years -0.2733
## 6 to 12 months - More than 2 years -0.4448
## 1 to 2 years - More than 2 years -0.1715
##      SE  df z.ratio p.value
##      NA  NA      NA      NA
##      NA  NA      NA      NA
##      NA  NA      NA      NA
##      NA  NA      NA      NA
## 1.020 Inf   0.035  1.0000
## 0.931 Inf  -0.256  0.9941
## 0.921 Inf  -0.444  0.9707
## 0.691 Inf  -0.395  0.9790
## 0.658 Inf  -0.676  0.9063
## 0.526 Inf  -0.326  0.9880
##
## Results are averaged over the levels of: ai_understanding
## P value adjustment: tukey method for comparing a family of 4 estimates
```

Higher Jaccard similarity predicted lower odds of reporting primarily own-thinking epistemic ownership, with negative similarity slopes across tenure groups. The similarity-by-tenure interaction was not significant, chi-squared(3) = 0.58, $p = .900$. Hours per week, $p = .545$, and self-reported AI understanding, $p = .640$, were not significant predictors of epistemic ownership.

Altogether: There is a limited experience effect. **AI-use tenure predicted whether participants consulted the optional assistant** (driven by near-zero use in the no-experience group), but experience did not meaningfully predict how much they copied or converged with the AI response, nor did it change the link between uptake and epistemic ownership.

Analysis 3: Task-Category Analysis

This analysis asks whether different kinds of tasks invite different levels of AI use, uptake, ownership, or confidence.

These models use fixed effects rather than random effects because `assigned_category` is the thing I want to compare directly. This means the models estimate differences between the six task categories, rather than treating task category as background variation. Adding a random intercept for `task_id` within category would likely produce singular fits given the small number of observations per scenario.

```
library(emmeans)

task_categories <- c(
  "academic",
  "creative",
  "decisions",
  "info",
  "learning",
  "writing"
)

task_category_data <- subset(
  model_data,
  assigned_category %in% task_categories
)
```



```

task_category_data$assigned_category <- factor(
  task_category_data$assigned_category,
  levels = task_categories
)

task_category_similarity_data <- subset(
  task_category_data,
  ai_used == TRUE & !is.na(jaccard_similarity_to_ai)
)

task_category_ownership_data <- subset(
  task_category_data,
  ai_used == TRUE & !is.na(epistemic_ownership)
)

task_category_confidence_data <- subset(
  task_category_data,
  !is.na(confidence_level)
)

table(task_category_data$assigned_category, useNA = "ifany")

```

```

##
## academic creative decisions info learning writing
##      36      35      36      41      42      37

```

Analysis 3A: Does Task Category Predict Whether Participants Used AI?

Outcome: ai_used

Model type: logistic regression

Predictor: assigned_category

This model asks whether participants were more likely to open/use the optional AI assistant for some task categories than others.

```

binary_category_summary(task_category_data, "ai_used")

```

```

## assigned_category n proportion
## 1 academic 36 0.5833333
## 2 creative 35 0.6571429
## 3 decisions 36 0.6944444
## 4 info 41 0.7073171
## 5 learning 42 0.8571429
## 6 writing 37 0.6486486

```

```

model_task_ai_used <- glm(
  ai_used ~ assigned_category,
  data = task_category_data,
  family = binomial
)

```

```

summary(model_task_ai_used)

```

```
##
## Call:
## glm(formula = ai_used ~ assigned_category, family = binomial,
##      data = task_category_data)
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)      0.3365     0.3381   0.995  0.31959
## assigned_categorycreative  0.3141     0.4910   0.640  0.52235
## assigned_categorydecisions 0.4845     0.4952   0.978  0.32784
## assigned_categoryinfo      0.5459     0.4818   1.133  0.25715
## assigned_categorylearning  1.4553     0.5556   2.619  0.00882 **
## assigned_categorywriting    0.2766     0.4826   0.573  0.56648
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 278.84 on 226 degrees of freedom
## Residual deviance: 270.22 on 221 degrees of freedom
## AIC: 282.22
##
## Number of Fisher Scoring iterations: 4
```

```
tests_task_ai_used <- drop1(
  model_task_ai_used,
  test = "Chisq"
)

tests_task_ai_used
```

```
## Single term deletions
##
## Model:
## ai_used ~ assigned_category
##              Df Deviance   AIC    LRT Pr(>Chi)
## <none>          270.22 282.22
## assigned_category  5   278.84 280.84 8.6245   0.125
```

```
emm_task_ai_used <- emmeans(
  model_task_ai_used,
  ~ assigned_category,
  type = "response"
)

emm_task_ai_used
```

```
## assigned_category prob      SE df asymp.LCL asymp.UCL
## academic          0.583 0.0822 Inf    0.419    0.731
## creative           0.657 0.0802 Inf    0.488    0.794
## decisions          0.694 0.0768 Inf    0.528    0.822
## info              0.707 0.0711 Inf    0.552    0.826
## learning           0.857 0.0540 Inf    0.717    0.934
```

```
## writing          0.649 0.0785 Inf      0.485      0.784
##
## Confidence level used: 0.95
## Intervals are back-transformed from the logit scale
```

```
pairs_task_ai_used <- pairs(
  emm_task_ai_used,
  adjust = "tukey"
)

summary(pairs_task_ai_used, type = "response")
```

```
## contrast          odds.ratio      SE  df null z.ratio p.value
## academic / creative      0.730 0.359 Inf    1  -0.640  0.9880
## academic / decisions      0.616 0.305 Inf    1  -0.978  0.9249
## academic / info          0.579 0.279 Inf    1  -1.133  0.8677
## academic / learning      0.233 0.130 Inf    1  -2.619  0.0926
## academic / writing        0.758 0.366 Inf    1  -0.573  0.9928
## creative / decisions      0.843 0.428 Inf    1  -0.336  0.9994
## creative / info          0.793 0.392 Inf    1  -0.469  0.9972
## creative / learning      0.319 0.181 Inf    1  -2.013  0.3345
## creative / writing        1.038 0.514 Inf    1   0.076  1.0000
## decisions / info          0.940 0.469 Inf    1  -0.123  1.0000
## decisions / learning      0.379 0.216 Inf    1  -1.702  0.5304
## decisions / writing        1.231 0.615 Inf    1   0.416  0.9984
## info / learning          0.403 0.225 Inf    1  -1.627  0.5803
## info / writing            1.309 0.636 Inf    1   0.554  0.9938
## learning / writing        3.250 1.820 Inf    1   2.107  0.2837
##
## P value adjustment: tukey method for comparing a family of 6 estimates
## Tests are performed on the log odds ratio scale
```

Assigned task category did not significantly predict whether participants used the AI assistant, chi-squared(5) = 8.62, $p = .125$. Descriptively, AI use was highest in learning tasks (86%) and lowest in academic tasks (58%), with other categories between 65% and 71%. Tukey-adjusted pairwise comparisons were not significant.

Analysis 3B: Does Task Category Predict Similarity To The AI Response?

Outcome: `jaccard_similarity_to_ai`
 Model type: linear regression
 Predictor: `assigned_category`
 Sample: only participants who used the AI assistant

This model asks whether participants copied, borrowed from, or converged with the AI response more in some task categories than others.

```
continuous_category_summary(
  task_category_similarity_data,
  "jaccard_similarity_to_ai"
)
```

```
##   assigned_category n      mean      sd    median
## 1      academic 21 0.5854038 0.4585318 1.0000000
## 2      creative 23 0.7562262 0.4190480 1.0000000
## 3    decisions 25 0.2957821 0.3802970 0.0900000
## 4         info 29 0.4319651 0.4163731 0.1882353
## 5     learning 36 0.3623819 0.3481275 0.2230777
## 6        writing 24 0.7190525 0.3431838 0.9509569
```

```
model_task_similarity <- lm(
  jaccard_similarity_to_ai ~ assigned_category,
  data = task_category_similarity_data
)
```

```
summary(model_task_similarity)
```

```
##
## Call:
## lm(formula = jaccard_similarity_to_ai ~ assigned_category, data = task_category_similarity_data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.75623 -0.28297 -0.09218  0.28095  0.70422
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      0.58540    0.08552   6.845 1.76e-10 ***
## assigned_categorycreative  0.17082    0.11828   1.444  0.1507
## assigned_categorydecisions -0.28962    0.11600  -2.497  0.0136 *
## assigned_categoryinfo     -0.15344    0.11229  -1.366  0.1738
## assigned_categorylearning -0.22302    0.10761  -2.073  0.0399 *
## assigned_categorywriting   0.13365    0.11710   1.141  0.2555
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.3919 on 152 degrees of freedom
## Multiple R-squared:  0.1666, Adjusted R-squared:  0.1392
## F-statistic: 6.079 on 5 and 152 DF, p-value: 3.677e-05
```

```
tests_task_similarity <- anova(model_task_similarity)
```

```
tests_task_similarity
```

```
## Analysis of Variance Table
##
## Response: jaccard_similarity_to_ai
##           Df Sum Sq Mean Sq F value    Pr(>F)
## assigned_category    5  4.6682  0.93364  6.0792 3.677e-05 ***
## Residuals          152 23.3441  0.15358
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
emm_task_similarity <- emmeans(
  model_task_similarity,
  ~ assigned_category
)

emm_task_similarity
```

```
## assigned_category emmean      SE df lower.CL upper.CL
## academic          0.585 0.0855 152    0.416    0.754
## creative           0.756 0.0817 152    0.595    0.918
## decisions          0.296 0.0784 152    0.141    0.451
## info              0.432 0.0728 152    0.288    0.576
## learning           0.362 0.0653 152    0.233    0.491
## writing            0.719 0.0800 152    0.561    0.877
##
## Confidence level used: 0.95
```

```
pairs_task_similarity <- pairs(
  emm_task_similarity,
  adjust = "tukey"
)

pairs_task_similarity
```

```
## contrast          estimate      SE df t.ratio p.value
## academic - creative -0.1708 0.1180 152  -1.444  0.7000
## academic - decisions  0.2896 0.1160 152   2.497  0.1316
## academic - info       0.1534 0.1120 152   1.366  0.7469
## academic - learning   0.2230 0.1080 152   2.073  0.3073
## academic - writing    -0.1336 0.1170 152  -1.141  0.8633
## creative - decisions  0.4604 0.1130 152   4.067  0.0011
## creative - info       0.3243 0.1090 152   2.963  0.0405
## creative - learning   0.3938 0.1050 152   3.765  0.0032
## creative - writing     0.0372 0.1140 152   0.325  0.9995
## decisions - info     -0.1362 0.1070 152  -1.273  0.7991
## decisions - learning -0.0666 0.1020 152  -0.653  0.9866
## decisions - writing   -0.4233 0.1120 152  -3.779  0.0030
## info - learning       0.0696 0.0978 152   0.712  0.9803
## info - writing        -0.2871 0.1080 152  -2.655  0.0907
## learning - writing    -0.3567 0.1030 152  -3.454  0.0092
##
## P value adjustment: tukey method for comparing a family of 6 estimates
```

Among AI users, assigned task category significantly predicted Jaccard similarity, $F(5, 152) = 6.08$, $p < .001$. Creative ($M = .76$) and writing ($M = .72$) tasks showed the highest similarity to the AI response, while decisions ($M = .30$) and learning ($M = .36$) showed the lowest. Tukey-adjusted comparisons showed that creative tasks were significantly higher than decisions ($p = .001$), info ($p = .041$), and learning ($p = .003$), and writing tasks were significantly higher than decisions ($p = .003$) and learning ($p = .009$).

```
# Mean Jaccard by category
jac_cat_df <- do.call(rbind, lapply(
  split(task_category_similarity_data,
```

```

    task_category_similarity_data$assigned_category),
function(sub) {
  data.frame(
    assigned_category = as.character(sub$assigned_category[1]),
    mean_jac         = mean(sub$jaccard_similarity_to_ai, na.rm = TRUE),
    n                 = sum(!is.na(sub$jaccard_similarity_to_ai)),
    stringsAsFactors = FALSE
  )
}
))

jac_cat_df$assigned_category <- factor(
  jac_cat_df$assigned_category,
  levels = jac_cat_df$assigned_category[order(jac_cat_df$mean_jac)]
)

p_jac <- ggplot(jac_cat_df, aes(x = assigned_category, y = mean_jac)) +
  geom_col(fill = "#4472C4", width = 0.65) +
  geom_text(
    aes(label = sprintf("%.3f", mean_jac)),
    hjust = -0.10, size = 3.2, fontface = "bold"
  ) +
  geom_text(
    aes(label = paste0("n=", n), y = 0.012),
    hjust = 0, size = 3, color = "white"
  ) +
  coord_flip() +
  scale_y_continuous(limits = c(0, 0.98), expand = c(0, 0)) +
  labs(title = "Mean Similarity to AI", x = NULL, y = "Mean Jaccard Similarity") +
  theme_classic(base_size = 11) +
  theme(
    plot.title = element_text(face = "bold", size = 11),
    axis.text.y = element_blank(),
    axis.ticks.y = element_blank()
  )
)

# AI use rate by category, using the same category order
use_cat_df <- do.call(rbind, lapply(
  split(task_category_data, task_category_data$assigned_category),
  function(sub) {
    data.frame(
      assigned_category = as.character(sub$assigned_category[1]),
      pct_used          = mean(sub$ai_used, na.rm = TRUE),
      n                 = nrow(sub),
      stringsAsFactors = FALSE
    )
  }
))

use_cat_df$assigned_category <- factor(
  use_cat_df$assigned_category,
  levels = levels(jac_cat_df$assigned_category)
)


```

```

p_use <- ggplot(use_cat_df, aes(x = assigned_category, y = pct_used)) +
  geom_col(fill = "#70AD47", width = 0.65) +
  geom_text(
    aes(label = paste0(round(pct_used * 100), "%")),
    hjust = -0.10, size = 3.2, fontface = "bold"
  ) +
  geom_text(
    aes(label = paste0("n=", n), y = 0.02),
    hjust = 0, size = 3, color = "white"
  ) +
  coord_flip() +
  scale_y_continuous(
    limits = c(0, 1.1), expand = c(0, 0),
    labels = function(x) paste0(round(x * 100), "%")
  ) +
  labs(title = "AI Use Rate", x = NULL, y = "% Used AI") +
  theme_classic(base_size = 11) +
  theme(
    plot.title = element_text(face = "bold", size = 11)
  )

grid.arrange(
  p_use, p_jac,
  ncol = 2,
  top = grid::textGrob(
    "AI Use Rate & Mean Jaccard Similarity to AI by Task Category",
    gp = grid::gpar(fontface = "bold", fontsize = 13)
  )
)

```



Analysis 3C: Does Task Category Predict Epistemic Ownership?

Outcome: `epistemic_ownership_binary` (0 = AI's thinking, 1 = own thinking)

Model type: logistic regression

Predictor: `assigned_category`

Sample: only participants who used the AI assistant

This model asks whether epistemic ownership ratings differed across task categories. Epistemic ownership

is recoded as binary (see Model 3).

```
binary_category_summary(  
  task_category_ownership_data,  
  "epistemic_ownership_binary"  
)
```

```
##   assigned_category  n proportion  
## 1      academic 20  0.4500000  
## 2      creative 23  0.3478261  
## 3    decisions 24  0.4166667  
## 4         info 23  0.2608696  
## 5      learning 32  0.2500000  
## 6      writing 22  0.3181818
```

```
model_task_ownership <- glm(  
  epistemic_ownership_binary ~ assigned_category,  
  data = task_category_ownership_data,  
  family = binomial  
)
```

```
summary(model_task_ownership)
```

```
##  
## Call:  
## glm(formula = epistemic_ownership_binary ~ assigned_category,  
##      family = binomial, data = task_category_ownership_data)  
##  
## Coefficients:  
##              Estimate Std. Error z value Pr(>|z|)  
## (Intercept)      -0.2007    0.4495  -0.446   0.655  
## assigned_categorycreative -0.4279    0.6274  -0.682   0.495  
## assigned_categorydecisions -0.1358    0.6111  -0.222   0.824  
## assigned_categoryinfo    -0.8408    0.6538  -1.286   0.198  
## assigned_categorylearning -0.8979    0.6072  -1.479   0.139  
## assigned_categorywriting  -0.5615    0.6415  -0.875   0.381  
##  
## (Dispersion parameter for binomial family taken to be 1)  
##  
##      Null deviance: 183.32  on 143  degrees of freedom  
## Residual deviance: 179.76  on 138  degrees of freedom  
##      (14 observations deleted due to missingness)  
## AIC: 191.76  
##  
## Number of Fisher Scoring iterations: 4
```

```
tests_task_ownership <- drop1(model_task_ownership, test = "Chisq")
```

```
tests_task_ownership
```

```
## Single term deletions  
##
```



```
## Model:
## epistemic_ownership_binary ~ assigned_category
##              Df Deviance   AIC    LRT Pr(>Chi)
## <none>              179.76 191.76
## assigned_category  5   183.32 185.32 3.5557    0.615
```

```
emm_task_ownership <- emmeans(
  model_task_ownership,
  ~ assigned_category,
  type = "response"
)

emm_task_ownership
```

```
## assigned_category prob    SE  df asymp.LCL asymp.UCL
## academic          0.450 0.1110 Inf    0.253    0.664
## creative           0.348 0.0993 Inf    0.184    0.557
## decisions          0.417 0.1010 Inf    0.241    0.617
## info              0.261 0.0916 Inf    0.122    0.472
## learning           0.250 0.0765 Inf    0.130    0.426
## writing            0.318 0.0993 Inf    0.160    0.534
##
## Confidence level used: 0.95
## Intervals are back-transformed from the logit scale
```

```
pairs_task_ownership <- pairs(
  emm_task_ownership,
  adjust = "tukey"
)

summary(pairs_task_ownership, type = "response")
```

```
## contrast          odds.ratio    SE  df null z.ratio p.value
## academic / creative      1.534 0.963 Inf    1  0.682 0.9840
## academic / decisions     1.145 0.700 Inf    1  0.222 0.9999
## academic / info          2.318 1.520 Inf    1  1.286 0.7929
## academic / learning      2.455 1.490 Inf    1  1.479 0.6779
## academic / writing        1.753 1.120 Inf    1  0.875 0.9525
## creative / decisions      0.747 0.450 Inf    1 -0.485 0.9967
## creative / info          1.511 0.976 Inf    1  0.639 0.9881
## creative / learning      1.600 0.958 Inf    1  0.785 0.9701
## creative / writing        1.143 0.724 Inf    1  0.211 0.9999
## decisions / info         2.024 1.280 Inf    1  1.119 0.8737
## decisions / learning     2.143 1.250 Inf    1  1.311 0.7793
## decisions / writing       1.531 0.945 Inf    1  0.690 0.9832
## info / learning          1.059 0.663 Inf    1  0.091 1.0000
## info / writing            0.756 0.499 Inf    1 -0.423 0.9983
## learning / writing        0.714 0.438 Inf    1 -0.549 0.9941
##
## P value adjustment: tukey method for comparing a family of 6 estimates
## Tests are performed on the log odds ratio scale
```

Assigned task category did not significantly predict binary epistemic ownership, $\chi^2(5) = 3.56$, $p = .615$. Estimated probabilities of primarily own-thinking ownership varied descriptively from .25 in learning to .45 in academic tasks, but Tukey-adjusted pairwise comparisons were not significant.

Analysis 3D: Does Task Category Predict Confidence?

Outcome: `confidence_level`
 Model type: linear regression
 Predictor: `assigned_category`

This model asks whether participants felt more confident in their final response for some task categories than others. The confidence scale is treated as approximately continuous.

```
continuous_category_summary(
  task_category_confidence_data,
  "confidence_level"
)
```

```
##   assigned_category  n    mean      sd median
## 1      academic 36 5.305556 1.3270686      6
## 2      creative 35 5.400000 1.3974767      6
## 3    decisions 36 6.138889 0.8669413      6
## 4         info 41 5.170732 1.4474536      5
## 5      learning 42 4.785714 1.4739314      5
## 6       writing 37 5.864865 1.0317774      6
```

```
model_task_confidence <- lm(
  confidence_level ~ assigned_category,
  data = task_category_confidence_data
)
```

```
summary(model_task_confidence)
```

```
##
## Call:
## lm(formula = confidence_level ~ assigned_category, data = task_category_confidence_data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -4.4000 -0.8649  0.1351  0.8611  2.2143
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      5.30556    0.21435   24.752 < 2e-16 ***
## assigned_categorycreative  0.09444    0.30529    0.309  0.75734
## assigned_categorydecisions  0.83333    0.30314    2.749  0.00647 **
## assigned_categoryinfo     -0.13482    0.29375   -0.459  0.64670
## assigned_categorylearning -0.51984    0.29211   -1.780  0.07651 .
## assigned_categorywriting   0.55931    0.30108    1.858  0.06455 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.286 on 221 degrees of freedom
```

```
## Multiple R-squared:  0.1115, Adjusted R-squared:  0.09136
## F-statistic: 5.545 on 5 and 221 DF,  p-value: 7.83e-05
```

```
tests_task_confidence <- anova(model_task_confidence)
```

```
tests_task_confidence
```

```
## Analysis of Variance Table
```

```
##
```

```
## Response: confidence_level
```

```
##              Df Sum Sq Mean Sq F value    Pr(>F)
## assigned_category  5  45.86   9.1712   5.5447 7.83e-05 ***
## Residuals        221 365.55   1.6541
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
emm_task_confidence <- emmeans(
  model_task_confidence,
  ~ assigned_category
)
```

```
emm_task_confidence
```

```
## assigned_category emmean    SE df lower.CL upper.CL
## academic          5.31 0.214 221     4.88     5.73
## creative           5.40 0.217 221     4.97     5.83
## decisions          6.14 0.214 221     5.72     6.56
## info              5.17 0.201 221     4.77     5.57
## learning           4.79 0.198 221     4.39     5.18
## writing            5.86 0.211 221     5.45     6.28
##
```

```
## Confidence level used: 0.95
```

```
pairs_task_confidence <- pairs(
  emm_task_confidence,
  adjust = "tukey"
)
```

```
pairs_task_confidence
```

```
## contrast          estimate    SE df t.ratio p.value
## academic - creative  -0.0944 0.305 221  -0.309  0.9996
## academic - decisions -0.8333 0.303 221  -2.749  0.0699
## academic - info      0.1348 0.294 221   0.459  0.9974
## academic - learning   0.5198 0.292 221   1.780  0.4811
## academic - writing    -0.5593 0.301 221  -1.858  0.4312
## creative - decisions  -0.7389 0.305 221  -2.420  0.1538
## creative - info       0.2293 0.296 221   0.775  0.9715
## creative - learning   0.6143 0.294 221   2.087  0.2979
## creative - writing    -0.4649 0.303 221  -1.533  0.6434
## decisions - info      0.9682 0.294 221   3.296  0.0144
## decisions - learning  1.3532 0.292 221   4.632 <0.0001
```

```
## decisions - writing      0.2740 0.301 221    0.910  0.9436
## info - learning        0.3850 0.282 221    1.364  0.7487
## info - writing         -0.6941 0.292 221   -2.380  0.1677
## learning - writing     -1.0792 0.290 221   -3.722  0.0034
##
## P value adjustment: tukey method for comparing a family of 6 estimates
```

Assigned task category significantly predicted confidence, $F(5, 221) = 5.54$, $p < .001$. Confidence was highest for decisions ($M = 6.14$) and writing ($M = 5.86$), and lowest for learning ($M = 4.79$). Tukey-adjusted comparisons showed significantly higher confidence for decisions than info ($p = .014$) and learning ($p < .001$), and higher confidence for writing than learning ($p = .003$).

Altogether: Task category mattered most for uptake and confidence. **Among AI users, creative and writing tasks produced much higher AI similarity than decisions and learning tasks, $F(5, 152) = 6.08$, $p < .001$.** Task category did not significantly predict whether participants opened AI or whether they reported primarily own-thinking ownership, but it did significantly predict confidence, $F(5, 221) = 5.54$, $p < .001$.

Read Intent Codes

```
intent_path <- "survey_responses_with_ai_intent_codes.csv"

intent_raw <- read.csv(
  intent_path,
  stringsAsFactors = FALSE,
  na.strings = c("", "NA"),
  fileEncoding = "UTF-8"
)

intent <- intent_raw
```

Analysis 4: Asking/Doing/Expressing Intent Coding

This analysis uses the coded `question_for_ai` text, following the ChatGPT paper's Appendix A.2 Asking/Doing/Expressing scheme.

- **Asking:** seeking information or advice to be better informed or make a decision
- **Doing:** requesting that the AI perform a task or produce an output
- **Expressing:** neither asking for information/advice nor requesting a task

This analysis only includes participants who used the AI assistant and had a non-missing `ai_intent_code`. Participants who did not type a prompt are not assigned an intent code because there is no message to classify. Intent codes were assigned by a language model; human validation is a planned next step.

```
intent$ai_used <- tolower(intent$ai_used) == "true"
intent$jaccard_similarity_to_ai <- suppressWarnings(
  as.numeric(intent$jaccard_similarity_to_ai)
)
intent$epistemic_ownership <- suppressWarnings(
  as.numeric(intent$epistemic_ownership)
)
```

```

intent$task_id <- factor(intent$task_id)

intent$ai_intent_code <- factor(
  intent$ai_intent_code,
  levels = c("Asking", "Doing", "Expressing")
)

emm_options(lmer.df = "satterthwaite")

# Binary ownership and reading depth recodes (mirrors model-setup chunk).
intent$epistemic_ownership_binary <- ifelse(
  intent$epistemic_ownership <= 3, 0L,
  ifelse(intent$epistemic_ownership >= 5, 1L, NA_integer_)
)
intent$read_carefully <- as.integer(
  intent$ai_reading_depth == "read_carefully"
)

intent_model_data <- subset(
  intent,
  ai_used == TRUE &
  !is.na(ai_intent_code) &
  !is.na(task_id)
)

intent_similarity_data <- subset(
  intent_model_data,
  !is.na(jaccard_similarity_to_ai)
)

table(intent_model_data$ai_intent_code, useNA = "ifany")

```

```

##
##      Asking      Doing Expressing
##         40         86          6

```

```

table(intent_model_data$task_id, intent_model_data$ai_intent_code)

```

```

##
##           Asking Doing Expressing
## academic_1      1      3         2
## academic_2      0      7         0
## academic_3      1      5         0
## creative_1      0      4         0
## creative_2      0      8         0
## creative_3      0      5         0
## decisions_1     2      4         1
## decisions_2     3      4         1
## decisions_3     3      1         1
## info_1          4      4         0
## info_2          3      5         0
## info_3          4      5         0

```

```
## learning_1      7      5      0
## learning_2      6      6      0
## learning_3      5      4      0
## writing_1        0      5      0
## writing_2        0      6      1
## writing_3        1      5      0
```

Analysis 4A: Does AI Intent Predict Similarity To The AI Response?

Outcome: jaccard_similarity_to_ai
 Model type: linear mixed regression
 Predictor: ai_intent_code
 Random intercept: task_id
 Sample: AI users with a coded question_for_ai

This model asks whether people who used the AI for Asking, Doing, or Expressing differed in how similar their final response was to the pre-written AI response.

```
model_intent_similarity <- lmer(
  jaccard_similarity_to_ai ~ ai_intent_code + (1 | task_id),
  data = intent_similarity_data
)
```

```
## boundary (singular) fit: see help('isSingular')
```

```
summary(model_intent_similarity)
```

```
## Linear mixed model fit by REML. t-tests use Satterthwaite's method [
## lmerModLmerTest]
## Formula: jaccard_similarity_to_ai ~ ai_intent_code + (1 | task_id)
## Data: intent_similarity_data
##
## REML criterion at convergence: 98
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -2.05264 -0.41205  0.03366  0.84441  2.50158
##
## Random effects:
## Groups Name Variance Std.Dev.
## task_id (Intercept) 0.0000  0.0000
## Residual          0.1159  0.3404
## Number of obs: 132, groups: task_id, 18
##
## Fixed effects:
##              Estimate Std. Error    df t value Pr(>|t|)
## (Intercept)      0.14850   0.05382 129.00000   2.759  0.00664 **
## ai_intent_codeDoing      0.56408   0.06514 129.00000   8.659 1.67e-14 ***
## ai_intent_codeExpressing  0.24035   0.14902 129.00000   1.613  0.10921
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Correlation of Fixed Effects:
```

```
##          (Intr) a_nt_D
## a_ntnt_cdDn -0.826
## a_ntnt_cdEx -0.361  0.298
## optimizer (nloptwrap) convergence code: 0 (OK)
## boundary (singular) fit: see help('isSingular')

isSingular(model_intent_similarity)

## [1] TRUE

tests_intent_similarity <- anova(model_intent_similarity)

tests_intent_similarity

## Type III Analysis of Variance Table with Satterthwaite's method
##              Sum Sq Mean Sq NumDF DenDF F value    Pr(>F)
## ai_intent_code 8.8067  4.4034      2   129  38.005 1.056e-13 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

emm_intent_similarity <- emmeans(
  model_intent_similarity,
  ~ ai_intent_code
)

emm_intent_similarity

## ai_intent_code emmean      SE df lower.CL upper.CL
## Asking         0.148 0.0538 129    0.042    0.255
## Doing          0.713 0.0367 129    0.640    0.785
## Expressing     0.389 0.1390 129    0.114    0.664
##
## Degrees-of-freedom method: satterthwaite
## Confidence level used: 0.95

pairs_intent_similarity <- pairs(
  emm_intent_similarity,
  adjust = "tukey"
)

pairs_intent_similarity

## contrast          estimate      SE df t.ratio p.value
## Asking - Doing      -0.564 0.0651 129  -8.659 <0.0001
## Asking - Expressing -0.240 0.1490 129  -1.613  0.2439
## Doing - Expressing   0.324 0.1440 129   2.252  0.0664
##
## Degrees-of-freedom method: satterthwaite
## P value adjustment: tukey method for comparing a family of 3 estimates
```

AI intent significantly predicted Jaccard similarity to the AI response, $F(2, 129) = 38.01$, $p < .001$. Doing prompts produced much higher similarity ($M = .713$) than Asking prompts ($M = .148$); Tukey-adjusted Asking - Doing = -0.564, $p < .001$. Expressing was descriptively intermediate ($M = .389$), but did not differ significantly from Asking ($p = .244$) or Doing ($p = .066$).

Analysis 4B: Does AI Intent Predict Epistemic Ownership?

Outcome: `epistemic_ownership_binary` (0 = AI's thinking, 1 = own thinking)

Model type: mixed logistic regression

Predictor: `ai_intent_code`

Random intercept: `task_id`

Sample: AI users with a coded `question_for_ai`

This model asks whether people who used the AI for Asking, Doing, or Expressing differed in how much AI influence they reported in their final answer. Epistemic ownership is recoded as binary (see Model 3).

```
# Exclude Expressing (n <= 3) -- too sparse for stable binary logistic regression.
intent_ownership_data <- subset(
  intent_model_data,
  !is.na(epistemic_ownership_binary) &
  ai_intent_code %in% c("Asking", "Doing")
)
intent_ownership_data$ai_intent_code <- droplevels(intent_ownership_data$ai_intent_code)

model_intent_ownership <- glmer(
  epistemic_ownership_binary ~ ai_intent_code + (1 | task_id),
  data = intent_ownership_data,
  family = binomial,
  control = glmerControl(optimizer = "bobyqa")
)
```

```
## boundary (singular) fit: see help('isSingular')
```

```
summary(model_intent_ownership)
```

```
## Generalized linear mixed model fit by maximum likelihood (Laplace
## Approximation) [glmerMod]
## Family: binomial ( logit )
## Formula: epistemic_ownership_binary ~ ai_intent_code + (1 | task_id)
## Data: intent_ownership_data
## Control: glmerControl(optimizer = "bobyqa")
##
##           AIC          BIC      logLik -2*log(L)  df.resid
##       132.9       141.1      -63.5     126.9      111
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -0.7386 -0.5195 -0.5195  0.8856  1.9251
##
## Random effects:
##  Groups Name            Variance Std.Dev.
## task_id (Intercept) 0          0
## Number of obs: 114, groups: task_id, 18
##
## Fixed effects:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    -0.6061    0.3589  -1.689   0.0912 .
## ai_intent_codeDoing -0.7038    0.4511  -1.560   0.1187
```



```

## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Correlation of Fixed Effects:
##          (Intr)
## a_ntnt_cdDn -0.796
## optimizer (bobyqa) convergence code: 0 (OK)
## boundary (singular) fit: see help('isSingular')

isSingular(model_intent_ownership)

## [1] TRUE

tests_intent_ownership <- drop1(model_intent_ownership, test = "Chisq")

## boundary (singular) fit: see help('isSingular')

tests_intent_ownership

## Single term deletions
##
## Model:
## epistemic_ownership_binary ~ ai_intent_code + (1 | task_id)
##               npar      AIC      LRT Pr(Chi)
## <none>              132.91
## ai_intent_code      1 133.30 2.3904  0.1221

emm_intent_ownership <- emmeans(
  model_intent_ownership,
  ~ ai_intent_code,
  type = "response"
)

emm_intent_ownership

## ai_intent_code prob  SE df asymp.LCL asymp.UCL
## Asking          0.353 0.0820 Inf      0.213  0.524
## Doing           0.212 0.0457 Inf      0.136  0.316
##
## Confidence level used: 0.95
## Intervals are back-transformed from the logit scale

pairs_intent_ownership <- pairs(
  emm_intent_ownership,
  adjust = "tukey"
)

summary(pairs_intent_ownership, type = "response")

## contrast      odds.ratio    SE df null z.ratio p.value
## Asking / Doing      2.02 0.912 Inf   1  1.560 0.1187
##
## Tests are performed on the log odds ratio scale

```

The Expressing group ($n \leq 3$) is excluded because it is too sparse for stable binary logistic regression. The Asking vs. Doing comparison showed no significant difference in epistemic ownership.

Altogether: Intent coding showed a clear effect for similarity – Doing prompts led to more AI-like final responses than Asking prompts (Jaccard difference = 0.564, $p < .001$). Evidence for differences in binary epistemic ownership was not significant, $\chi^2(1) = 2.39$, $p = .122$.

Analysis 5: AI Reading Depth

This analysis asks whether how carefully participants read the AI response predicted textual similarity and epistemic ownership.

Because 57% of AI users reported reading the AI response carefully (`read_carefully`), the ordered 4-level scale collapses to a binary: `read_carefully` (1) vs. anything less (0). This avoids treating ordinal levels with very few observations as evenly spaced.

Outcome: `jaccard_similarity_to_ai`, `epistemic_ownership_binary`

Model type: linear/logistic mixed regression

Predictor: `read_carefully` (binary)

Random intercept: `task_id`

Sample: AI users with a valid `ai_reading_depth`

```
reading_depth_data <- subset(
  model_data,
  ai_used == TRUE & !is.na(ai_reading_depth) & !is.na(task_id)
)
```

```
# read_carefully inherited from model_data
table(reading_depth_data$ai_reading_depth, useNA = "ifany")
```

```
##
##      not_read      glanced      skimmed read_carefully
##           9         14         45          90
```

```
table(reading_depth_data$read_carefully, useNA = "ifany")
```

```
##
##  0  1
## 68 90
```

```
model_reading_similarity <- lmer(
  jaccard_similarity_to_ai ~ read_carefully + (1 | task_id),
  data = subset(reading_depth_data, !is.na(jaccard_similarity_to_ai))
)
```

```
summary(model_reading_similarity)
```

```
## Linear mixed model fit by REML. t-tests use Satterthwaite's method [
## lmerModLmerTest]
## Formula: jaccard_similarity_to_ai ~ read_carefully + (1 | task_id)
## Data: subset(reading_depth_data, !is.na(jaccard_similarity_to_ai))
##
```

```
## REML criterion at convergence: 170.6
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -1.8733 -0.8683 -0.1436  0.9368  1.6588
##
## Random effects:
##   Groups   Name      Variance Std.Dev.
## task_id   (Intercept) 0.02479  0.1574
## Residual                0.15011  0.3874
## Number of obs: 158, groups: task_id, 18
##
## Fixed effects:
##              Estimate Std. Error      df t value Pr(>|t|)
## (Intercept)    0.62804    0.06096  36.87491  10.303  2.1e-12 ***
## read_carefully -0.19060    0.06476 152.29432  -2.943  0.00376 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Correlation of Fixed Effects:
##              (Intr)
## read_crflly -0.606
```

```
fixed_effects_table(model_reading_similarity)
```

```
##              term      estimate std_error      p_value
## 1      (Intercept) 0.6280359 0.06095777 2.103428e-12
## 2 read_carefully -0.1906014 0.06475940 3.757009e-03
```

```
model_reading_ownership <- glmer(
  epistemic_ownership_binary ~ read_carefully + (1 | task_id),
  data = subset(reading_depth_data, !is.na(epistemic_ownership_binary)),
  family = binomial,
  control = glmerControl(optimizer = "bobyqa")
)
```

```
## boundary (singular) fit: see help('isSingular')
```

```
summary(model_reading_ownership)
```

```
## Generalized linear mixed model fit by maximum likelihood (Laplace
##   Approximation) [glmerMod]
##   Family: binomial ( logit )
## Formula: epistemic_ownership_binary ~ read_carefully + (1 | task_id)
##   Data: subset(reading_depth_data, !is.na(epistemic_ownership_binary))
## Control: glmerControl(optimizer = "bobyqa")
##
##              AIC              BIC      logLik -2*log(L)  df.resid
##          187.6           196.5         -90.8       181.6       141
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
```

```
## -0.7796 -0.7796 -0.6146 1.2826 1.6270
##
## Random effects:
## Groups Name Variance Std.Dev.
## task_id (Intercept) 0 0
## Number of obs: 144, groups: task_id, 18
##
## Fixed effects:
## Estimate Std. Error z value Pr(>|z|)
## (Intercept) -0.9734 0.2847 -3.419 0.000628 ***
## read_carefully 0.4756 0.3646 1.305 0.192037
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Correlation of Fixed Effects:
## (Intr)
## read_crflly -0.781
## optimizer (bobyqa) convergence code: 0 (OK)
## boundary (singular) fit: see help('isSingular')
```

```
fixed_effects_table(model_reading_ownership, exponentiate = TRUE)
```

```
## term estimate std_error p_value odds_ratio
## 1 (Intercept) -0.9734491 0.2846854 0.0006276261 0.3777778
## 2 read_carefully 0.4756107 0.3645705 0.1920365532 1.6089965
```


Reading carefully significantly predicted lower Jaccard similarity, $b = -0.191$, $p = .0038$, but did not significantly predict binary epistemic ownership, $OR = 1.61$, $p = .192$. Among AI users, 57% reported reading the response carefully, 28% skimmed it, and the remaining 15% glanced or did not read it. The negative similarity result suggests that careful readers may have edited or paraphrased more rather than copying verbatim.

```
# Distribution of Jaccard similarity: read_carefully vs. everyone else
density_data <- subset(
  reading_depth_data,
  !is.na(jaccard_similarity_to_ai)
)
density_data$reading_group <- factor(
  ifelse(density_data$read_carefully == 1, "Read carefully", "Skimmed / glanced / didn't read"),
  levels = c("Read carefully", "Skimmed / glanced / didn't read")
)
group_means <- aggregate(
  jaccard_similarity_to_ai ~ reading_group,
  data = density_data,
  FUN = mean, na.rm = TRUE
)
group_ns <- table(density_data$reading_group)
group_means$n <- as.integer(group_ns[as.character(group_means$reading_group)])
group_means$label <- sprintf(
  "%s: mean = %.2f (n = %d)",
  group_means$reading_group, group_means$jaccard_similarity_to_ai, group_means$n
)
```

```

ggplot(density_data, aes(x = jaccard_similarity_to_ai, fill = reading_group)) +
  geom_density(alpha = 0.45, color = NA, bw = 0.08) +
  geom_vline(
    data = group_means,
    aes(xintercept = jaccard_similarity_to_ai, color = reading_group),
    linetype = "dashed", linewidth = 0.9, show.legend = FALSE
  ) +
  scale_fill_manual(values = c("Read carefully" = "#4472C4",
                                "Skimmed / glanced / didn't read" = "#C00000")) +
  scale_color_manual(values = c("Read carefully" = "#4472C4",
                                "Skimmed / glanced / didn't read" = "#C00000")) +
  scale_x_continuous(limits = c(0, 1), expand = c(0, 0)) +
  labs(
    title = "Distribution of Similarity to AI by Reading Depth",
    subtitle = paste0(
      group_means$label[group_means$reading_group == "Read carefully"], " | ",
      group_means$label[group_means$reading_group == "Skimmed / glanced / didn't read"]
    ),
    x = "Jaccard Similarity to AI Response",
    y = "Density",
    fill = NULL
  ) +
  theme_classic(base_size = 12) +
  theme(
    plot.title = element_text(face = "bold", size = 12),
    plot.subtitle = element_text(size = 9.5),
    legend.position = "bottom"
  )

```



Create Local Embedding Cosine Similarity Dataset

This section creates a new dataset in R called `embedding_similarity_data`.

The goal is to measure semantic similarity between each participant's final response and the AI response they were shown. This is different from `jaccard_similarity_to_ai`, which mostly captures exact word overlap. Embedding cosine similarity is better for shared meaning, including synonyms or paraphrased ideas.

Here, we use a free local embedding model `sentence-transformers/all-MiniLM-L6-v2`. Embedding size: 384 dimensions, not 3072 like the OpenAI model.

```
library(reticulate)

use_virtualenv("r-reticulate-py312", required = TRUE)

py_config()

torch <- import("torch")
sentence_transformers <- import("sentence_transformers")
```

Compute Embedding Cosine Similarity

```
library(reticulate)
```

```
## Warning: package 'reticulate' was built under R version 4.4.3
```

```
use_virtualenv("r-reticulate-py312", required = TRUE)
py_config()
```

```
## python:      C:/Users/lengo/Documents/.virtualenvs/r-reticulate-py312/Scripts/python.exe
## libpython:   C:/Users/lengo/AppData/Local/r-reticulate/r-reticulate/pyenv/pyenv-win/versions/3.12
## pythonhome:  C:/Users/lengo/Documents/.virtualenvs/r-reticulate-py312
## version:     3.12.8 (tags/v3.12.8:2dc476b, Dec  3 2024, 19:30:04) [MSC v.1942 64 bit (AMD64)]
## Architecture: 64bit
## numpy:       C:/Users/lengo/Documents/.virtualenvs/r-reticulate-py312/Lib/site-packages/numpy
## numpy_version: 2.4.6
##
## NOTE: Python version was forced by use_python() function
```

```
sentence_transformers <- tryCatch(
  import("sentence_transformers", delay_load = FALSE),
  error = function(e) {
    python_config <- paste(capture.output(py_config()), collapse = "\n")
    stop(
      "R could not import the Python module 'sentence_transformers'.\n\n",
      "Python configuration being used by reticulate:\n",
      python_config,
      "\n\nOriginal Python import error:\n",
      conditionMessage(e),
      "\n\nTry restarting R, running the one-time setup chunk, ",
      "and then running this chunk again.",
    )
  }
)
```

```

    call. = FALSE
  )
}
)

embedding_model_name <- "sentence-transformers/all-MiniLM-L6-v2"
embedding_model <- sentence_transformers$SentenceTransformer(
  embedding_model_name
)

embedding_similarity_data <- intent

embedding_similarity_data$embedding_cosine_similarity_to_ai <- NA_real_
embedding_similarity_data$embedding_model_name <- embedding_model_name

embedding_rows <- which(
  embedding_similarity_data$ai_used == TRUE &
  !is.na(embedding_similarity_data$participant_response) &
  !is.na(embedding_similarity_data$ai_response_shown) &
  trimws(embedding_similarity_data$participant_response) != "" &
  trimws(embedding_similarity_data$ai_response_shown) != ""
)

participant_texts <- as.character(
  embedding_similarity_data$participant_response[embedding_rows]
)

ai_texts <- as.character(
  embedding_similarity_data$ai_response_shown[embedding_rows]
)

texts_to_embed <- c(participant_texts, ai_texts)

# normalize_embeddings = TRUE makes cosine similarity equal to the dot product.
embeddings <- embedding_model$encode(
  texts_to_embed,
  convert_to_numpy = TRUE,
  normalize_embeddings = TRUE,
  show_progress_bar = TRUE
)

embeddings <- py_to_r(embeddings)

n_pairs <- length(embedding_rows)

participant_embeddings <- embeddings[seq_len(n_pairs), , drop = FALSE]
ai_embeddings <- embeddings[n_pairs + seq_len(n_pairs), , drop = FALSE]

embedding_cosine_similarity <- rowSums(
  participant_embeddings * ai_embeddings
)

embedding_similarity_data$embedding_cosine_similarity_to_ai[

```

```

embedding_rows
] <- embedding_cosine_similarity

summary(embedding_similarity_data$embedding_cosine_similarity_to_ai)

```

```

##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.     NA's
## 0.05209 0.58189 0.86238 0.77955 1.00000 1.00000      70

```

Embedding Similarity Robustness Checks

Because Jaccard similarity and embedding cosine similarity were strongly correlated, the embedding measure is best used as a robustness check and semantic extension. These models ask whether the main findings still hold when similarity means shared meaning rather than exact word overlap.

```

embedding_analysis_data <- subset(
  embedding_similarity_data,
  ai_used == TRUE &
    !is.na(embedding_cosine_similarity_to_ai) &
    !is.na(task_id)
)

embedding_analysis_data$task_id <- factor(embedding_analysis_data$task_id)

embedding_analysis_data$ai_intent_code <- factor(
  embedding_analysis_data$ai_intent_code,
  levels = c("Asking", "Doing", "Expressing")
)

embedding_analysis_data$ai_kind_choice <- factor(
  embedding_analysis_data$ai_kind_choice
)

embedding_analysis_data$assigned_category <- factor(
  embedding_analysis_data$assigned_category
)

embedding_analysis_data$z_user_decides <- z(
  embedding_analysis_data$exp_ai_does_work_vs_user_decides
)

embedding_analysis_data$z_independent <- z(
  -embedding_analysis_data$exp_independent_vs_reliant
)

embedding_analysis_data$z_hours_per_week <- z(
  embedding_analysis_data$hours_per_week
)

embedding_analysis_data$z_instrument_vs_point_of_view <- z(
  embedding_analysis_data$exp_instrument_vs_point_of_view
)

embedding_analysis_data$z_existing_info_vs_creating_new <- z(

```



```

embedding_analysis_data$exp_existing_info_vs_creating_new
)

embedding_analysis_data$z_embedding_cosine_similarity_to_ai <- z(
  embedding_analysis_data$embedding_cosine_similarity_to_ai
)
table(embedding_analysis_data$ai_intent_code, useNA = "ifany")

```

```

##
##      Asking      Doing Expressing      <NA>
##          40          86           6         25

```

```

summary(embedding_analysis_data$embedding_cosine_similarity_to_ai)

```

```

##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## 0.05209 0.58189 0.86238 0.77955 1.00000 1.00000

```

Embedding Robustness 1: Jaccard vs. Embedding Similarity

```

similarity_check_data <- subset(
  embedding_analysis_data,
  !is.na(jaccard_similarity_to_ai) &
  !is.na(embedding_cosine_similarity_to_ai)
)

embedding_jaccard_correlation <- cor(
  similarity_check_data$jaccard_similarity_to_ai,
  similarity_check_data$embedding_cosine_similarity_to_ai,
  use = "complete.obs"
)

embedding_jaccard_correlation

```

```


## [1] 0.8283619

```

```

plot(
  similarity_check_data$jaccard_similarity_to_ai,
  similarity_check_data$embedding_cosine_similarity_to_ai,
  xlab = "Jaccard similarity to AI",
  ylab = "Embedding cosine similarity to AI",
  main = "Lexical vs. semantic similarity"
)

```



Jaccard similarity to the AI response was strongly positively associated with embedding cosine similarity to the AI response, Pearson $r = .83$. This suggests that the lexical overlap measure captured meaningful semantic convergence with the AI response, not only exact-word copying.

Embedding Robustness 2: Does AI Intent Predict Semantic Similarity?

Outcome: `embedding_cosine_similarity_to_ai`

Model type: linear mixed regression

Predictor: `ai_intent_code`

Random intercept: `task_id`

Sample: AI users with embedding similarity scores

This is the semantic version of Analysis 4A. It asks whether Doing prompts led to responses that were closer in meaning to the AI response, not just closer in wording.

```
embedding_intent_data <- subset(
  embedding_analysis_data,
  !is.na(ai_intent_code)
)

model_embedding_intent <- lmer(
  embedding_cosine_similarity_to_ai ~ ai_intent_code + (1 | task_id),
  data = embedding_intent_data
)

summary(model_embedding_intent)
```

```
## Linear mixed model fit by REML. t-tests use Satterthwaite's method [
## lmerModLmerTest]
## Formula: embedding_cosine_similarity_to_ai ~ ai_intent_code + (1 | task_id)
## Data: embedding_intent_data
##
## REML criterion at convergence: -26.7
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -3.1151 -0.3742  0.4133  0.5825  1.8866
##
## Random effects:
## Groups Name Variance Std.Dev.
## task_id (Intercept) 0.002806 0.05297
## Residual 0.041897 0.20469
## Number of obs: 132, groups: task_id, 18
##
## Fixed effects:
##              Estimate Std. Error      df t value Pr(>|t|)
## (Intercept)    0.59546    0.03618  51.53334  16.456 < 2e-16 ***
## ai_intent_codeDoing    0.28883    0.04085  126.46581   7.071 9.26e-11 ***
## ai_intent_codeExpressing 0.02322    0.09232  128.99347   0.251  0.802
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Correlation of Fixed Effects:
##              (Intr) a_nt_D
## a_ntnt_cdDn -0.786
## a_ntnt_cdEx -0.352  0.306
```

```
isSingular(model_embedding_intent)
```

```
## [1] FALSE
```

```
tests_embedding_intent <- anova(model_embedding_intent)
```

```
tests_embedding_intent
```

```
## Type III Analysis of Variance Table with Satterthwaite's method
##              Sum Sq Mean Sq NumDF DenDF F value Pr(>F)
## ai_intent_code 2.2633  1.1317     2 128.05  27.011 1.633e-10 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
emm_embedding_intent <- emmeans(
  model_embedding_intent,
  ~ ai_intent_code
)
```

```
emm_embedding_intent
```

```
## ai_intent_code emmean      SE      df lower.CL upper.CL
```

```
## Asking          0.595 0.0362 51.5    0.523    0.668
## Doing           0.884 0.0256 27.8    0.832    0.937
## Expressing      0.619 0.0865 127.2    0.448    0.790
##
## Degrees-of-freedom method: satterthwaite
## Confidence level used: 0.95
```

```
pairs_embedding_intent <- pairs(
  emm_embedding_intent,
  adjust = "tukey"
)

pairs_embedding_intent
```

```
## contrast          estimate      SE df t.ratio p.value
## Asking - Doing      -0.2888 0.0408 126  -7.071 <0.0001
## Asking - Expressing -0.0232 0.0923 129  -0.251  0.9657
## Doing - Expressing   0.2656 0.0888 129   2.991  0.0093
##
## Degrees-of-freedom method: satterthwaite
## P value adjustment: tukey method for comparing a family of 3 estimates
```

AI intent significantly predicted embedding-based semantic similarity to the AI response, $F(2, 128.05) = 27.01$, $p < .001$. Doing prompts were more semantically similar to the AI response ($M = .884$) than Asking prompts ($M = .595$), Tukey-adjusted Asking - Doing = -0.289 , $p < .001$. Doing was also higher than Expressing ($M = .619$), Doing - Expressing = 0.266 , $p = .009$, while Asking and Expressing did not differ ($p = .966$).

Embedding Robustness 3: Do Pre-Task Independence Beliefs Predict Semantic Similarity?

Outcome: `embedding_cosine_similarity_to_ai`
 Model type: linear mixed regression
 Predictors: `z_user_decides`, `z_independent`
 Random intercept: `task_id`
 Sample: AI users with embedding similarity scores

This is the semantic version of the earlier stated-independence similarity model. It asks whether participants who described themselves as more independent on either slider produced responses that were less semantically similar to the AI response.

```
model_embedding_independence <- lmer(
  embedding_cosine_similarity_to_ai ~ z_user_decides +
    z_independent +
    z_hours_per_week +
    z_instrument_vs_point_of_view +
    (1 | task_id),
  data = embedding_analysis_data
)

summary(model_embedding_independence)
```

```
## Linear mixed model fit by REML. t-tests use Satterthwaite's method [
```

```
## lmerModLmerTest]
## Formula: embedding_cosine_similarity_to_ai ~ z_user_decides + z_independent +
##       z_hours_per_week + z_instrument_vs_point_of_view + (1 | task_id)
## Data: embedding_analysis_data
##
## REML criterion at convergence: 31.9
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -3.0158 -0.7210  0.2456  0.7842  1.6869
##
## Random effects:
## Groups Name Variance Std.Dev.
## task_id (Intercept) 0.008454 0.09194
## Residual 0.055897 0.23643
## Number of obs: 157, groups: task_id, 18
##
## Fixed effects:
##
##              Estimate Std. Error      df t value Pr(>|t|)
## (Intercept)    0.778696   0.028912  14.653318  26.933 6.85e-14
## z_user_decides -0.028370   0.020553  148.722718  -1.380   0.170
## z_independent  -0.007937   0.020655  149.186153  -0.384   0.701
## z_hours_per_week  0.005288   0.019841  145.197136   0.267   0.790
## z_instrument_vs_point_of_view  0.008407   0.020611  148.265054   0.408   0.684
##
## (Intercept)          ***
## z_user_decides
## z_independent
## z_hours_per_week
## z_instrument_vs_point_of_view
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Correlation of Fixed Effects:
##      (Intr) z_sr_d z_ndpn z_hr__
## z_user_dcds -0.005
## z_indepndnt  0.014 -0.228
## z_hrs_pr_wk -0.015 -0.080 -0.078
## z_nstrm____ -0.006 -0.186  0.216 -0.127
```

```
isSingular(model_embedding_independence)
```

```
## [1] FALSE
```

```
fixed_effects_table(model_embedding_independence)
```

```
##              term      estimate std_error    p_value
## 1      (Intercept)  0.778695623 0.02891202 6.849850e-14
## 2      z_user_decides -0.028369873 0.02055262 1.695492e-01
## 3      z_independent -0.007937422 0.02065540 7.013194e-01
## 4      z_hours_per_week  0.005287936 0.01984099 7.902202e-01
## 5 z_instrument_vs_point_of_view  0.008407357 0.02061130 6.839352e-01
```

Neither independence predictor significantly predicted embedding-based semantic similarity to the AI response: $z_{\text{user_decides}}$, $b = -0.028$, $p = .170$; $z_{\text{independent}}$, $b = -0.008$, $p = .701$. Hours per week ($b = 0.005$, $p = .790$) and instrument-vs-point-of-view beliefs ($b = 0.008$, $p = .684$) were also not significant. Participants who described themselves as more independent did not reliably produce responses that were less semantically similar to the AI response.

Embedding Robustness 4: Does AI Conceptualization Predict Semantic Similarity?

Outcome: `embedding_cosine_similarity_to_ai`

Model type: linear mixed regression

Predictor: `ai_kind_choice`

Random intercept: `task_id`

Sample: AI users with embedding similarity scores, excluding `Writing` or `print` because that group had only 3 participants overall

This is the semantic version of Analysis 1B. It asks whether people's conceptualization of AI predicted meaning-level convergence with the AI response.

```
embedding_concept_data <- subset(
  embedding_analysis_data,
  !is.na(ai_kind_choice) & ai_kind_choice != "Writing or print"
)

embedding_concept_data$ai_kind_choice <- droplevels(
  embedding_concept_data$ai_kind_choice
)

contrasts(embedding_concept_data$ai_kind_choice) <- contr.sum(
  nlevels(embedding_concept_data$ai_kind_choice)
)

model_embedding_ai_kind_similarity <- lmer(
  embedding_cosine_similarity_to_ai ~ ai_kind_choice + (1 | task_id),
  data = embedding_concept_data
)

summary(model_embedding_ai_kind_similarity)
```

```
## Linear mixed model fit by REML. t-tests use Satterthwaite's method [
## lmerModLmerTest]
## Formula: embedding_cosine_similarity_to_ai ~ ai_kind_choice + (1 | task_id)
## Data: embedding_concept_data
##
## REML criterion at convergence: 27.3
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -2.9971 -0.7779  0.2943  0.7361  1.4732
##
## Random effects:
## Groups   Name                Variance Std.Dev.
## task_id  (Intercept)  0.005923  0.07696
## Residual                                0.058927  0.24275
```

```
## Number of obs: 152, groups: task_id, 18
##
## Fixed effects:
##              Estimate Std. Error      df t value Pr(>|t|)
## (Intercept)   0.780751   0.032608  30.048925  23.944  <2e-16 ***
## ai_kind_choice1  0.037248   0.047923 143.480556   0.777    0.438
## ai_kind_choice2 -0.001736   0.036323 146.089782  -0.048    0.962
## ai_kind_choice3 -0.017423   0.034843 144.089550  -0.500    0.618
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Correlation of Fixed Effects:
##              (Intr) a_kn_1 a_kn_2
## ai_knd_chc1  0.031
## ai_knd_chc2 -0.401 -0.175
## ai_knd_chc3 -0.450 -0.163  0.179
```

```
anova(model_embedding_ai_kind_similarity)
```

```
## Type III Analysis of Variance Table with Satterthwaite's method
##              Sum Sq Mean Sq NumDF DenDF F value Pr(>F)
## ai_kind_choice 0.045363 0.015121      3 144.24  0.2566 0.8565
```

```
emm_embedding_ai_kind_similarity <- emmeans(
  model_embedding_ai_kind_similarity,
  ~ ai_kind_choice
)

emm_embedding_ai_kind_similarity
```

```
## ai_kind_choice      emmean      SE      df lower.CL upper.CL
## A calculator or specialized tool  0.818 0.0588 114.2    0.702    0.934
## A colleague or advisor           0.779 0.0379  46.5    0.703    0.855
## A library or search engine        0.763 0.0354  37.7    0.692    0.835
## A market or aggregator           0.763 0.0840 147.1    0.597    0.929
##
## Degrees-of-freedom method: satterthwaite
## Confidence level used: 0.95
```

```
pairs_embedding_ai_kind_similarity <- pairs(
  emm_embedding_ai_kind_similarity,
  adjust = "tukey"
)

pairs_embedding_ai_kind_similarity
```

```
## contrast      estimate      SE
## A calculator or specialized tool - A colleague or advisor  0.038984 0.0650
## A calculator or specialized tool - A library or search engine 0.054670 0.0637
## A calculator or specialized tool - A market or aggregator  0.055336 0.0993
## A colleague or advisor - A library or search engine 0.015686 0.0456
## A colleague or advisor - A market or aggregator 0.016352 0.0890
```

```
## A library or search engine - A market or aggregator          0.000666 0.0876
## df t.ratio p.value
## 145 0.600 0.9320
## 145 0.859 0.8261
## 141 0.557 0.9444
## 147 0.344 0.9860
## 142 0.184 0.9978
## 140 0.008 1.0000
##
## Degrees-of-freedom method: satterthwaite
## P value adjustment: tukey method for comparing a family of 4 estimates
```

AI conceptualization did not significantly predict semantic similarity to the AI response, $F(3, 144.24) = 0.26$, $p = .857$. Tukey-adjusted pairwise comparisons showed no significant differences between conceptualization groups, all $p \geq .826$, consistent with the Jaccard-based result in Analysis 1B.

Embedding Robustness 5: Do Learning-Curve Variables Predict Semantic Similarity?

Outcome: embedding_cosine_similarity_to_ai
 Model type: linear mixed regression
 Predictors: ai_use_tenure, z_hours_per_week, ai_understanding
 Random intercept: task_id
 Sample: AI users with embedding similarity scores

This is the semantic version of Analysis 2B. It asks whether more experienced participants produced responses that were closer in meaning to the AI response.

```
embedding_learning_data <- subset(
  embedding_analysis_data,
  !is.na(ai_use_tenure) &
  !is.na(z_hours_per_week) &
  !is.na(ai_understanding)
)

model_embedding_learning_similarity <- lmer(
  embedding_cosine_similarity_to_ai ~ ai_use_tenure +
    z_hours_per_week +
    ai_understanding +
    (1 | task_id),
  data = embedding_learning_data
)

summary(model_embedding_learning_similarity)
```

```
## Linear mixed model fit by REML. t-tests use Satterthwaite's method [
## lmerModLmerTest]
## Formula:
## embedding_cosine_similarity_to_ai ~ ai_use_tenure + z_hours_per_week +
## ai_understanding + (1 | task_id)
## Data: embedding_learning_data
##
## REML criterion at convergence: 36
##
```



```

## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -2.9325 -0.6558  0.2366  0.7327  1.4598
##
## Random effects:
##   Groups   Name      Variance Std.Dev.
##   task_id  (Intercept) 0.007623 0.08731
##   Residual              0.057196 0.23916
## Number of obs: 157, groups: task_id, 18
##
## Fixed effects:
##
## (Intercept)
## ai_use_tenure6 to 12 months
## ai_use_tenureI do not use them, or am just starting out.
## ai_use_tenureLess than 6 months
## ai_use_tenureMore than 2 years
## z_hours_per_week
## ai_understandingI just use it. I do not think much about how it works.
## ai_understandingI understand it reasonably well. I know about training, prompting, and limitations.
## ai_understandingI understand it technically. I work with or study AI systems.
##
## (Intercept)
## ai_use_tenure6 to 12 months
## ai_use_tenureI do not use them, or am just starting out.
## ai_use_tenureLess than 6 months
## ai_use_tenureMore than 2 years
## z_hours_per_week
## ai_understandingI just use it. I do not think much about how it works.
## ai_understandingI understand it reasonably well. I know about training, prompting, and limitations.
## ai_understandingI understand it technically. I work with or study AI systems.
##
## (Intercept)
## ai_use_tenure6 to 12 months
## ai_use_tenureI do not use them, or am just starting out.
## ai_use_tenureLess than 6 months
## ai_use_tenureMore than 2 years
## z_hours_per_week
## ai_understandingI just use it. I do not think much about how it works.
## ai_understandingI understand it reasonably well. I know about training, prompting, and limitations.
## ai_understandingI understand it technically. I work with or study AI systems.
##
## (Intercept)
## ai_use_tenure6 to 12 months
## ai_use_tenureI do not use them, or am just starting out.
## ai_use_tenureLess than 6 months
## ai_use_tenureMore than 2 years
## z_hours_per_week
## ai_understandingI just use it. I do not think much about how it works.
## ai_understandingI understand it reasonably well. I know about training, prompting, and limitations.
## ai_understandingI understand it technically. I work with or study AI systems.
##
## (Intercept)
## ai_use_tenure6 to 12 months

```

```
## ai_use_tenureI do not use them, or am just starting out.
## ai_use_tenureLess than 6 months
## ai_use_tenureMore than 2 years
## z_hours_per_week
## ai_understandingI just use it. I do not think much about how it works.
## ai_understandingI understand it reasonably well. I know about training, prompting, and limitations.
## ai_understandingI understand it technically. I work with or study AI systems.
##
## (Intercept)
## ai_use_tenure6 to 12 months
## ai_use_tenureI do not use them, or am just starting out.
## ai_use_tenureLess than 6 months
## ai_use_tenureMore than 2 years
## z_hours_per_week
## ai_understandingI just use it. I do not think much about how it works.
## ai_understandingI understand it reasonably well. I know about training, prompting, and limitations.
## ai_understandingI understand it technically. I work with or study AI systems.
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Correlation of Fixed Effects:
##          (Intr) a__t1m adnutoajso a__t6m a__t2y z_hr__ ajuiIdntmahiw
## a_s_tn6t12m   -0.470
## a_dnutoajso   -0.045  0.092
## a_s_tnrLt6m   -0.317  0.264  0.061
## a_s_tnrMt2y   -0.447  0.382  0.085      0.251
## z_hrs_pr_wk    0.096  0.069  0.054      0.026 -0.186
## ajuiIdntmahiw -0.303  0.084 -0.004      0.007  0.007 -0.174
## auirwIkatpal  -0.588  0.114 -0.090      0.063  0.021 -0.078  0.307
## auitIwosAs    -0.267 -0.016 -0.025      0.068 -0.043 -0.112  0.169
##
##          auirwIkatpal
## a_s_tn6t12m
## a_dnutoajso
## a_s_tnrLt6m
## a_s_tnrMt2y
## z_hrs_pr_wk
## ajuiIdntmahiw
## auirwIkatpal
## auitIwosAs    0.299
```

```
anova(model_embedding_learning_similarity)
```

```
## Type III Analysis of Variance Table with Satterthwaite's method
##              Sum Sq Mean Sq NumDF DenDF F value Pr(>F)
## ai_use_tenure  0.172450 0.043113      4 141.87  0.7538 0.5571
## z_hours_per_week 0.004093 0.004093      1 141.38  0.0716 0.7895
## ai_understanding 0.059564 0.019855      3 140.40  0.3471 0.7913
```

```
emm_embedding_learning_tenure <- emmeans(
  model_embedding_learning_similarity,
  ~ ai_use_tenure
)
```

```
emm_embedding_learning_tenure
```

```
## ai_use_tenure          emmean    SE    df lower.CL
## 1 to 2 years           0.759 0.0441  65.8    0.671
## 6 to 12 months        0.825 0.0517  93.2    0.722
## I do not use them, or am just starting out. 0.802 0.2510 144.9    0.307
## Less than 6 months    0.787 0.0759 137.3    0.637
## More than 2 years     0.835 0.0466  75.8    0.742
## upper.CL
##      0.847
##      0.927
##      1.297
##      0.938
##      0.928
##
## Results are averaged over the levels of: ai_understanding
## Degrees-of-freedom method: satterthwaite
## Confidence level used: 0.95
```

AI-use tenure, $F(4, 141.87) = 0.75$, $p = .557$, hours per week, $F(1, 141.38) = 0.07$, $p = .790$, and AI understanding, $F(3, 140.40) = 0.35$, $p = .791$, did not significantly predict embedding-based semantic similarity to the AI response. This pattern is consistent with the Jaccard-based result in Analysis 2B.

Embedding Robustness 6: Does Task Category Predict Semantic Similarity?

Outcome: `embedding_cosine_similarity_to_ai`

Model type: linear regression

Predictor: `assigned_category`

Sample: AI users with embedding similarity scores

This is the semantic version of Analysis 3B. It uses fixed effects for task category, matching the task-category analysis above.

```
embedding_task_category_data <- subset(
  embedding_analysis_data,
  assigned_category %in% task_categories &
  !is.na(embedding_cosine_similarity_to_ai)
)

embedding_task_category_data$assigned_category <- factor(
  embedding_task_category_data$assigned_category,
  levels = task_categories
)

continuous_category_summary(
  embedding_task_category_data,
  "embedding_cosine_similarity_to_ai"
)
```

```
## assigned_category n      mean      sd    median
## 1      academic 21 0.8014962 0.2390487 0.9981539
## 2      creative 23 0.8479434 0.2936878 0.9999999
## 3     decisions 24 0.6022662 0.2684879 0.5655401
```

```
## 4          info 29 0.7316817 0.2624763 0.7841010
## 5      learning 36 0.8105107 0.1952642 0.8629885
## 6      writing 24 0.8835004 0.1494671 0.9892157
```

```
model_embedding_task_category_similarity <- lm(
  embedding_cosine_similarity_to_ai ~ assigned_category,
  data = embedding_task_category_data
)
```

```
summary(model_embedding_task_category_similarity)
```

```
##
## Call:
## lm(formula = embedding_cosine_similarity_to_ai ~ assigned_category,
##     data = embedding_task_category_data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.78024 -0.14966  0.05841  0.15206  0.39773
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      0.801496   0.051700  15.503 < 2e-16 ***
## assigned_categorycreative  0.046447   0.071508   0.650  0.51698
## assigned_categorydecisions -0.199230   0.070793  -2.814  0.00554 **
## assigned_categoryinfo     -0.069814   0.067885  -1.028  0.30540
## assigned_categorylearning  0.009015   0.065054   0.139  0.88998
## assigned_categorywriting   0.082004   0.070793   1.158  0.24854
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.2369 on 151 degrees of freedom
## Multiple R-squared:  0.1269, Adjusted R-squared:  0.09803
## F-statistic: 4.391 on 5 and 151 DF, p-value: 0.0009236
```

```
anova(model_embedding_task_category_similarity)
```

```
## Analysis of Variance Table
##
## Response: embedding_cosine_similarity_to_ai
##              Df Sum Sq Mean Sq F value    Pr(>F)
## assigned_category    5 1.2323  0.246461   4.3908 0.0009236 ***
## Residuals          151 8.4758  0.056131
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
emm_embedding_task_category_similarity <- emmeans(
  model_embedding_task_category_similarity,
  ~ assigned_category
)
```

```
emm_embedding_task_category_similarity
```

```
## assigned_category emmean      SE df lower.CL upper.CL
## academic          0.801 0.0517 151    0.699    0.904
## creative           0.848 0.0494 151    0.750    0.946
## decisions          0.602 0.0484 151    0.507    0.698
## info              0.732 0.0440 151    0.645    0.819
## learning           0.811 0.0395 151    0.732    0.889
## writing            0.884 0.0484 151    0.788    0.979
##
## Confidence level used: 0.95
```

```
pairs_embedding_task_category_similarity <- pairs(
  emm_embedding_task_category_similarity,
  adjust = "tukey"
)

pairs_embedding_task_category_similarity
```

```
## contrast          estimate      SE df t.ratio p.value
## academic - creative -0.04645 0.0715 151  -0.650  0.9869
## academic - decisions  0.19923 0.0708 151   2.814  0.0606
## academic - info       0.06981 0.0679 151   1.028  0.9077
## academic - learning -0.00901 0.0651 151  -0.139  1.0000
## academic - writing    -0.08200 0.0708 151  -1.158  0.8557
## creative - decisions  0.24568 0.0691 151   3.554  0.0066
## creative - info       0.11626 0.0662 151   1.758  0.4963
## creative - learning   0.03743 0.0632 151   0.592  0.9914
## creative - writing    -0.03556 0.0691 151  -0.514  0.9956
## decisions - info      -0.12942 0.0654 151  -1.979  0.3590
## decisions - learning -0.20824 0.0624 151  -3.335  0.0134
## decisions - writing    -0.28123 0.0684 151  -4.112  0.0009
## info - learning       -0.07883 0.0591 151  -1.333  0.7660
## info - writing         -0.15182 0.0654 151  -2.322  0.1917
## learning - writing     -0.07299 0.0624 151  -1.169  0.8508
##
## P value adjustment: tukey method for comparing a family of 6 estimates
```

Task category significantly predicted embedding-based semantic similarity, $F(5, 151) = 4.39$, $p < .001$. Writing tasks had the highest estimated semantic similarity ($M = .884$) and decisions tasks had the lowest ($M = .602$). Tukey-adjusted pairwise comparisons showed that decisions were significantly lower than creative tasks (difference = -0.246, $p = .0066$), learning tasks (difference = -0.208, $p = .013$), and writing tasks (difference = -0.281, $p < .001$).

Embedding Robustness 7: Does Reading Depth Predict Semantic Similarity?

Outcome: `embedding_cosine_similarity_to_ai`

Model type: linear mixed regression

Predictor: `read_carefully` (binary)

Random intercept: `task_id`

Sample: AI users with embedding similarity scores and a valid `ai_reading_depth`

This is the semantic version of the reading-depth similarity model in Analysis 5.

```

embedding_reading_depth_data <- subset(
  embedding_analysis_data,
  !is.na(ai_reading_depth)
)

embedding_reading_depth_data$read_carefully <- as.integer(
  embedding_reading_depth_data$ai_reading_depth == "read_carefully"
)

model_embedding_reading_similarity <- lmer(
  embedding_cosine_similarity_to_ai ~ read_carefully + (1 | task_id),
  data = embedding_reading_depth_data
)

summary(model_embedding_reading_similarity)

## Linear mixed model fit by REML. t-tests use Satterthwaite's method [
## lmerModLmerTest]
## Formula: embedding_cosine_similarity_to_ai ~ read_carefully + (1 | task_id)
## Data: embedding_reading_depth_data
##
## REML criterion at convergence: 10
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -3.3647 -0.6968  0.3560  0.7494  1.5124
##
## Random effects:
## Groups Name Variance Std.Dev.
## task_id (Intercept) 0.008758 0.09359
## Residual 0.053663 0.23165
## Number of obs: 157, groups: task_id, 18
##
## Fixed effects:
##              Estimate Std. Error      df t value Pr(>|t|)
## (Intercept)  0.82938    0.03637  34.88666  22.807  <2e-16 ***
## read_carefully -0.08902    0.03881 151.06051  -2.294  0.0232 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Correlation of Fixed Effects:
##              (Intr)
## read_crflly -0.605

fixed_effects_table(model_embedding_reading_similarity)

```

```

##              term      estimate std_error      p_value
## 1 (Intercept)  0.82938155 0.03636554 1.512924e-22
## 2 read_carefully -0.08902257 0.03880571 2.316537e-02

```

Reading carefully significantly predicted lower embedding-based semantic similarity, $b = -0.089$, $p = .023$. This is consistent with Analysis 5: careful readers showed lower lexical and semantic similarity, suggesting more selective editing or paraphrasing rather than direct copying.

Cases Where Lexical And Semantic Similarity Diverge

This table helps identify cases where participants did not share many exact words with the AI response but still had high semantic similarity.

```
embedding_analysis_data$jaccard_z <- z(
  embedding_analysis_data$jaccard_similarity_to_ai
)

embedding_analysis_data$embedding_similarity_z <- z(
  embedding_analysis_data$embedding_cosine_similarity_to_ai
)

embedding_analysis_data$semantic_minus_lexical <-
  embedding_analysis_data$embedding_similarity_z -
  embedding_analysis_data$jaccard_z

embedding_divergence_cases <- embedding_analysis_data[
  order(-embedding_analysis_data$semantic_minus_lexical),
  c(
    "submission_id",
    "assigned_category",
    "ai_intent_code",
    "jaccard_similarity_to_ai",
    "embedding_cosine_similarity_to_ai",
    "semantic_minus_lexical",
    "participant_response",
    "ai_response_shown"
  )
]

embedding_divergence_cases
```

##	submission_id	assigned_category	ai_intent_code
## 200	ai-d07f7fe1-634e-4b7a-8715-38695fe51985	learning	Asking
## 48	ai-fb25bda4-b9d2-4548-a930-4367a7620940	learning	Doing
## 183	ai-0194fcbf-b43b-48ee-909a-0d5b523e377e	learning	Asking
## 168	ai-00113ca9-1216-4a5f-a4c7-9107e42f0948	learning	<NA>
## 187	ai-4684f7ed-a24f-4e26-8d48-b636360f7b63	info	Asking
## 13	ai-802fec76-fc09-429b-b4f5-74f6643a19ec	learning	Doing
## 203	ai-8bbfa3d6-4e68-4150-afd0-91febbcd78f0	learning	Asking
## 68	ai-e3007e63-495b-4e2a-9c90-7b613d0bba23	learning	Asking
## 55	ai-f32f8340-4cd5-44ff-a028-1cc1a787d7b7	learning	Asking
## 75	ai-7de8404f-ff3d-4cac-ab7f-121868326d22	learning	Doing
## 95	ai-658b6665-1e89-43ac-8293-8912d94f181c	learning	Asking
## 185	ai-375b84d5-de93-44e5-85f6-ee8b0fbace0c	info	Doing
## 34	ai-25a61e97-2444-4095-b07a-7f763e4b440b	learning	Doing
## 197	ai-d7cc626b-1951-4ec8-b50c-3412cb5f981a	learning	Doing
## 11	ai-c8b66e6b-7ac7-4889-995b-1fa326f42f26	learning	Asking
## 105	ai-c82af804-e29b-4c2e-bc99-678ad2a59c94	info	Asking
## 35	ai-cfc1376a-921a-4e11-a818-35a4038be455	academic	Doing
## 212	ai-8676c9ec-ae2e-4fe3-8384-513dc06313b6	learning	Asking
## 14	ai-480b82e9-6a2b-4b09-bd0b-58e20601ca0c	creative	<NA>
## 88	ai-41df940e-deea-43c9-8a35-3f8a65228a5e	learning	Asking

## 63	ai-57f22ce2-b49d-48db-8c44-0e619737f951	learning	<NA>
## 9	ai-f41b68bb-6e9e-46e5-8f45-9aa7bae58ab9	academic	Doing
## 166	ai-043a2017-3c73-41de-85eb-2824ad332b38	info	Doing
## 132	ai-7fc47df5-97e6-4e5f-b31a-44aea4ce424c	learning	Doing
## 146	ai-a4d0a785-bea7-4b5b-8f90-9cc3a5216012	learning	Asking
## 4	ai-2c713b0c-b1a6-46a7-92ca-d737c06830bf	decisions	Doing
## 80	ai-3de92f98-3c76-4079-9e0c-74bdb130a8fc	learning	<NA>
## 19	ai-a19188e1-a1de-4a59-acb3-52e0153f56eb	info	Asking
## 40	ai-7f22fa45-3ff8-4e9f-9e48-ed938199c411	creative	<NA>
## 70	ai-74efcb75-a131-45b9-9ddd-2be2a3e26c70	info	Asking
## 69	ai-506e63d2-9d4d-4049-a9c0-621d5d0306b4	writing	<NA>
## 62	ai-fc01ab3a-1580-4992-83c5-c1161cd2f3c1	learning	Doing
## 72	ai-fa6af708-08a0-43bc-b150-b0c473526855	writing	<NA>
## 142	ai-036b9943-41ef-4153-837e-97f59cef9f2f	writing	Doing
## 90	ai-d3fffa42-0196-4a02-ac79-145cf281b5d5	info	Doing
## 79	ai-fad2aa5b-4252-4a19-808e-597378908dd7	decisions	<NA>
## 184	ai-a2410e08-e9d0-4261-9265-ccb057246e6b	academic	Doing
## 123	ai-98a9f8b0-a38c-437c-a1d3-0ef8b4556755	learning	Asking
## 33	ai-c47c4c3e-71e5-4427-97ec-d799ea4ade11	academic	Doing
## 83	ai-1780384949407-83z90pym18g	decisions	Asking
## 53	ai-e72053f5-bc7a-49be-b10f-26816f6e0ead	academic	Doing
## 121	ai-84d1e4a7-b9f2-4141-ade7-0dc4eb2dcd63	writing	Doing
## 165	ai-f214ca17-23e9-43b6-86bd-7d30fc3c8f3b	info	Doing
## 213	ai-79050dc5-cdf7-4c59-b6cf-eb67420ab638	decisions	Asking
## 51	ai-2dea9244-b173-4d30-a2ee-43fb7f38e4fc	learning	Asking
## 57	ai-1780382223533-1e3jjyghs5d	learning	Doing
## 52	ai-db819a6a-0bfa-48fd-9041-dc47bf3fbc7	info	Doing
## 170	ai-7a02838d-03c5-467f-9c28-31a9ec31bd58	learning	Asking
## 32	ai-53d14de3-84d0-4a54-9094-537b083a4ae2	info	Doing
## 135	ai-4fc82e3c-a8f9-40df-9f64-12289cf675bd	learning	Asking
## 41	ai-1780382055100-pq94eggrdj9	decisions	<NA>
## 179	ai-80b80b6b-e15a-4c1b-8d01-f64def255968	learning	Asking
## 61	ai-7327a28a-cfdd-4295-be40-d82bde104f0a	decisions	Asking
## 25	ai-380c80db-8894-40a5-b665-de02e34a8dca	writing	<NA>
## 156	ai-2bc96f0e-2c97-4d8e-bd85-358ecf9951d7	decisions	Expressing
## 8	ai-63ff3c65-a7d6-480a-b00e-5dac58c41dc8	info	<NA>
## 18	ai-d2bf71f2-9eff-46a0-a8a8-53e578fb760d	writing	<NA>
## 161	ai-5ceb0401-c7f6-454e-9738-eb5dbabf72da	learning	Asking
## 155	ai-3639d520-8bb0-4c0b-a552-b68b3d5a4bfd	decisions	Doing
## 130	ai-3d2ab6d2-5f92-4eca-89f1-81176f1dd244	info	Asking
## 49	ai-2434a9cb-1da0-49cd-aabe-c3180a6aaafd	academic	Asking
## 133	ai-644b27be-a736-455d-888b-fa4973460df7	writing	Asking
## 6	ai-4719caa3-7a55-4c33-bd88-c3e639489708	info	Asking
## 225	ai-6e3d867c-f4df-4470-8205-bd2d154b069f	decisions	Asking
## 150	ai-5788b73b-bc5c-4bc9-a01b-eee88aa034f3	writing	Doing
## 58	ai-15cf7924-5bd5-42c1-8e33-c1321d94f46c	decisions	Asking
## 125	ai-541e2136-5a21-4ae8-b7de-d35fbc31c86d	writing	Doing
## 16	ai-05b7871a-ef13-48b4-a0de-388551d4749d	info	<NA>
## 169	ai-f6f48ce2-5f6e-4491-847e-9f74a7c591d2	decisions	Doing
## 137	ai-04f36787-468d-4c09-9452-4d63906e467b	info	Asking
## 117	ai-894cf2d9-1671-4771-9dc1-77f10ef0478f	writing	Doing
## 93	ai-1780386396588-tqlpk93nsc	writing	<NA>
## 159	ai-31c8400c-a2d3-4fdd-a6af-c29726676721	writing	Doing
## 60	ai-a04ac4f3-8a89-4928-9f55-dc334773e45f	creative	Doing

## 100	ai-5a46312c-0103-4c62-b370-2d2b718063d7	writing	Doing
## 17	ai-50e57cb2-6cfb-4bc0-92a6-48169a4453cd	writing	Doing
## 78	ai-64101525-589f-4b2a-b2d0-716a0958b87f	academic	<NA>
## 84	ai-b3e0b5a5-3721-47e5-86af-d92e787d8eb2	creative	Doing
## 59	ai-b559fa87-7123-43e3-9a85-c63cbb8471ec	writing	Doing
## 77	ai-1780383992866-v3cn2v08wgi	academic	Doing
## 54	ai-1c27e5ae-866e-4349-9db4-7ec42752d43d	academic	Expressing
## 199	ai-df1fb0c8-3356-4f21-a0ac-643fcec9f0ac	writing	Expressing
## 45	ai-c0a9d318-d72f-43e5-8bf2-71a9b2fc7647	info	<NA>
## 139	ai-1668fdcd-a675-488e-bf06-77b879fee727	info	Asking
## 50	ai-f0740d95-0b49-49f3-92a6-31168e14e10a	writing	Doing
## 24	ai-cce5b570-8aee-4ffb-bbc6-7a4d9337ccd7	info	Doing
## 145	ai-e2feda33-be10-408b-9202-59190975df89	learning	Doing
## 147	ai-667e3853-b8b2-4cf6-ac50-fea11083b967	learning	Doing
## 23	ai-85186445-5b06-41b0-a628-22c6e0927db4	academic	Doing
## 108	ai-4a158d13-7a0b-4d89-b9c7-22122eca67e0	learning	Doing
## 226	ai-2e2f1669-1c7a-40ef-a382-6a184af445f7	learning	Doing
## 92	ai-5865e427-b54f-46eb-acd1-0a606f5b4e31	info	Asking
## 107	ai-f487f5b5-1e43-4c50-a1a4-43b83313aca4	info	Doing
## 224	ai-7ecfef37-065b-42a7-8136-f8f5501add5d	info	Doing
## 46	ai-4676ece5-b290-422b-916c-18427b38f067	decisions	Doing
## 149	ai-9898e418-d2bf-4042-830f-f4e4b1e747f2	decisions	Doing
## 37	ai-76be743e-4c44-4477-826d-b3f1784b95f6	academic	Doing
## 91	ai-25a62fbe-faa4-4d2c-a117-8db9b579db77	academic	Doing
## 208	ai-ef69d085-7c3b-41e2-991f-9c117a6f5883	academic	<NA>
## 36	ai-ef7f768d-4b70-457b-a2f4-02537384463a	info	Doing
## 148	ai-57fc84ce-fea3-42a3-b26e-27b262f862fd	info	Doing
## 201	ai-702a40b4-475a-4a70-9f40-ff834a8aa076	info	Doing
## 1	ai-05e682ab-248a-4f40-9b3a-b74f583a72c8	learning	Doing
## 5	ai-b412fe12-44c2-4976-9a4e-aed83f18f7cf	learning	Doing
## 220	ai-9987ba30-607e-4a41-836d-9c56b0141daa	learning	Doing
## 98	ai-cbeda50f-64f7-4185-962b-9d155e1be1be	academic	Doing
## 167	ai-94345d83-3eb9-4ed4-9745-351c45165fa8	academic	Expressing
## 172	ai-3bc4dc10-e3b0-4274-bf09-22cf680b8a41	academic	Doing
## 2	ai-ab784340-49a7-4104-92ca-d547be96dd4f	creative	Doing
## 30	ai-f7e06690-c0e7-43ff-b639-9ec468aa5799	creative	Doing
## 103	ai-faccebf0-68d5-44cf-a96c-bc07c48dee7b	creative	Doing
## 113	ai-257878db-84a4-4695-979c-de2e1b7f6dbe	creative	Doing
## 128	ai-2deb3418-5c3c-45f0-986a-e7948e00c5dc	creative	Doing
## 163	ai-ddaca5e9-35e3-4267-aa27-44326fa6b702	creative	Doing
## 109	ai-391cad87-9a13-4343-b940-e5d560df149c	creative	Doing
## 190	ai-4b327991-4786-435c-9ed2-fdc881b355e5	creative	Doing
## 195	ai-311527bb-108a-438d-b60e-d05ad53db747	creative	Doing
## 222	ai-1e2c1956-950f-4567-98be-717561906681	creative	Doing
## 74	ai-4790c4bd-f7a1-4555-a54e-37ff1653afa1	writing	Doing
## 86	ai-bb078926-7e23-4be3-942c-f53bb4c3e699	writing	Doing
## 140	ai-d72a192b-c0e6-4916-aa2b-f1d33cb95d4f	writing	<NA>
## 198	ai-acd9a310-9901-44f0-866f-68d00e582abf	writing	Doing
## 206	ai-45d87509-95df-4ba3-8684-bf4c47114357	writing	Doing
## 110	ai-75a1f5ad-5a10-4338-b298-1c27400165f2	decisions	Doing
## 28	ai-c40c3831-5a60-4414-9aaa-ce98e0436784	info	<NA>
## 116	ai-9c92e553-a702-47dd-850f-7c30b013b9cf	info	Doing
## 27	ai-7fb8b321-79aa-4e6a-907d-11427d683586	writing	Doing
## 144	ai-38ee6a8e-b11b-4d64-8b0c-ffae3e9d810b	decisions	Doing

## 66	ai-dc688d91-f689-4178-aee9-dfd03ae8d0b3	academic	Doing
## 87	ai-cb4d9ae9-eebf-415f-aaff-5488c1680abc	academic	Doing
## 209	ai-55542a1e-5319-410a-bfe1-21e9adbfdab4	academic	Doing
## 81	ai-e412922d-abea-41e5-891f-8db0a799fc1d	creative	<NA>
## 82	ai-1780384840980-qjcy7t5141k	creative	Doing
## 174	ai-513a1235-0fa4-4d1c-9b33-98640ebb5206	creative	Doing
## 211	ai-31ec013c-4843-41d9-8de2-eae2bcbdl9ea	creative	<NA>
## 215	ai-d5022e49-dc2d-4f15-b18b-c348087fc1e7	creative	<NA>
## 106	ai-df9e2071-c4b3-4686-ba0a-6ec41505e2f3	academic	Doing
## 44	ai-3951d5d6-f3d8-4baf-a22b-e511fe50633d	decisions	Doing
## 21	ai-ccb3f336-7ef2-4a73-8c1c-67cad746792f	writing	Doing
## 71	ai-99d0e015-dce5-4c20-93b0-f6b8fa6201af	info	Doing
## 221	ai-950ef311-6489-4b30-93f7-f4ebaf0e9bfe	decisions	<NA>
## 219	ai-f1d710a4-c334-4aa5-a196-f629b8740b8a	decisions	Asking
## 194	ai-580297a3-a980-491c-98fd-3ffce776fefa	creative	Doing
## 120	ai-17ff6899-7836-41d6-8eae-34a95814dffc	decisions	Expressing
## 56	ai-162f6f01-48ee-4777-a486-ae23dcc85f88	decisions	Doing
## 173	ai-88a80b3e-f515-4568-942a-5e5b8877b0d7	learning	Asking
## 15	ai-d654fb95-6b07-4461-9e67-b9559ee8cd55	decisions	Asking
## 96	ai-b768b11f-81a2-4bc3-9226-527d4db25c37	decisions	<NA>
## 216	ai-e55ed7e9-0128-4486-bbbe-76bb1c519a36	creative	Doing
## 43	ai-e06bfee9-6ad4-4045-8b89-1be0929e2d6a	academic	Asking
## 162	ai-28cf02f6-7f90-47ad-9afe-685556e6baf5	creative	Doing
## 136	ai-47a74470-e9db-45fc-864a-62c8cac63c36	info	Asking
## 127	ai-26bcd679-30ca-4067-814d-0da75d2f6434	decisions	Expressing
## 112	ai-74a5dc89-0f9d-437c-844d-7710520ab12e	learning	Asking
## 65	ai-ca973dc9-9ed3-4e89-ab3e-d29a66f54760	creative	<NA>
## 204	ai-82458211-0cb5-4c13-a29a-cf67131c4304	info	Asking
## 126	ai-002afe42-6bae-47f0-a433-654c09bf2d9b	decisions	Asking
##	jaccard_similarity_to_ai embedding_cosine_similarity_to_ai		
## 200	0.21951219	0.92747154	
## 48	0.15217391	0.87027097	
## 183	0.16049383	0.86238122	
## 168	0.17272727	0.86892432	
## 187	0.06185567	0.80321242	
## 13	0.15384615	0.85613012	
## 203	0.08333333	0.80262313	
## 68	0.23750000	0.88992872	
## 55	0.08256881	0.79644042	
## 75	0.32911392	0.94123846	
## 95	0.15094340	0.83572366	
## 185	0.22916667	0.88032279	
## 34	0.30864198	0.92320698	
## 197	0.22093023	0.86359574	
## 11	0.31168831	0.91618313	
## 105	0.09278351	0.78410096	
## 35	0.12605042	0.78592362	
## 212	0.15942029	0.80028588	
## 14	0.23214286	0.82759778	
## 88	0.31914894	0.87511480	
## 63	0.22522523	0.81141141	
## 9	0.10679612	0.73965127	
## 166	0.07216495	0.70820976	
## 132	0.25000000	0.81205038	

## 146	0.11811024	0.72505046
## 4	0.30000000	0.81514832
## 80	0.26804124	0.79375868
## 19	0.18823529	0.74557059
## 40	0.09677419	0.67650601
## 70	0.17391304	0.71505866
## 69	0.20634921	0.73383604
## 62	0.42857143	0.86404002
## 72	0.18867924	0.71530161
## 142	0.06250000	0.62896121
## 90	0.58823529	0.92822698
## 79	0.20560748	0.69615750
## 184	0.07291667	0.61436780
## 123	0.07692308	0.61142447
## 33	0.06896552	0.60565974
## 83	0.02941176	0.57556314
## 53	0.04878049	0.58604164
## 121	0.47058824	0.83333394
## 165	0.53012048	0.86401546
## 213	0.06172839	0.57719992
## 51	0.08080808	0.58188596
## 57	0.44230769	0.79480718
## 52	0.60919540	0.89048543
## 170	0.07142857	0.57155796
## 32	0.14606742	0.60833183
## 135	0.79104478	0.98414635
## 41	0.12222222	0.58762524
## 179	0.04587156	0.53361043
## 61	0.10869565	0.56979391
## 25	0.09230769	0.54662596
## 156	0.08247423	0.53835117
## 8	0.24038461	0.62767349
## 18	0.60975610	0.84562434
## 161	0.15596330	0.57629191
## 155	0.61666667	0.84191849
## 130	0.06060606	0.51074529
## 49	0.07377049	0.51786934
## 133	0.18032787	0.57847486
## 6	0.14851485	0.55871269
## 225	0.08695652	0.52030264
## 150	0.53191489	0.78305764
## 58	0.17777778	0.56128624
## 125	0.65909091	0.83516472
## 16	0.02970297	0.46104758
## 169	0.03614458	0.46458904
## 137	0.08737864	0.49214858
## 117	0.62500000	0.80871253
## 93	0.94736842	0.99663703
## 159	0.84782609	0.93687419
## 60	0.04494382	0.46146694
## 100	0.95454545	0.99651358
## 17	0.95454545	0.99624594
## 78	0.57291667	0.76856086
## 84	0.97468354	0.99598547

## 59	0.97674419	0.99608053
## 77	0.04166667	0.44250570
## 54	0.16161616	0.51126033
## 199	0.97674419	0.99265851
## 45	0.03260870	0.43369648
## 139	0.07777778	0.45829505
## 50	0.97297297	0.98577294
## 24	1.00000000	1.00000014
## 145	1.00000000	1.00000014
## 147	1.00000000	1.00000014
## 23	1.00000000	1.00000013
## 108	1.00000000	1.00000011
## 226	1.00000000	1.00000010
## 92	1.00000000	1.00000009
## 107	1.00000000	1.00000009
## 224	1.00000000	1.00000009
## 46	1.00000000	1.00000009
## 149	1.00000000	1.00000009
## 37	1.00000000	1.00000009
## 91	1.00000000	1.00000009
## 208	1.00000000	1.00000009
## 36	1.00000000	1.00000008
## 148	1.00000000	1.00000008
## 201	1.00000000	1.00000008
## 1	1.00000000	1.00000006
## 5	1.00000000	1.00000006
## 220	1.00000000	1.00000006
## 98	1.00000000	1.00000005
## 167	1.00000000	1.00000005
## 172	1.00000000	1.00000005
## 2	1.00000000	1.00000005
## 30	1.00000000	1.00000005
## 103	1.00000000	1.00000005
## 113	1.00000000	1.00000005
## 128	1.00000000	1.00000005
## 163	1.00000000	1.00000005
## 109	1.00000000	1.00000004
## 190	1.00000000	1.00000004
## 195	1.00000000	1.00000004
## 222	1.00000000	1.00000004
## 74	1.00000000	1.00000004
## 86	1.00000000	1.00000004
## 140	1.00000000	1.00000004
## 198	1.00000000	1.00000004
## 206	1.00000000	1.00000004
## 110	1.00000000	1.00000003
## 28	1.00000000	1.00000003
## 116	1.00000000	1.00000003
## 27	1.00000000	0.99999999
## 144	1.00000000	0.99999999
## 66	1.00000000	0.99999993
## 87	1.00000000	0.99999993
## 209	1.00000000	0.99999993
## 81	1.00000000	0.99999990

## 82	1.00000000	0.99999990
## 174	1.00000000	0.99999990
## 211	1.00000000	0.99999990
## 215	1.00000000	0.99999990
## 106	1.00000000	0.99815390
## 44	1.00000000	0.99692608
## 21	1.00000000	0.99413293
## 71	0.07751938	0.44758251
## 221	0.02380952	0.41516187
## 219	0.06976744	0.43202401
## 194	1.00000000	0.97817635
## 120	0.08695652	0.43216898
## 56	0.20547945	0.49751966
## 173	0.02941176	0.38607977
## 15	0.05319149	0.38319880
## 96	0.09000000	0.34668298
## 216	0.03076923	0.28742891
## 43	0.02000000	0.26142506
## 162	0.01388889	0.20783077
## 136	0.06976744	0.23416447
## 127	0.02531646	0.15068111
## 112	0.00000000	0.10275160
## 65	0.00000000	0.06770594
## 204	0.01098901	0.06716727
## 126	0.01234568	0.05208959
##	semantic_minus_lexical	
## 200	1.279242418	
## 48	1.209591204	
## 183	1.158239353	
## 168	1.155465303	
## 187	1.154903619	
## 13	1.148941226	
## 203	1.101622355	
## 68	1.086101775	
## 55	1.078650669	
## 75	1.074586429	
## 95	1.074021056	
## 185	1.067351718	
## 34	1.050839517	
## 197	1.019825972	
## 11	1.015461261	
## 105	1.004969452	
## 35	0.933407009	
## 212	0.911867104	
## 14	0.848940670	
## 88	0.833145905	
## 63	0.800455696	
## 9	0.793566089	
## 166	0.749631819	
## 132	0.744281362	
## 146	0.708213419	
## 4	0.638160409	
## 80	0.628184710	
## 19	0.624219396	

## 40	0.564199844
## 70	0.535863339
## 69	0.534235639
## 62	0.529333587
## 72	0.501829673
## 142	0.454867101
## 90	0.408106267
## 79	0.384954760
## 184	0.371671791
## 123	0.350374689
## 33	0.346131704
## 83	0.319259100
## 53	0.315344421
## 121	0.306631284
## 165	0.288483825
## 213	0.249204462
## 51	0.222755161
## 57	0.219238855
## 52	0.207122283
## 170	0.203590867
## 32	0.174050835
## 135	0.151448050
## 41	0.147577655
## 179	0.112063434
## 61	0.108167069
## 25	0.054147680
## 156	0.044290139
## 8	0.027978311
## 18	0.025961431
## 161	0.022153541
## 155	-0.005277479
## 130	-0.014526868
## 49	-0.017179257
## 133	-0.026859021
## 6	-0.030656370
## 225	-0.038686326
## 150	-0.040300434
## 58	-0.089716147
## 125	-0.132929741
## 16	-0.140481986
## 169	-0.141557277
## 137	-0.152546336
## 117	-0.158144511
## 93	-0.169091670
## 159	-0.172665088
## 60	-0.174933779
## 100	-0.186601768
## 17	-0.187674632
## 78	-0.195619194
## 84	-0.236461926
## 59	-0.240966194
## 77	-0.243172971
## 54	-0.251935824
## 199	-0.254683819

## 45	-0.257011342
## 139	-0.265491088
## 50	-0.273344791
## 24	-0.280388538
## 145	-0.280388564
## 147	-0.280388564
## 23	-0.280388580
## 108	-0.280388683
## 226	-0.280388692
## 92	-0.280388746
## 107	-0.280388746
## 224	-0.280388746
## 46	-0.280388747
## 149	-0.280388747
## 37	-0.280388766
## 91	-0.280388766
## 208	-0.280388766
## 36	-0.280388773
## 148	-0.280388773
## 201	-0.280388773
## 1	-0.280388858
## 5	-0.280388858
## 220	-0.280388858
## 98	-0.280388898
## 167	-0.280388898
## 172	-0.280388898
## 2	-0.280388900
## 30	-0.280388900
## 103	-0.280388900
## 113	-0.280388900
## 128	-0.280388900
## 163	-0.280388900
## 109	-0.280388951
## 190	-0.280388951
## 195	-0.280388951
## 222	-0.280388951
## 74	-0.280388952
## 86	-0.280388952
## 140	-0.280388952
## 198	-0.280388952
## 206	-0.280388952
## 110	-0.280388998
## 28	-0.280389011
## 116	-0.280389011
## 27	-0.280389145
## 144	-0.280389149
## 66	-0.280389396
## 87	-0.280389396
## 209	-0.280389396
## 81	-0.280389530
## 82	-0.280389530
## 174	-0.280389530
## 211	-0.280389530
## 215	-0.280389530

## 106	-0.287789450
## 44	-0.292711319
## 21	-0.303908034
## 71	-0.307821072
## 221	-0.310448792
## 219	-0.351811190
## 194	-0.367872077
## 120	-0.391981733
## 56	-0.411007701
## 173	-0.440309953
## 15	-0.508235374
## 96	-0.741879066
## 216	-0.838983197
## 43	-0.917691486
## 162	-1.118043111
## 136	-1.144957218
## 127	-1.374227383
## 112	-1.506339195
## 65	-1.646824340
## 204	-1.675036297
## 126	-1.738693526
##	
## 200	
## 48	
## 183	
## 168	
## 187	
## 13	
## 203	
## 68	
## 55	
## 75	
## 95	
## 185	
## 34	
## 197	
## 11	
## 105	
## 35	
## 212	
## 14	
## 88	
## 63	
## 9	
## 166	
## 132	
## 146	
## 4	
## 80	
## 19	
## 40	
## 70	
## 69	
## 62	

Having eight months of s

72
142
90
79
184
123
33
83
53
121
165
213
51
57
52
170
32
135
41
179
61
25
156
8
18
161
155
130
49
133
6
225
150
58
125
16
169
137
117
93
159
60
100
17
78
84
59
77
54
199
45
139
50
24

145 General anesthesia has been used in surgery since
 ## 147 General anesthesia has been used in surgery since
 ## 23 The Norwegian study's within-family design is its most important methodology
 ## 108
 ## 226
 ## 92
 ## 107
 ## 224
 ## 46
 ## 149
 ## 37 The co
 ## 91 The c
 ## 208 The c
 ## 36
 ## 148
 ## 201
 ## 1 Research in epigenetics has shown that certain experiences - particularly stress, starvation, or
 ## 5 Research in epigenetics has shown that certain experiences - particularly stress, starvation, o
 ## 220 Research in epigenetics has shown that certain experiences - particularly stress, starvation, o
 ## 98
 ## 167
 ## 172
 ## 2
 ## 30
 ## 103
 ## 113
 ## 128
 ## 163
 ## 109
 ## 190
 ## 195
 ## 222
 ## 74
 ## 86
 ## 140
 ## 198
 ## 206
 ## 110
 ## 28
 ## 116
 ## 27
 ## 144
 ## 66 The Norwegian study's within-family design is its m
 ## 87 The Norwegian study's within-family design is its m
 ## 209 The Norwegian study's within-family design is its m
 ## 81
 ## 82
 ## 174
 ## 211
 ## 215
 ## 106 The
 ## 44
 ## 21
 ## 71

221
 ## 219
 ## 194
 ## 120
 ## 56
 ## 173
 ## 15
 ## 96
 ## 216
 ## 43
 ## 162
 ## 136
 ## 127
 ## 112
 ## 65
 ## 204
 ## 126
 ##
 ## 200
 ## 48 General anesthesia has been used in surgery since
 ## 183
 ## 168
 ## 187
 ## 13 General anesthesia has been used in surgery since
 ## 203 General anesthesia has been used in surgery since
 ## 68
 ## 55 General anesthesia has been used in surgery since
 ## 75
 ## 95 General anesthesia has been used in surgery since
 ## 185
 ## 34
 ## 197
 ## 11
 ## 105
 ## 35 The co
 ## 212
 ## 14
 ## 88 General anesthesia has been used in surgery since
 ## 63 General anesthesia has been used in surgery since
 ## 9
 ## 166
 ## 132 General anesthesia has been used in surgery since
 ## 146 Research in epigenetics has shown that certain experiences - particularly stress, starvation, or
 ## 4
 ## 80 Research in epigenetics has shown that certain experiences - particularly stress, starvation, or
 ## 19
 ## 40
 ## 70
 ## 69
 ## 62 General anesthesia has been used in surgery since
 ## 72
 ## 142
 ## 90
 ## 79

```

## 184
## 123 General anesthesia has been used in surgery since
## 33
## 83
## 53 The Norwegian study's within-family design is its most important methodologic
## 121
## 165
## 213
## 51 Research in epigenetics has shown that certain experiences - particularly stress, starvation, or
## 57 Research in epigenetics has shown that certain experiences - particularly stress, starvation, or
## 52
## 170 General anesthesia has been used in surgery since
## 32
## 135
## 41
## 179 Research in epigenetics has shown that certain experiences - particularly stress, starvation, or
## 61
## 25
## 156
## 8
## 18
## 161 Research in epigenetics has shown that certain experiences - particularly stress, starvation, or
## 155
## 130
## 49
## 133
## 6
## 225
## 150
## 58
## 125
## 16
## 169
## 137
## 117
## 93
## 159
## 60
## 100
## 17
## 78 The temporal pattern here is gen
## 84
## 59
## 77
## 54 The temporal pattern here is gen
## 199
## 45
## 139
## 50
## 24
## 145 General anesthesia has been used in surgery since
## 147 General anesthesia has been used in surgery since
## 23 The Norwegian study's within-family design is its most important methodologic
## 108

```

```

## 226
## 92
## 107
## 224
## 46
## 149
## 37
## 91
## 208
## 36
## 148
## 201
## 1   Research in epigenetics has shown that certain experiences - particularly stress, starvation, or
## 5   Research in epigenetics has shown that certain experiences - particularly stress, starvation, or
## 220 Research in epigenetics has shown that certain experiences - particularly stress, starvation, or
## 98
## 167
## 172
## 2
## 30
## 103
## 113
## 128
## 163
## 109
## 190
## 195
## 222
## 74
## 86
## 140
## 198
## 206
## 110
## 28
## 116
## 27
## 144
## 66
## 87
## 209
## 81
## 82
## 174
## 211
## 215
## 106
## 44
## 21
## 71
## 221
## 219
## 194
## 120

```

The co

The co

The co

The Norwegian study's within-family design is its mo

The Norwegian study's within-family design is its mo

The Norwegian study's within-family design is its mo

The co

```
## 56
## 173 Research in epigenetics has shown that certain experiences - particularly stress, starvation, or
## 15
## 96
## 216
## 43          The Norwegian study's within-family design is its most important methodologic
## 162
## 136
## 127
## 112
## 65
## 204
## 126
```

```
summary(embedding_divergence_cases)
```


```
## submission_id      assigned_category    ai_intent_code
## Length:157         academic :21        Asking      :40
## Class :character   creative :23        Doing       :86
## Mode :character    decisions:24      Expressing: 6
##                   info      :29        NA's       :25
##                   learning :36
##                   writing   :24
## jaccard_similarity_to_ai embedding_cosine_similarity_to_ai
## Min. :0.00000      Min. :0.05209
## 1st Qu.:0.08738    1st Qu.:0.58189
## Median :0.31169    Median :0.86238
## Mean :0.50899      Mean :0.77955
## 3rd Qu.:1.00000    3rd Qu.:1.00000
## Max. :1.00000      Max. :1.00000
## semantic_minus_lexical participant_response ai_response_shown
## Min. :-1.7387      Length:157      Length:157
## 1st Qu.: -0.2804    Class :character Class :character
## Median : -0.2410    Mode :character Mode :character
## Mean : 0.0000
## 3rd Qu.: 0.3193
## Max. : 1.2792
```

High positive cases mean relatively high semantic similarity but lower lexical overlap so possible paraphrasing / same ideas in different words.

High negative cases mean relatively high lexical overlap but lower semantic similarity so possible shared keywords, copied phrases, or task-specific vocabulary without the whole response being as semantically close.

```
hist(
  embedding_analysis_data$semantic_minus_lexical,
  breaks = 15,
  xlab = "Semantic minus lexical similarity",
  main = "Distribution of semantic-lexical divergence"
)
```


Distribution of semantic–lexical divergence



The divergence distribution was centered slightly below zero, suggesting that lexical and semantic similarity generally aligned, with a small tendency for lexical overlap to be relatively stronger. However, there were more moderately positive divergence cases than strongly negative ones, indicating several responses that were semantically close to the AI response despite lower exact word overlap.

```
boxplot(  
  semantic_minus_lexical ~ ai_intent_code,  
  data = embedding_analysis_data,  
  xlab = "AI intent code",  
  ylab = "Semantic minus lexical similarity",  
  main = "Divergence by AI intent"  
)
```

Divergence by AI intent



Asking has the highest and most variable semantic_minus_lexical values. Many Asking cases are positive or near positive, meaning people who asked for information/advice may have taken up the AI's ideas while using their own wording. Doing is centered lower, around negative values. That suggests Doing prompts tended to produce responses where lexical overlap was relatively strong compared with semantic-only divergence.

```
par(mfrow = c(1, 2))

plot(
  embedding_analysis_data$z_user_decides,
  embedding_analysis_data$semantic_minus_lexical,
  xlab = "z_user_decides",
  ylab = "Semantic minus lexical similarity",
  main = "Divergence by z_user_decides"
)
abline(
  lm(
    semantic_minus_lexical ~ z_user_decides,
    data = embedding_analysis_data
  ),
  col = "red",
  lwd = 2
)

plot(
  embedding_analysis_data$z_independent,
  embedding_analysis_data$semantic_minus_lexical,
```




```


xlab = "z_independent",
ylab = "Semantic minus lexical similarity",
main = "Divergence by z_independent"
)
abline(
  lm(
    semantic_minus_lexical ~ z_independent,
    data = embedding_analysis_data
  ),
  col = "blue",
  lwd = 2
)

```

Divergence by z_user_decides



Divergence by z_independent



```

par(mfrow = c(1, 1))

summary(lm(
  semantic_minus_lexical ~ z_user_decides + z_independent,
  data = embedding_analysis_data
))

```

```

##
## Call:
## lm(formula = semantic_minus_lexical ~ z_user_decides + z_independent,
##     data = embedding_analysis_data)
##

```

```
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.7931 -0.2804 -0.1841  0.3090  1.2711
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   1.828e-16  4.654e-02   0.000  1.0000
## z_user_decides -2.058e-02  4.769e-02  -0.432  0.6666
## z_independent  8.858e-02  4.769e-02   1.857  0.0652 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.5832 on 154 degrees of freedom
## Multiple R-squared:  0.02193,    Adjusted R-squared:  0.009227
## F-statistic: 1.726 on 2 and 154 DF,  p-value: 0.1813
```

Neither independence predictor significantly predicted semantic-lexical divergence. Participants who described themselves as more independent on either slider were not more likely to show idea-level uptake without word-level overlap.

Altogether: Embedding similarity analyses suggest that participants' overlap with the AI response also reflected meaningful semantic convergence besides shared words. Asking prompts appear to involve using AI as an informational or advisory resource, where participants may incorporate AI-generated ideas while still rephrasing or integrating them into their own response. Doing prompts, by contrast, more directly delegate the task of producing the response to the AI. This was reflected in higher lexical and semantic similarity to the AI response among Doing prompts, suggesting a stronger form of cognitive offloading. However, neither pre-task independence predictor reliably predicted lower semantic uptake or greater paraphrasing. Actual prompting behavior predicted AI uptake more clearly than prior independence beliefs.

Exploratory Clustering: AI Use Profiles

This exploratory analysis asks whether participants fall into AI-use profiles based on how often they use AI for different purposes. Because the eight frequency items are on the same 1-7 scale but may vary in spread, I standardize them before clustering. The cluster profiles are then described using the original 1-7 means so they remain interpretable.

```
cluster_frequency_items <- c(
  "freq_information_guidance",
  "freq_writing_communication",
  "freq_coding_math_technical",
  "freq_creative_work",
  "freq_learning_understanding",
  "freq_advice_decisions",
  "freq_summarizing_opinions",
  "freq_automation_repetitive"
)

cluster_data <- embedding_similarity_data

for (item in cluster_frequency_items) {
  cluster_data[[item]] <- suppressWarnings(as.numeric(cluster_data[[item]]))
}
```

```

cluster_complete_data <- cluster_data[
  complete.cases(cluster_data[, cluster_frequency_items]),
]

cluster_matrix <- scale(cluster_complete_data[, cluster_frequency_items])

set.seed(123)

cluster_fit_table <- data.frame()
cluster_models <- list()

for (k in 2:4) {
  fit <- kmeans(
    cluster_matrix,
    centers = k,
    nstart = 50
  )

  cluster_models[[as.character(k)]] <- fit

  cluster_fit_table <- rbind(
    cluster_fit_table,
    data.frame(
      n_clusters = k,
      total_withinss = fit$tot.withinss,
      between_total_ratio = fit$betweenss / fit$totss
    )
  )
}

cluster_fit_table

```

```

##   n_clusters total_withinss between_total_ratio
## 1         2      1163.5507         0.3564432
## 2         3       986.3183         0.4544700
## 3         4       876.1107         0.5154255


```

```

plot(
  cluster_fit_table$n_clusters,
  cluster_fit_table$total_withinss,
  type = "b",
  xlab = "Number of clusters",
  ylab = "Total within-cluster sum of squares",
  main = "Elbow plot for AI-use frequency clusters"
)

```

Elbow plot for AI-use frequency clusters



```
# Choose the number of clusters after inspecting the elbow plot and profiles.  
# Three clusters is the default because it is interpretable and comparable to  
# the earlier exploratory LCA result.
```

```
best_n_clusters <- 3  
best_cluster_model <- cluster_models[[as.character(best_n_clusters)]]
```

```
cluster_complete_data$ai_use_cluster <- factor(  
  best_cluster_model$cluster  
)
```

```
table(cluster_complete_data$ai_use_cluster)
```

```
##  
##  1  2  3  
## 51 87 89
```

```
cluster_profiles <- aggregate(  
  cluster_complete_data[, cluster_frequency_items],  
  by = list(ai_use_cluster = cluster_complete_data$ai_use_cluster),  
  FUN = mean,  
  na.rm = TRUE  
)
```

```
cluster_profiles
```

```
##   ai_use_cluster freq_information_guidance freq_writing_communication
## 1             1             3.352941             2.862745
## 2             2             5.977011             4.160920
## 3             3             6.078652             5.505618
##   freq_coding_math_technical freq_creative_work freq_learning_understanding
## 1             2.843137             2.352941             2.862745
## 2             2.195402             3.793103             5.310345
## 3             5.123596             5.629213             5.898876
##   freq_advice_decisions freq_summarizing_opinions freq_automation_repetitive
## 1             1.647059             2.411765             2.176471
## 2             4.379310             4.505747             2.816092
## 3             5.280899             5.786517             5.550562
```

The elbow plot suggested that a 3-cluster solution provided a reasonable balance between fit and interpretability. It explained approximately 47% of the variation in AI-use frequency patterns.

The clusters reflected distinct AI-use profiles.

Cluster 1: Low-frequency AI users, with relatively low use across all purposes.

Cluster 2: Broad high-frequency AI users, with high use across information, writing, coding, creative work, learning, advice, summarizing, and automation.

Cluster 3: Information/learning/advice users, with relatively high use for information, learning, advice, and summarizing, but lower use for coding, creative work, and automation.

```
model_cluster_ai_used <- glm(
  ai_used ~ ai_use_cluster,
  data = cluster_complete_data,
  family = binomial
)

summary(model_cluster_ai_used)
```

```
##
## Call:
## glm(formula = ai_used ~ ai_use_cluster, family = binomial, data = cluster_complete_data)
##
## Coefficients:
##             Estimate Std. Error z value Pr(>|z|)
## (Intercept)   -0.1967     0.2814  -0.699  0.484544
## ai_use_cluster2  1.4057     0.3796   3.703  0.000213 ***
## ai_use_cluster3  1.3717     0.3762   3.646  0.000266 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 278.84 on 226 degrees of freedom
## Residual deviance: 261.27 on 224 degrees of freedom
## AIC: 267.27
##
## Number of Fisher Scoring iterations: 4
```

```
drop1(model_cluster_ai_used, test = "Chisq")
```

```
## Single term deletions
##
## Model:
## ai_used ~ ai_use_cluster
##           Df Deviance    AIC    LRT Pr(>Chi)
## <none>           261.27 267.27
## ai_use_cluster  2   278.84 280.84 17.566 0.0001533 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
emmeans(model_cluster_ai_used, ~ ai_use_cluster, type = "response")
```

```
## ai_use_cluster prob  SE df asymp.LCL asymp.UCL
## 1             0.451 0.0697 Inf     0.321    0.588
## 2             0.770 0.0451 Inf     0.670    0.847
## 3             0.764 0.0450 Inf     0.665    0.841
##
## Confidence level used: 0.95
## Intervals are back-transformed from the logit scale
```

```
cluster_ai_users <- subset(
  cluster_complete_data,
  ai_used == TRUE & !is.na(task_id)
)

cluster_ai_users$task_id <- factor(cluster_ai_users$task_id)

model_cluster_jaccard <- lmer(
  jaccard_similarity_to_ai ~ ai_use_cluster + (1 | task_id),
  data = subset(cluster_ai_users, !is.na(jaccard_similarity_to_ai))
)

summary(model_cluster_jaccard)
```

```
## Linear mixed model fit by REML. t-tests use Satterthwaite's method [
## lmerModLmerTest]
## Formula: jaccard_similarity_to_ai ~ ai_use_cluster + (1 | task_id)
## Data: subset(cluster_ai_users, !is.na(jaccard_similarity_to_ai))
##
## REML criterion at convergence: 180.6
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -1.6785 -0.8369 -0.3121  0.9796  1.6087
##
## Random effects:
## Groups   Name                Variance Std.Dev.
## task_id  (Intercept)  0.02013   0.1419
## Residual                    0.16098   0.4012
```

```

## Number of obs: 158, groups: task_id, 18
##
## Fixed effects:
##              Estimate Std. Error      df t value Pr(>|t|)
## (Intercept)    0.44048    0.09266 109.28109   4.754 6.13e-06 ***
## ai_use_cluster2 0.08937    0.09915 148.21167   0.901   0.369
## ai_use_cluster3 0.09142    0.10192 154.96608   0.897   0.371
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Correlation of Fixed Effects:
##              (Intr) a_s_c2
## ai_s_clstr2 -0.800
## ai_s_clstr3 -0.811  0.734

anova(model_cluster_jaccard)

## Type III Analysis of Variance Table with Satterthwaite's method
##              Sum Sq Mean Sq NumDF DenDF F value Pr(>F)
## ai_use_cluster 0.15008 0.075041      2 154.06  0.4662 0.6283

emmeans(model_cluster_jaccard, pairwise ~ ai_use_cluster, adjust = "tukey")

## $emmeans
## ai_use_cluster emmean      SE      df lower.CL upper.CL
## 1              0.440 0.0927 109.3     0.257     0.624
## 2              0.530 0.0609  39.6     0.407     0.653
## 3              0.532 0.0605  38.6     0.409     0.654
##
## Degrees-of-freedom method: satterthwaite
## Confidence level used: 0.95
##
## $contrasts
## contrast              estimate      SE df t.ratio p.value
## ai_use_cluster1 - ai_use_cluster2 -0.08937 0.0991 148  -0.901  0.6403
## ai_use_cluster1 - ai_use_cluster3 -0.09142 0.1020 155  -0.897  0.6431
## ai_use_cluster2 - ai_use_cluster3 -0.00205 0.0733 155  -0.028  0.9996
##
## Degrees-of-freedom method: satterthwaite
## P value adjustment: tukey method for comparing a family of 3 estimates

model_cluster_embedding <- lmer(
  embedding_cosine_similarity_to_ai ~ ai_use_cluster + (1 | task_id),
  data = subset(cluster_ai_users, !is.na(embedding_cosine_similarity_to_ai))
)

summary(model_cluster_embedding)

## Linear mixed model fit by REML. t-tests use Satterthwaite's method [
## lmerModLmerTest]
## Formula: embedding_cosine_similarity_to_ai ~ ai_use_cluster + (1 | task_id)
## Data: subset(cluster_ai_users, !is.na(embedding_cosine_similarity_to_ai))

```

```
##
## REML criterion at convergence: 18.2
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -3.0747 -0.7351  0.2317  0.7835  1.5862
##
## Random effects:
##   Groups   Name      Variance Std.Dev.
## task_id (Intercept) 0.007559 0.08694
## Residual          0.056248 0.23717
## Number of obs: 157, groups: task_id, 18
##
## Fixed effects:
##              Estimate Std. Error      df t value Pr(>|t|)
## (Intercept)    0.74446    0.05510 104.01351  13.511   <2e-16 ***
## ai_use_cluster2 0.03418    0.05866 146.26727   0.583    0.561
## ai_use_cluster3 0.04610    0.06044 153.84268   0.763    0.447
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Correlation of Fixed Effects:
##              (Intr) a_s_c2
## ai_s_clstr2 -0.796
## ai_s_clstr3 -0.805  0.732
```

```
anova(model_cluster_embedding)
```

```
## Type III Analysis of Variance Table with Satterthwaite's method
##              Sum Sq Mean Sq NumDF DenDF F value Pr(>F)
## ai_use_cluster 0.032791 0.016396      2 152.76  0.2915 0.7476
```

```
emmeans(model_cluster_embedding, pairwise ~ ai_use_cluster, adjust = "tukey")
```

```
## $emmeans
##   ai_use_cluster emmean      SE      df lower.CL upper.CL
## 1                0.744 0.0551 104.0    0.635    0.854
## 2                0.779 0.0365  36.1    0.705    0.853
## 3                0.791 0.0364  35.8    0.717    0.864
##
## Degrees-of-freedom method: satterthwaite
## Confidence level used: 0.95
##
## $contrasts
##   contrast              estimate      SE      df t.ratio p.value
## ai_use_cluster1 - ai_use_cluster2 -0.0342 0.0587 146  -0.583  0.8296
## ai_use_cluster1 - ai_use_cluster3 -0.0461 0.0604 154  -0.763  0.7264
## ai_use_cluster2 - ai_use_cluster3 -0.0119 0.0436 154  -0.273  0.9596
##
## Degrees-of-freedom method: satterthwaite
## P value adjustment: tukey method for comparing a family of 3 estimates
```



```

cluster_ai_users$epistemic_ownership_binary <- ifelse(
  cluster_ai_users$epistemic_ownership <= 3, 0L,
  ifelse(cluster_ai_users$epistemic_ownership >= 5, 1L, NA_integer_)
)

model_cluster_ownership <- glmer(
  epistemic_ownership_binary ~ ai_use_cluster + (1 | task_id),
  data = subset(cluster_ai_users, !is.na(epistemic_ownership_binary)),
  family = binomial,
  control = glmerControl(optimizer = "bobyqa")
)

## boundary (singular) fit: see help('isSingular')

summary(model_cluster_ownership)

## Generalized linear mixed model fit by maximum likelihood (Laplace
## Approximation) [glmerMod]
## Family: binomial ( logit )
## Formula: epistemic_ownership_binary ~ ai_use_cluster + (1 | task_id)
## Data: subset(cluster_ai_users, !is.na(epistemic_ownership_binary))
## Control: glmerControl(optimizer = "bobyqa")
##
##          AIC          BIC      logLik -2*log(L)  df.resid
##      188.9       200.8      -90.4      180.9       140
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -0.8165 -0.8165 -0.6100  1.2247  1.6394
##
## Random effects:
##  Groups Name            Variance Std.Dev.
## task_id (Intercept) 0          0
## Number of obs: 144, groups: task_id, 18
##
## Fixed effects:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)   -0.8473    0.4880  -1.736   0.0825 .
## ai_use_cluster2 -0.1413    0.5691  -0.248   0.8039
## ai_use_cluster3  0.4418    0.5497   0.804   0.4216
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Correlation of Fixed Effects:
##              (Intr) a_s_c2
## ai_s_clstr2 -0.857
## ai_s_clstr3 -0.888  0.761
## optimizer (bobyqa) convergence code: 0 (OK)
## boundary (singular) fit: see help('isSingular')

drop1(model_cluster_ownership, test = "Chisq")

## boundary (singular) fit: see help('isSingular')

```

```
## Single term deletions
##
## Model:
## epistemic_ownership_binary ~ ai_use_cluster + (1 | task_id)
##               npar      AIC      LRT Pr(Chi)
## <none>                188.89
## ai_use_cluster      2 187.32 2.4266 0.2972

emmeans(model_cluster_ownership, pairwise ~ ai_use_cluster, adjust = "tukey", type = "response")

## $emmeans
##   ai_use_cluster  prob      SE df asymp.LCL asymp.UCL
## 1                0.300 0.1020 Inf    0.141    0.527
## 2                0.271 0.0579 Inf    0.173    0.398
## 3                0.400 0.0608 Inf    0.289    0.523
##
## Confidence level used: 0.95
## Intervals are back-transformed from the logit scale
##
## $contrasts
##   contrast                odds.ratio      SE df null z.ratio p.value
## ai_use_cluster1 / ai_use_cluster2      1.152 0.655 Inf    1    0.248 0.9666
## ai_use_cluster1 / ai_use_cluster3      0.643 0.353 Inf    1   -0.804 0.7007
## ai_use_cluster2 / ai_use_cluster3      0.558 0.216 Inf    1   -1.506 0.2879
##
## P value adjustment: tukey method for comparing a family of 3 estimates
## Tests are performed on the log odds ratio scale
```

Cluster membership significantly predicted whether participants used the optional AI assistant, $\chi^2(2) = 17.57$, $p < .001$. Participants in Cluster 1, the low-frequency AI-use group, had an estimated AI-use probability of .45. Participants in Cluster 2, the broad high-frequency AI-use group, had a higher estimated probability of AI use (.77) and were significantly more likely to use AI than Cluster 1, $b = 1.41$, $SE = 0.38$, $z = 3.70$, $p < .001$. Participants in Cluster 3, the information/learning/advice group, also had a higher estimated probability of AI use (.76) and were significantly more likely to use AI than Cluster 1, $b = 1.37$, $SE = 0.38$, $z = 3.65$, $p < .001$.

However, among participants who used AI, cluster membership did not significantly predict lexical similarity to the AI response, $F(2, 154.06) = 0.47$, $p = .628$. Estimated Jaccard similarity was similar across clusters: Cluster 1, $M = .440$; Cluster 2, $M = .530$; Cluster 3, $M = .532$. Tukey-adjusted pairwise comparisons were all non-significant, all $p \geq .640$.

Cluster membership also did not significantly predict embedding-based semantic similarity to the AI response, $F(2, 152.76) = 0.29$, $p = .748$. Estimated semantic similarity was similar across clusters: Cluster 1, $M = .744$; Cluster 2, $M = .779$; Cluster 3, $M = .791$. Tukey-adjusted pairwise comparisons were again non-significant, all $p \geq .726$.

Finally, cluster membership did not significantly predict binary epistemic ownership, $\chi^2(2) = 2.43$, $p = .297$. Estimated probabilities of primarily own-thinking ownership were .30 for Cluster 1, .27 for Cluster 2, and .40 for Cluster 3. Tukey-adjusted pairwise comparisons were not significant, all $p \geq .288$.

```
# -----
# 1. Cluster profile heatmap
# -----
```

```


cluster_profile_long <- do.call(rbind, lapply(
  seq_len(nrow(cluster_profiles)),
  function(i) {
    data.frame(
      ai_use_cluster = cluster_profiles$ai_use_cluster[i],
      frequency_item = cluster_frequency_items,
      mean_frequency = as.numeric(cluster_profiles[i, cluster_frequency_items]),
      stringsAsFactors = FALSE
    )
  }
))

cluster_profile_long$frequency_label <- factor(
  cluster_profile_long$frequency_item,
  levels = cluster_frequency_items,
  labels = c(
    "Info / guidance",
    "Writing / communication",
    "Coding / technical",
    "Creative work",
    "Learning",
    "Advice / decisions",
    "Summarizing",
    "Automation"
  )
)

ggplot(cluster_profile_long, aes(
  x = frequency_label,
  y = ai_use_cluster,
  fill = mean_frequency
)) +
  geom_tile(color = "white", linewidth = 0.6) +
  geom_text(
    aes(label = sprintf("%.1f", mean_frequency)),
    size = 3.4,
    fontface = "bold"
  ) +
  scale_fill_gradient(
    low = "#f2f2f2",
    high = "#4472C4",
    limits = c(1, 7)
  ) +
  labs(
    title = "AI Use Frequency Profiles by Cluster",
    x = NULL,
    y = "AI use cluster",
    fill = "Mean\nfrequency"
  ) +
  theme_classic(base_size = 11) +
  theme(
    plot.title = element_text(face = "bold", size = 13),
    axis.text.x = element_text(angle = 35, hjust = 1)
  )

```

)



```
# -----
# 2. PCA visualization of clusters
# -----

cluster_pca <- prcomp(cluster_matrix, center = FALSE, scale. = FALSE)

cluster_pca_data <- data.frame(
  pc1 = cluster_pca$x[, 1],
  pc2 = cluster_pca$x[, 2],
  ai_use_cluster = cluster_complete_data$ai_use_cluster
)


pca_var <- round(100 * summary(cluster_pca)$importance[2, 1:2], 1)

ggplot(cluster_pca_data, aes(
  x = pc1,
  y = pc2,
  color = ai_use_cluster
)) +
  geom_point(alpha = 0.75, size = 2.2) +
  stat_ellipse(linewidth = 0.8, alpha = 0.6) +
  labs(
```

```

title = "Participant Clusters Based on AI Use Frequencies",
x = paste0("PC1 (", pca_var[1], "% variance)"),
y = paste0("PC2 (", pca_var[2], "% variance)"),
color = "Cluster"
) +
theme_classic(base_size = 11) +
theme(
  plot.title = element_text(face = "bold", size = 13)
)

```



```

# -----
# 3. Cluster relationships with outcomes
# -----

cluster_outcome_summary <- do.call(rbind, lapply(
  split(cluster_complete_data, cluster_complete_data$ai_use_cluster),
  function(sub) {
    data.frame(
      ai_use_cluster = as.character(sub$ai_use_cluster[1]),
      ai_use_rate = mean(sub$ai_used, na.rm = TRUE),
      mean_jaccard = mean(
        sub$jaccard_similarity_to_ai[sub$ai_used == TRUE],
        na.rm = TRUE
      ),
    ),
  )

```

```

    mean_embedding = mean(
      sub$embedding_cosine_similarity_to_ai[sub$ai_used == TRUE],
      na.rm = TRUE
    ),
    mean_epistemic_ownership = mean(
      sub$epistemic_ownership[sub$ai_used == TRUE],
      na.rm = TRUE
    ),
    n_all = nrow(sub),
    n_ai_users = sum(sub$ai_used == TRUE, na.rm = TRUE),
    stringsAsFactors = FALSE
  )
}
))

cluster_outcome_summary$ai_use_cluster <- factor(
  cluster_outcome_summary$ai_use_cluster,
  levels = sort(unique(cluster_outcome_summary$ai_use_cluster))
)

p_cluster_use <- ggplot(
  cluster_outcome_summary,
  aes(x = ai_use_cluster, y = ai_use_rate)
) +
  geom_col(fill = "#70AD47", width = 0.6) +
  geom_text(
    aes(label = paste0(round(ai_use_rate * 100), "%")),
    vjust = -0.35, size = 3.5, fontface = "bold"
  ) +
  geom_text(
    aes(label = paste0("n=", n_all), y = 0.05),
    color = "white", size = 3.1
  ) +
  scale_y_continuous(
    limits = c(0, 1.05),
    expand = c(0, 0),
    labels = function(x) paste0(round(x * 100), "%")
  ) +
  labs(
    title = "AI Use Rate",
    x = "Cluster",
    y = "% Used AI"
  ) +
  theme_classic(base_size = 11) +
  theme(plot.title = element_text(face = "bold", size = 11))

p_cluster_jaccard <- ggplot(
  cluster_outcome_summary,
  aes(x = ai_use_cluster, y = mean_jaccard)
) +
  geom_col(fill = "#4472C4", width = 0.6) +
  geom_text(
    aes(label = sprintf("%.3f", mean_jaccard)),

```

```

    vjust = -0.35, size = 3.5, fontface = "bold"
  ) +
  geom_text(
    aes(label = paste0("n=", n_ai_users), y = 0.035),
    color = "white", size = 3.1
  ) +
  scale_y_continuous(limits = c(0, 1), expand = c(0, 0)) +
  labs(
    title = "Mean Jaccard Similarity",
    x = "Cluster",
    y = "Mean Jaccard"
  ) +
  theme_classic(base_size = 11) +
  theme(plot.title = element_text(face = "bold", size = 11))

p_cluster_embedding <- ggplot(
  cluster_outcome_summary,
  aes(x = ai_use_cluster, y = mean_embedding)
) +
  geom_col(fill = "#5B9BD5", width = 0.6) +
  geom_text(
    aes(label = sprintf("%.3f", mean_embedding)),
    vjust = -0.35, size = 3.5, fontface = "bold"
  ) +
  geom_text(
    aes(label = paste0("n=", n_ai_users), y = 0.035),
    color = "white", size = 3.1
  ) +
  scale_y_continuous(limits = c(0, 1.1), expand = c(0, 0)) +
  labs(
    title = "Mean Embedding Similarity",
    x = "Cluster",
    y = "Mean embedding cosine"
  ) +
  theme_classic(base_size = 11) +
  theme(plot.title = element_text(face = "bold", size = 11))

p_cluster_ownership <- ggplot(
  cluster_outcome_summary,
  aes(x = ai_use_cluster, y = mean_epistemic_ownership)
) +
  geom_col(fill = "#A64D79", width = 0.6) +
  geom_text(
    aes(label = sprintf("%.2f", mean_epistemic_ownership)),
    vjust = -0.35, size = 3.5, fontface = "bold"
  ) +
  geom_text(
    aes(label = paste0("n=", n_ai_users), y = 0.35),
    color = "white", size = 3.1
  ) +
  scale_y_continuous(limits = c(0, 7), expand = c(0, 0)) +
  labs(
    title = "Mean Epistemic Ownership",


```

```

x = "Cluster",
y = "Ownership"
) +
theme_classic(base_size = 11) +
theme(plot.title = element_text(face = "bold", size = 11))

grid.arrange(
  p_cluster_use,
  p_cluster_jaccard,
  p_cluster_embedding,
  p_cluster_ownership,
  ncol = 2,
  top = grid::textGrob(
    "AI Use Clusters and Behavioral Outcomes",
    gp = grid::gpar(fontface = "bold", fontsize = 13)
  )
)

```



Altogether: Participants' general AI-use profiles predicted whether they chose to consult the AI assistant, $\chi^2(2) = 17.57$, $p < .001$, but not how much they incorporated the AI response once they used it, whether uptake was measured lexically, semantically, or through binary epistemic ownership.

Analysis 7: Domain-Specific AI Use Frequency vs. Task Behavior

This analysis asks whether self-reported AI use frequency in the specific domain of the assigned task predicts actual AI use and uptake in that task. For example, does someone who says they frequently use AI for writing tasks actually consult AI and copy more in a writing task?

Rather than constructing a category-to-item mapping, this analysis uses `self_reported_freq_in_assigned_category`, which directly records how often each participant reported using AI for the type of task they were assigned. This is a more direct measure and avoids the assumption that a single frequency item maps cleanly to each task category.

```
domain_freq_data <- model_data

domain_freq_data$z_self_reported_freq <- z(
  domain_freq_data$self_reported_freq_in_assigned_category
)

summary(domain_freq_data$self_reported_freq_in_assigned_category)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##      2.000   6.000   6.000   6.084   7.000   7.000
```

```
table(domain_freq_data$assigned_category, useNA = "ifany")
```

```
##
## academic creative decisions info learning writing
##          36          35          36          41          42          37
```

```
model_dfreq_ai_used <- glmer(
  ai_used ~ z_self_reported_freq +
    z_hours_per_week +
    (1 | task_id),
  data = subset(domain_freq_data, !is.na(z_self_reported_freq)),
  family = binomial,
  control = glmerControl(optimizer = "bobyqa")
)
```

```
## boundary (singular) fit: see help('isSingular')
```

```
summary(model_dfreq_ai_used)
```

```
## Generalized linear mixed model fit by maximum likelihood (Laplace
## Approximation) [glmerMod]
## Family: binomial ( logit )
## Formula: ai_used ~ z_self_reported_freq + z_hours_per_week + (1 | task_id)
## Data: subset(domain_freq_data, !is.na(z_self_reported_freq))
## Control: glmerControl(optimizer = "bobyqa")
##
##      AIC      BIC    logLik -2*log(L)  df.resid
##      267.0    280.7    -129.5    259.0      223
##
## Scaled residuals:
```

```
##      Min      1Q  Median      3Q      Max
## -2.4987 -0.8759  0.5059  0.6595  1.5757
##
## Random effects:
##   Groups Name      Variance Std.Dev.
##   task_id (Intercept) 0          0
## Number of obs: 227, groups: task_id, 18
##
## Fixed effects:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)      0.8861    0.1530   5.792 6.97e-09 ***
## z_self_reported_freq 0.5998    0.1584   3.785 0.000153 ***
## z_hours_per_week    0.1333    0.1771   0.752 0.451842
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Correlation of Fixed Effects:
##              (Intr) z_sl__
## z_slf_rprt_   0.103
## z_hrs_pr_wk  0.073 -0.268
## optimizer (bobyqa) convergence code: 0 (OK)
## boundary (singular) fit: see help('isSingular')
```

```
fixed_effects_table(model_dfreq_ai_used, exponentiate = TRUE)
```

```
##              term estimate std_error      p_value odds_ratio
## 1      (Intercept) 0.8860543 0.1529866 6.967258e-09   2.425540
## 2 z_self_reported_freq 0.5997989 0.1584489 1.534364e-04   1.821752
## 3      z_hours_per_week 0.1332590 0.1771243 4.518422e-01   1.142546
```

Self-reported domain-specific AI use frequency significantly predicted whether participants used the AI assistant during the task, OR = 1.82, $p < .001$. A one-standard-deviation increase in reported AI use for the assigned task domain was associated with higher odds of consulting the optional assistant, controlling for hours per week. Hours per week was not significant, OR = 1.14, $p = .452$.

```
domain_freq_ai_users <- subset(
  domain_freq_data,
  ai_used == TRUE & !is.na(z_self_reported_freq) & !is.na(jaccard_similarity_to_ai)
)

model_dfreq_similarity <- lmer(
  jaccard_similarity_to_ai ~ z_self_reported_freq +
    z_hours_per_week +
    (1 | task_id),
  data = domain_freq_ai_users
)

summary(model_dfreq_similarity)
```

```
## Linear mixed model fit by REML. t-tests use Satterthwaite's method [
## lmerModLmerTest]
## Formula: jaccard_similarity_to_ai ~ z_self_reported_freq + z_hours_per_week +
```

```
##      (1 | task_id)
##      Data: domain_freq_ai_users
##
## REML criterion at convergence: 184.6
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -1.6443 -0.8696 -0.2867  0.9693  1.6106
##
## Random effects:
##      Groups   Name      Variance Std.Dev.
## task_id (Intercept) 0.01895  0.1377
## Residual          0.16237  0.4030
## Number of obs: 158, groups: task_id, 18
##
## Fixed effects:
##              Estimate Std. Error      df t value Pr(>|t|)
## (Intercept)      0.51906    0.04673   16.89077   11.108 3.48e-09 ***
## z_self_reported_freq -0.01318    0.04250  154.22434   -0.310   0.757
## z_hours_per_week     0.01125    0.03238  149.82863    0.347   0.729
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Correlation of Fixed Effects:
##              (Intr) z_sl__
## z_slf_rprt_ -0.176
## z_hrs_pr_wk -0.036 -0.186
```

```
fixed_effects_table(model_dfreq_similarity)
```

```
##              term      estimate std_error      p_value
## 1      (Intercept) 0.51905897 0.04672640 3.483237e-09
## 2 z_self_reported_freq -0.01317750 0.04250093 7.569400e-01
## 3      z_hours_per_week 0.01125076 0.03238021 7.287349e-01
```

Among participants who used the assistant, domain-specific AI use frequency did not predict lexical similarity to the AI response, $b = -0.013$, $p = .757$. Hours per week was also not significant, $b = 0.011$, $p = .729$. This is an AI-user-only model, so it asks whether domain frequency predicts copying after participants have already chosen to use AI.

```
# Use self_reported_freq_in_assigned_category on embedding_analysis_data.
embedding_analysis_data$z_self_reported_freq <- z(
  embedding_analysis_data$self_reported_freq_in_assigned_category
)

domain_freq_embedding_data <- subset(
  embedding_analysis_data,
  !is.na(z_self_reported_freq) & !is.na(embedding_cosine_similarity_to_ai)
)

model_dfreq_embedding <- lmer(
  embedding_cosine_similarity_to_ai ~ z_self_reported_freq +
    z_hours_per_week +
```

```

    (1 | task_id),
    data = domain_freq_embedding_data
  )

summary(model_dfreq_embedding)

```

```

## Linear mixed model fit by REML. t-tests use Satterthwaite's method [
## lmerModLmerTest]
## Formula:
## embedding_cosine_similarity_to_ai ~ z_self_reported_freq + z_hours_per_week +
## (1 | task_id)
## Data: domain_freq_embedding_data
##
## REML criterion at convergence: 22.1
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -3.0282 -0.7537  0.2385  0.7331  1.5901
##
## Random effects:
## Groups Name Variance Std.Dev.
## task_id (Intercept) 0.006661 0.08162
## Residual 0.056757 0.23824
## Number of obs: 157, groups: task_id, 18
##
## Fixed effects:
##              Estimate Std. Error      df t value Pr(>|t|)
## (Intercept)    0.779022   0.027215  14.802363  28.625 2.26e-14 ***
## z_self_reported_freq -0.012016   0.020262  152.960807  -0.593   0.554
## z_hours_per_week    0.004629   0.019954  148.418030   0.232   0.817
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Correlation of Fixed Effects:
##              (Intr) z_sl__
## z_slf_rprt_ -0.010
## z_hrs_pr_wk -0.013 -0.188

```

```

fixed_effects_table(model_dfreq_embedding)

```

```

##              term      estimate std_error    p_value
## 1 (Intercept) 0.779021717 0.02721507 2.259196e-14
## 2 z_self_reported_freq -0.012016468 0.02026217 5.540239e-01
## 3 z_hours_per_week 0.004628941 0.01995416 8.168740e-01

```

```

domain_freq_embedding_data$z_embedding_cosine_similarity_to_ai <- z(
  domain_freq_embedding_data$embedding_cosine_similarity_to_ai
)

```

Domain-specific AI use frequency also did not predict embedding-based semantic similarity to the AI response among AI users, $b = -0.012$, $p = .554$. Hours per week was not significant, $b = 0.005$, $p = .817$.

```

domain_freq_data$epistemic_ownership_binary <- ifelse(
  domain_freq_data$epistemic_ownership <= 3, 0L,
  ifelse(domain_freq_data$epistemic_ownership >= 5, 1L, NA_integer_)
)

domain_freq_ownership_data <- subset(
  domain_freq_data,
  ai_used == TRUE & !is.na(z_self_reported_freq) & !is.na(epistemic_ownership_binary)
)

model_dfreq_ownership <- glmer(
  epistemic_ownership_binary ~ z_jaccard_similarity_to_ai +
    z_self_reported_freq +
    z_hours_per_week +
    (1 | task_id),
  data = domain_freq_ownership_data,
  family = binomial,
  control = glmerControl(optimizer = "bobyqa")
)

```

```
## boundary (singular) fit: see help('isSingular')
```

```
summary(model_dfreq_ownership)
```

```

## Generalized linear mixed model fit by maximum likelihood (Laplace
## Approximation) [glmerMod]
## Family: binomial ( logit )
## Formula:
## epistemic_ownership_binary ~ z_jaccard_similarity_to_ai + z_self_reported_freq +
## z_hours_per_week + (1 | task_id)
## Data: domain_freq_ownership_data
## Control: glmerControl(optimizer = "bobyqa")
##
##      AIC      BIC    logLik -2*log(L)  df.resid
##    159.3    174.1    -74.6    149.3      139
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -1.7142 -0.6034 -0.3321  0.8239  3.5480
##
## Random effects:
## Groups Name Variance Std.Dev.
## task_id (Intercept) 0 0
## Number of obs: 144, groups: task_id, 18
##
## Fixed effects:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    -0.5557    0.2092  -2.656  0.00791 **
## z_jaccard_similarity_to_ai -1.1036    0.2189  -5.042 4.62e-07 ***
## z_self_reported_freq      0.1909    0.2462   0.775  0.43825
## z_hours_per_week        0.1364    0.1944   0.702  0.48289
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

```
##
## Correlation of Fixed Effects:
##          (Intr) z_j___ z_sl__
## z_jccrd_s__  0.058
## z_slf_rprt_ -0.272 -0.008
## z_hrs_pr_wk -0.045 -0.133 -0.173
## optimizer (bobyqa) convergence code: 0 (OK)
## boundary (singular) fit: see help('isSingular')
```

```
fixed_effects_table(model_dfreq_ownership, exponentiate = TRUE)
```

```
##              term      estimate std_error      p_value odds_ratio
## 1      (Intercept) -0.5556580  0.2092207  7.910960e-03  0.5736947
## 2 z_jaccard_similarity_to_ai -1.1035684  0.2188948  4.617809e-07  0.3316854
## 3      z_self_reported_freq  0.1908550  0.2462140  4.382456e-01  1.2102840
## 4      z_hours_per_week    0.1364096  0.1944096  4.828907e-01  1.1461512
```

```
domain_freq_embedding_data$epistemic_ownership_binary <- ifelse(
  domain_freq_embedding_data$epistemic_ownership <= 3, 0L,
  ifelse(domain_freq_embedding_data$epistemic_ownership >= 5, 1L, NA_integer_)
)
```

```
domain_freq_embedding_ownership_data <- subset(
  domain_freq_embedding_data,
  !is.na(epistemic_ownership_binary) &
  !is.na(z_embedding_cosine_similarity_to_ai)
)
```

```
model_dfreq_ownership_embedding <- glmer(
  epistemic_ownership_binary ~ z_embedding_cosine_similarity_to_ai +
    z_self_reported_freq +
    z_hours_per_week +
    (1 | task_id),
  data = domain_freq_embedding_ownership_data,
  family = binomial,
  control = glmerControl(optimizer = "bobyqa")
)
```

```
## boundary (singular) fit: see help('isSingular')
```

```
summary(model_dfreq_ownership_embedding)
```

```
## Generalized linear mixed model fit by maximum likelihood (Laplace
## Approximation) [glmerMod]
## Family: binomial ( logit )
## Formula: epistemic_ownership_binary ~ z_embedding_cosine_similarity_to_ai +
##          z_self_reported_freq + z_hours_per_week + (1 | task_id)
## Data: domain_freq_embedding_ownership_data
## Control: glmerControl(optimizer = "bobyqa")
##
##          AIC          BIC      logLik -2*log(L)  df.resid
##        165.3        180.1       -77.6     155.3      138
```

```
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -2.9137 -0.5415 -0.4500  0.8065  2.4598
##
## Random effects:
##   Groups Name            Variance Std.Dev.
##   task_id (Intercept) 0          0
## Number of obs: 143, groups: task_id, 18
##
## Fixed effects:
##                                     Estimate Std. Error z value Pr(>|z|)
## (Intercept)                      -0.78720    0.19851  -3.966 7.32e-05 ***
## z_embedding_cosine_similarity_to_ai -0.91190    0.20034  -4.552 5.32e-06 ***
## z_self_reported_freq               0.10533    0.19966   0.528  0.598
## z_hours_per_week                   0.07703    0.19955   0.386  0.699
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Correlation of Fixed Effects:
##              (Intr) z_----- z_sl__
## z_mbddn_---- 0.115
## z_slf_rprt_ -0.054 0.040
## z_hrs_pr_wk -0.016 -0.065 -0.162
## optimizer (bobyqa) convergence code: 0 (OK)
## boundary (singular) fit: see help('isSingular')
```

```
fixed_effects_table(model_dfreq_ownership_embedding, exponentiate = TRUE)
```

```
##              term      estimate std_error    p_value
## 1              (Intercept) -0.78720411 0.1985098 7.322144e-05
## 2 z_embedding_cosine_similarity_to_ai -0.91189777 0.2003436 5.322189e-06
## 3              z_self_reported_freq 0.10532841 0.1996556 5.978114e-01
## 4              z_hours_per_week 0.07702609 0.1995498 6.994972e-01
## odds_ratio
## 1 0.4551155
## 2 0.4017610
## 3 1.1110754
## 4 1.0800703
```

Domain-specific AI use frequency did not predict binary epistemic ownership whether AI uptake was measured by Jaccard similarity (OR = 1.21, $p = .438$) or by embedding similarity (OR = 1.11, $p = .598$). **In both models, similarity to the AI response was the main predictor of lower own-thinking ownership: Jaccard similarity, OR = 0.33, $p < .001$; embedding similarity, OR = 0.40, $p < .001$.**

Analysis 8: AI Engagement And Attention

This exploratory analysis asks whether participants who spent more time with the AI response, or reported reading it more carefully, showed greater uptake of the AI response or different epistemic ownership ratings. Because timing variables are usually skewed, millisecond variables are converted to seconds and log-transformed with `log1p()`.

```

engagement_data <- subset(
  embedding_similarity_data,
  ai_used == TRUE & !is.na(task_id)
)

engagement_data$task_id <- factor(engagement_data$task_id)

engagement_data$ai_reading_depth <- factor(
  engagement_data$ai_reading_depth,
  levels = c("not_read", "glanced", "skimmed", "read_carefully"),
  ordered = TRUE
)

# Binary recodings.
engagement_data$epistemic_ownership_binary <- ifelse(
  engagement_data$epistemic_ownership <= 3, 0L,
  ifelse(engagement_data$epistemic_ownership >= 5, 1L, NA_integer_)
)
engagement_data$read_carefully <- as.integer(
  engagement_data$ai_reading_depth == "read_carefully"
)

engagement_data$response_writing_time_ms <-
  engagement_data$response_submitted_ms - engagement_data$response_started_ms

engagement_data$response_writing_time_ms[
  engagement_data$response_writing_time_ms < 0
] <- NA_real_

engagement_data$log_time_before_ai_sec <- log1p(
  engagement_data$time_before_ai_ms / 1000
)

engagement_data$log_ai_visible_duration_sec <- log1p(
  engagement_data$ai_visible_duration_ms / 1000
)

engagement_data$log_response_writing_time_sec <- log1p(
  engagement_data$response_writing_time_ms / 1000
)

engagement_data$log_task_duration_sec <- log1p(
  engagement_data$task_duration_ms / 1000
)

engagement_data$z_log_time_before_ai_sec <- z(
  engagement_data$log_time_before_ai_sec
)

engagement_data$z_log_ai_visible_duration_sec <- z(
  engagement_data$log_ai_visible_duration_sec
)

```



```
engagement_data$z_log_response_writing_time_sec <- z(
  engagement_data$log_response_writing_time_sec
)

engagement_data$z_log_task_duration_sec <- z(
  engagement_data$log_task_duration_sec
)

table(engagement_data$ai_reading_depth, useNA = "ifany")
```

```
##
##      not_read      glanced      skimmed read_carefully
##           9         14         45         90
```

```
summary(engagement_data$ai_visible_duration_ms)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##  14993   70970  149338  217684  255126  4152051
```

```
summary(engagement_data$response_writing_time_ms)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.   NA's
##    473    3595   46510   80431  134905   340073     1
```

Analysis 8A: Does Self-Reported Reading Depth Match Viewing Time?

This check asks whether people who reported reading the AI response carefully also had the AI response visible for longer.

```
model_reading_depth_visible_time <- lmer(
  log_ai_visible_duration_sec ~ read_carefully +
    (1 | task_id),
  data = subset(
    engagement_data,
    !is.na(log_ai_visible_duration_sec) & !is.na(read_carefully)
  )
)

summary(model_reading_depth_visible_time)
```

```
## Linear mixed model fit by REML. t-tests use Satterthwaite's method [
## lmerModLmerTest]
## Formula: log_ai_visible_duration_sec ~ read_carefully + (1 | task_id)
## Data: subset(engagement_data, !is.na(log_ai_visible_duration_sec) &
##       !is.na(read_carefully))
##
## REML criterion at convergence: 424.2
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
```


```
## -2.0736 -0.6348 0.0786 0.6408 3.8822
##
## Random effects:
## Groups Name Variance Std.Dev.
## task_id (Intercept) 0.09308 0.3051
## Residual 0.77752 0.8818
## Number of obs: 158, groups: task_id, 18
##
## Fixed effects:
## Estimate Std. Error df t value Pr(>|t|)
## (Intercept) 4.6816 0.1311 40.7406 35.711 <2e-16 ***
## read_carefully 0.3499 0.1465 153.6784 2.388 0.0182 *
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Correlation of Fixed Effects:
## (Intr)
## read_crflly -0.637
```

```
fixed_effects_table(model_reading_depth_visible_time)
```

```
## term estimate std_error p_value
## 1 (Intercept) 4.6816073 0.1310980 2.276878e-32
## 2 read_carefully 0.3499444 0.1465427 1.815458e-02
```

```
boxplot(
  log_ai_visible_duration_sec ~ ai_reading_depth,
  data = engagement_data,
  xlab = "AI reading depth",
  ylab = "Log AI visible duration (seconds)",
  main = "AI visible time by self-reported reading depth"
)
```

AI visible time by self-reported reading depth



People who reported reading carefully had the AI response visible significantly longer than other AI users, $b = 0.350$, $p = .018$. This supports the self-reported reading-depth item as behaviorally meaningful, although it is still an imperfect proxy for attention.

Analysis 8B: Does Engagement Predict Lexical Similarity?

Outcome: jaccard_similarity_to_ai

Model type: linear mixed regression

Predictors: ai_reading_depth, ai_visible_duration_ms, time_before_ai_ms

Random intercept: task_id

Sample: AI users

```
engagement_jaccard_data <- subset(
  engagement_data,
  !is.na(jaccard_similarity_to_ai) &
  !is.na(read_carefully) &
  !is.na(z_log_ai_visible_duration_sec) &
  !is.na(z_log_time_before_ai_sec)
)

model_engagement_jaccard <- lmer(
  jaccard_similarity_to_ai ~ read_carefully +
  z_log_ai_visible_duration_sec +
  z_log_time_before_ai_sec +
  (1 | task_id),
```

```

data = engagement_jaccard_data
)

summary(model_engagement_jaccard)

## Linear mixed model fit by REML. t-tests use Satterthwaite's method [
## lmerModLmerTest]
## Formula:
## jaccard_similarity_to_ai ~ read_carefully + z_log_ai_visible_duration_sec +
##      z_log_time_before_ai_sec + (1 | task_id)
## Data: engagement_jaccard_data
##
## REML criterion at convergence: 155.3
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -2.67357 -0.71625 -0.06888  0.80542  2.72179
##
## Random effects:
## Groups Name Variance Std.Dev.
## task_id (Intercept) 0.01634 0.1278
## Residual 0.13129 0.3623
## Number of obs: 158, groups: task_id, 18
##
## Fixed effects:
##              Estimate Std. Error      df t value Pr(>|t|)
## (Intercept)      0.57931    0.05497   39.49285   10.539 4.86e-13
## read_carefully    -0.11486    0.06219  151.70802    -1.847  0.0667
## z_log_ai_visible_duration_sec -0.15075    0.03143  153.17424    -4.796 3.80e-06
## z_log_time_before_ai_sec    -0.04092    0.03116  153.92847    -1.313  0.1912
##
## (Intercept)          ***
## read_carefully        .
## z_log_ai_visible_duration_sec ***
## z_log_time_before_ai_sec
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Correlation of Fixed Effects:
##              (Intr) rd_crf z_lg__v__
## read_crflly -0.646
## z_lg__vsb__  0.120 -0.156
## z_lg_tm_b__  0.092 -0.163 -0.168

fixed_effects_table(model_engagement_jaccard)

```

```

##              term      estimate std_error      p_value
## 1              (Intercept) 0.5793135 0.05496926 4.862625e-13
## 2              read_carefully -0.1148629 0.06218993 6.669956e-02
## 3 z_log_ai_visible_duration_sec -0.1507500 0.03143097 3.800506e-06
## 4              z_log_time_before_ai_sec -0.0409166 0.03116472 1.911658e-01

```

Reading carefully was marginally associated with lower lexical similarity, $b = -0.115$, $p = .067$. **Longer AI visible duration significantly predicted lower Jaccard similarity**, $b = -0.151$, $p < .001$. This

negative association should be interpreted cautiously: it is likely a reverse-causality or strategy-difference pattern. Participants who copied the AI response verbatim (Jaccard = 1.0, $n = 55$) had no reason to keep the AI visible, whereas participants who spent longer viewing the response may have selected, edited, or paraphrased more. Time before viewing AI was not significant, $b = -0.041$, $p = .191$.

Analysis 8C: Does Engagement Predict Semantic Similarity?

Outcome: embedding_cosine_similarity_to_ai

Model type: linear mixed regression

Predictors: ai_reading_depth, ai_visible_duration_ms, time_before_ai_ms

Random intercept: task_id

Sample: AI users with embedding similarity scores

```
engagement_embedding_data <- subset(
  engagement_data,
  !is.na(embedding_cosine_similarity_to_ai) &
  !is.na(read_carefully) &
  !is.na(z_log_ai_visible_duration_sec) &
  !is.na(z_log_time_before_ai_sec)
)

model_engagement_embedding <- lmer(
  embedding_cosine_similarity_to_ai ~ read_carefully +
    z_log_ai_visible_duration_sec +
    z_log_time_before_ai_sec +
    (1 | task_id),
  data = engagement_embedding_data
)

summary(model_engagement_embedding)
```



```
## Linear mixed model fit by REML. t-tests use Satterthwaite's method [
## lmerModLmerTest]
## Formula:
## embedding_cosine_similarity_to_ai ~ read_carefully + z_log_ai_visible_duration_sec +
##      z_log_time_before_ai_sec + (1 | task_id)
##      Data: engagement_embedding_data
##
## REML criterion at convergence: 5.1
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -3.6177 -0.4728  0.2252  0.6311  1.9004
##
## Random effects:
##  Groups   Name                Variance Std.Dev.
## task_id  (Intercept)  0.007374  0.08587
## Residual                    0.048927  0.22119
## Number of obs: 157, groups: task_id, 18
##
## Fixed effects:
##
##              Estimate Std. Error    df t value Pr(>|t|)
## (Intercept)    0.80509    0.03467 36.55770  23.219 < 2e-16
```

```
## read_carefully          -0.04677    0.03831 149.84274 -1.221 0.22405
## z_log_ai_visible_duration_sec -0.05927    0.01937 152.83035 -3.059 0.00262
## z_log_time_before_ai_sec    -0.04566    0.01923 152.68718 -2.374 0.01883
##
## (Intercept)                ***
## read_carefully
## z_log_ai_visible_duration_sec **
## z_log_time_before_ai_sec    *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Correlation of Fixed Effects:
##          (Intr) rd_crf z_lg__v__
## read_crfly -0.626
## z_lg__vsb__ 0.120 -0.164
## z_lg_tm_b__ 0.089 -0.171 -0.164
```

```
fixed_effects_table(model_engagement_embedding)
```

```
##          term      estimate std_error    p_value
## 1          (Intercept) 0.80509151 0.03467334 1.767199e-23
## 2      read_carefully -0.04676631 0.03830563 2.240517e-01
## 3 z_log_ai_visible_duration_sec -0.05926574 0.01937441 2.623346e-03
## 4      z_log_time_before_ai_sec -0.04565783 0.01923084 1.883013e-02
```

Reading carefully did not significantly predict semantic similarity, $b = -0.047$, $p = .224$. **Longer AI visible duration predicted lower embedding similarity, $b = -0.059$, $p = .0026$** , and longer time before viewing AI also predicted lower embedding similarity, $b = -0.046$, $p = .019$. These timing effects should be interpreted as strategy differences rather than evidence that engagement itself causally reduces uptake.

Analysis 8D: Does Engagement Predict Epistemic Ownership?

Outcome: `epistemic_ownership_binary` (0 = AI's thinking, 1 = own thinking)

Model type: mixed logistic regression

Predictors: `read_carefully`, `z_log_ai_visible_duration_sec`, `z_log_time_before_ai_sec`

Random intercept: `task_id`

Sample: AI users

```
engagement_ownership_data <- subset(
  engagement_data,
  !is.na(epistemic_ownership_binary) &
  !is.na(read_carefully) &
  !is.na(z_log_ai_visible_duration_sec) &
  !is.na(z_log_time_before_ai_sec)
)

model_engagement_ownership <- glmer(
  epistemic_ownership_binary ~ read_carefully +
    z_log_ai_visible_duration_sec +
    z_log_time_before_ai_sec +
    (1 | task_id),
  data = engagement_ownership_data,
```

```
family = binomial,
control = glmerControl(optimizer = "bobyqa")
)
```

```
## boundary (singular) fit: see help('isSingular')
```

```
summary(model_engagement_ownership)
```

```
## Generalized linear mixed model fit by maximum likelihood (Laplace
## Approximation) [glmerMod]
## Family: binomial ( logit )
## Formula:
## epistemic_ownership_binary ~ read_carefully + z_log_ai_visible_duration_sec +
## z_log_time_before_ai_sec + (1 | task_id)
## Data: engagement_ownership_data
## Control: glmerControl(optimizer = "bobyqa")
##
##          AIC          BIC      logLik -2*log(L)  df.resid
##        187.6        202.4       -88.8      177.6       139
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -1.2342 -0.7408 -0.5713  1.1989  2.1889
##
## Random effects:
## Groups Name          Variance Std.Dev.
## task_id (Intercept) 0          0
## Number of obs: 144, groups: task_id, 18
##
## Fixed effects:
##
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    -0.92435    0.29101  -3.176  0.00149 **
## read_carefully     0.36822    0.37854   0.973  0.33068
## z_log_ai_visible_duration_sec  0.36582    0.18611   1.966  0.04934 *
## z_log_time_before_ai_sec    -0.03064    0.17971  -0.170  0.86464
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Correlation of Fixed Effects:
##          (Intr) rd_crf z_lg__v__
## read_crflly -0.782
## z_lg__vsb__  0.036 -0.124
## z_lg_tm_b__  0.109 -0.158 -0.113
## optimizer (bobyqa) convergence code: 0 (OK)
## boundary (singular) fit: see help('isSingular')
```

```
fixed_effects_table(model_engagement_ownership, exponentiate = TRUE)
```

```
##              term      estimate std_error    p_value odds_ratio
## 1      (Intercept) -0.92434782 0.2910085 0.001491358 0.3967901
## 2      read_carefully 0.36822359 0.3785408 0.330680152 1.4451651
## 3 z_log_ai_visible_duration_sec 0.36581990 0.1861065 0.049339222 1.4416956
## 4      z_log_time_before_ai_sec -0.03063548 0.1797061 0.864636253 0.9698290
```

Reading carefully did not predict binary epistemic ownership, $OR = 1.45$, $p = .331$. **Longer AI visible duration predicted higher odds of primarily own-thinking ownership, $OR = 1.44$, $p = .049$.** Time before viewing AI was not significant, $OR = 0.97$, $p = .865$. Together with Analyses 8B-8C, this suggests that longer viewing time may mark a more selective or editing-oriented strategy: participants who kept the AI visible longer produced less AI-like responses and were more likely to report the final answer as primarily their own thinking.

Jaccard Similarity vs. Epistemic Ownership

```
# Scatter: Jaccard similarity vs. epistemic ownership (continuous scale)
# Uses AI users from model_data; n is computed below.
scatter_data <- subset(
  model_data,
  ai_used == TRUE &
  !is.na(jaccard_similarity_to_ai) &
  !is.na(epistemic_ownership)
)

rho_jac_own <- cor(
  scatter_data$jaccard_similarity_to_ai,
  scatter_data$epistemic_ownership,
  method = "spearman",
  use = "complete.obs"
)

n_scatter <- sum(complete.cases(
  scatter_data[, c("jaccard_similarity_to_ai", "epistemic_ownership")]
))

# Tertile splits using rank to avoid duplicate-break error from tied 0s/1s
n_s <- sum(!is.na(scatter_data$jaccard_similarity_to_ai))
scatter_data$sim_rank <- rank(scatter_data$jaccard_similarity_to_ai,
  ties.method = "average", na.last = "keep")
scatter_data$similarity_tertile <- cut(
  scatter_data$sim_rank,
  breaks = c(0, n_s / 3, 2 * n_s / 3, n_s + 0.1),
  labels = c("Low", "Medium", "High"),
  include.lowest = TRUE
)
tertile_means <- aggregate(
  epistemic_ownership ~ similarity_tertile,
  data = scatter_data,
  FUN = mean, na.rm = TRUE
)
tertile_ns <- table(scatter_data$similarity_tertile)

p_scatter <- ggplot(scatter_data,
  aes(x = jaccard_similarity_to_ai, y = epistemic_ownership)) +
  geom_jitter(height = 0.15, width = 0.008,
    alpha = 0.35, size = 1.8, color = "#4472C4") +
  geom_smooth(method = "lm", color = "#C00000",
    se = TRUE, linewidth = 1.1) +
```



```


annotate("text", x = 0.72, y = 6.7,
         label = sprintf("rho = %.3f\\nn = %d", rho_jac_own, n_scatter),
         hjust = 0, size = 3.8, fontface = "italic") +
scale_y_continuous(breaks = 1:7, limits = c(0.5, 7.5)) +
labs(
  title = "Similarity to AI vs. Epistemic Ownership",
  subtitle = "AI users only; 1 = AI's thinking, 7 = my own thinking",
  x = "Jaccard Similarity to AI Response",
  y = "Epistemic Ownership"
) +
theme_classic(base_size = 12) +
theme(
  plot.title = element_text(face = "bold", size = 12),
  plot.subtitle = element_text(size = 10)
)

p_tertile <- ggplot(tertile_means,
  aes(x = similarity_tertile, y = epistemic_ownership)) +
geom_col(fill = "#70AD47", width = 0.6) +
geom_text(
  aes(label = sprintf("%.2f\\nn=%d", epistemic_ownership,
                        as.integer(tertile_ns[as.character(similarity_tertile)]))),
  vjust = -0.3, size = 3.2, fontface = "bold", lineheight = 0.9
) +
scale_y_continuous(limits = c(0, 6), expand = c(0, 0)) +
labs(
  title = "Ownership by Similarity Tertile",
  x = "Similarity tertile",
  y = "Mean Epistemic Ownership"
) +
theme_classic(base_size = 12) +
theme(plot.title = element_text(face = "bold", size = 12))

grid.arrange(p_scatter, p_tertile, ncol = 2,
  top = grid::textGrob(
    sprintf("Jaccard Similarity to AI and Epistemic Ownership (AI Users, n = %d)", n_scatter),
    gp = grid::gpar(fontface = "bold", fontsize = 13)
  )
)

## 'geom_smooth()' using formula = 'y ~ x'

```



Exploratory Correlation Matrix

This focused matrix gives a broad view of associations among mental-model variables, domain-specific AI-use frequency, engagement measures, similarity measures, ownership, and confidence, **among AI users only**. Spearman correlations are used because several variables are ordinal or skewed. This table is exploratory and should be used to identify patterns worth following up, not as a set of separate confirmatory tests.

```
correlation_data <- engagement_data

if (!"self_reported_freq_in_assigned_category" %in% names(correlation_data)) {
  correlation_data$self_reported_freq_in_assigned_category <- NA_real_
}

correlation_data$z_self_reported_freq <- z(
  correlation_data$self_reported_freq_in_assigned_category
)

correlation_data$z_user_decides <- z(
  correlation_data$exp_ai_does_work_vs_user_decides
)

correlation_data$z_independent <- z(
  -correlation_data$exp_independent_vs_reliant
)

correlation_data$z_instrument_vs_point_of_view <- z(
  correlation_data$exp_instrument_vs_point_of_view
)

correlation_data$z_existing_info_vs_creating_new <- z(
```

```

correlation_data$exp_existing_info_vs_creating_new
)

correlation_vars <- c(
  user_decides = "z_user_decides",
  independent_not_reliant = "z_independent",
  instrument_vs_point_of_view = "z_instrument_vs_point_of_view",
  existing_info_vs_creating_new = "z_existing_info_vs_creating_new",
  domain_freq = "z_self_reported_freq",
  read_carefully = "read_carefully",
  log_ai_visible_time = "log_ai_visible_duration_sec",
  log_time_before_ai = "log_time_before_ai_sec",
  log_response_writing_time = "log_response_writing_time_sec",
  jaccard_similarity = "jaccard_similarity_to_ai",
  embedding_cosine_similarity_to_ai = "embedding_cosine_similarity_to_ai",
  epistemic_ownership = "epistemic_ownership",
  confidence = "confidence_level"
)

correlation_matrix_data <- correlation_data[, correlation_vars]
names(correlation_matrix_data) <- names(correlation_vars)

spearman_correlations <- round(
  cor(
    correlation_matrix_data,
    use = "pairwise.complete.obs",
    method = "spearman"
  ),
  2
)

spearman_p_values <- matrix(
  NA_real_,
  nrow = length(correlation_vars),
  ncol = length(correlation_vars),
  dimnames = list(names(correlation_vars), names(correlation_vars))
)

for (i in seq_along(correlation_vars)) {
  for (j in seq_along(correlation_vars)) {
    x <- correlation_matrix_data[[i]]
    y <- correlation_matrix_data[[j]]
    complete_rows <- complete.cases(x, y)
    if (sum(complete_rows) >= 3) {
      spearman_p_values[i, j] <- suppressWarnings(
        cor.test(
          x[complete_rows],
          y[complete_rows],
          method = "spearman",
          exact = FALSE
        )$p.value
      )
    }
  }
}

```

```

    }
}

spearman_p_values <- round(spearman_p_values, 3)

spearman_correlations

```

```

##               user_decides independent_not_reliant
## user_decides           1.00              0.22
## independent_not_reliant 0.22              1.00
## instrument_vs_point_of_view 0.06          -0.18
## existing_info_vs_creating_new 0.08         -0.10
## domain_freq            0.01          -0.10
## read_carefully         0.02           0.10
## log_ai_visible_time    0.12           0.22
## log_time_before_ai     0.01           0.04
## log_response_writing_time 0.12          0.16
## jaccard_similarity     -0.12         -0.09
## embedding_similarity    -0.14         -0.04
## epistemic_ownership     0.21           0.04
## confidence             0.00          -0.03
##               instrument_vs_point_of_view
## user_decides                0.06
## independent_not_reliant     -0.18
## instrument_vs_point_of_view 1.00
## existing_info_vs_creating_new 0.42
## domain_freq                0.19
## read_carefully             0.17
## log_ai_visible_time       -0.02
## log_time_before_ai        0.03
## log_response_writing_time -0.08
## jaccard_similarity         0.05
## embedding_similarity       -0.01
## epistemic_ownership        0.01
## confidence                 0.11
##               existing_info_vs_creating_new domain_freq
## user_decides                0.08          0.01
## independent_not_reliant     -0.10         -0.10
## instrument_vs_point_of_view 0.42          0.19
## existing_info_vs_creating_new 1.00          0.33
## domain_freq                0.33          1.00
## read_carefully             0.13          0.08
## log_ai_visible_time       -0.08         -0.03
## log_time_before_ai       -0.12         -0.05
## log_response_writing_time -0.18          0.02
## jaccard_similarity         0.07         -0.06
## embedding_similarity        0.05        -0.11
## epistemic_ownership        0.13          0.07
## confidence                 0.24          0.24
##               read_carefully log_ai_visible_time
## user_decides                0.02          0.12
## independent_not_reliant     0.10          0.22
## instrument_vs_point_of_view 0.17         -0.02

```

## existing_info_vs_creating_new	0.13	-0.08
## domain_freq	0.08	-0.03
## read_carefully	1.00	0.18
## log_ai_visible_time	0.18	1.00
## log_time_before_ai	0.19	0.13
## log_response_writing_time	0.23	0.63
## jaccard_similarity	-0.20	-0.36
## embedding_similarity	-0.22	-0.37
## epistemic_ownership	0.15	0.18
## confidence	0.34	0.16
##	log_time_before_ai	log_response_writing_time
## user_decides	0.01	0.12
## independent_not_reliant	0.04	0.16
## instrument_vs_point_of_view	0.03	-0.08
## existing_info_vs_creating_new	-0.12	-0.18
## domain_freq	-0.05	0.02
## read_carefully	0.19	0.23
## log_ai_visible_time	0.13	0.63
## log_time_before_ai	1.00	0.21
## log_response_writing_time	0.21	1.00
## jaccard_similarity	-0.18	-0.59
## embedding_similarity	-0.25	-0.58
## epistemic_ownership	0.10	0.39
## confidence	0.17	0.13
##	jaccard_similarity	embedding_similarity
## user_decides	-0.12	-0.14
## independent_not_reliant	-0.09	-0.04
## instrument_vs_point_of_view	0.05	-0.01
## existing_info_vs_creating_new	0.07	0.05
## domain_freq	-0.06	-0.11
## read_carefully	-0.20	-0.22
## log_ai_visible_time	-0.36	-0.37
## log_time_before_ai	-0.18	-0.25
## log_response_writing_time	-0.59	-0.58
## jaccard_similarity	1.00	0.94
## embedding_similarity	0.94	1.00
## epistemic_ownership	-0.54	-0.54
## confidence	-0.11	-0.16
##	epistemic_ownership	confidence
## user_decides	0.21	0.00
## independent_not_reliant	0.04	-0.03
## instrument_vs_point_of_view	0.01	0.11
## existing_info_vs_creating_new	0.13	0.24
## domain_freq	0.07	0.24
## read_carefully	0.15	0.34
## log_ai_visible_time	0.18	0.16
## log_time_before_ai	0.10	0.17
## log_response_writing_time	0.39	0.13
## jaccard_similarity	-0.54	-0.11
## embedding_similarity	-0.54	-0.16
## epistemic_ownership	1.00	0.19
## confidence	0.19	1.00

spearman_p_values

```
##                                user_decides independent_not_reliant
## user_decides                    0.000                    0.005
## independent_not_reliant          0.005                    0.000
## instrument_vs_point_of_view      0.432                    0.027
## existing_info_vs_creating_new     0.322                    0.226
## domain_freq                      0.857                    0.196
## read_carefully                    0.793                    0.234
## log_ai_visible_time               0.128                    0.005
## log_time_before_ai               0.875                    0.659
## log_response_writing_time         0.147                    0.045
## jaccard_similarity                0.138                    0.255
## embedding_similarity              0.077                    0.620
## epistemic_ownership               0.007                    0.661
## confidence                        0.952                    0.671
##                                instrument_vs_point_of_view
## user_decides                      0.432
## independent_not_reliant            0.027
## instrument_vs_point_of_view        0.000
## existing_info_vs_creating_new      0.000
## domain_freq                       0.015
## read_carefully                     0.037
## log_ai_visible_time                0.804
## log_time_before_ai                0.687
## log_response_writing_time          0.303
## jaccard_similarity                 0.554
## embedding_similarity               0.880
## epistemic_ownership               0.940
## confidence                         0.183
##                                existing_info_vs_creating_new domain_freq
## user_decides                      0.322                    0.857
## independent_not_reliant            0.226                    0.196
## instrument_vs_point_of_view        0.000                    0.015
## existing_info_vs_creating_new      0.000                    0.000
## domain_freq                       0.000                    0.000
## read_carefully                     0.105                    0.308
## log_ai_visible_time                0.325                    0.686
## log_time_before_ai                0.135                    0.537
## log_response_writing_time          0.027                    0.795
## jaccard_similarity                 0.377                    0.485
## embedding_similarity               0.529                    0.189
## epistemic_ownership               0.103                    0.412
## confidence                         0.003                    0.002
##                                read_carefully log_ai_visible_time
## user_decides                      0.793                    0.128
## independent_not_reliant            0.234                    0.005
## instrument_vs_point_of_view        0.037                    0.804
## existing_info_vs_creating_new      0.105                    0.325
## domain_freq                       0.308                    0.686
## read_carefully                     0.000                    0.021
## log_ai_visible_time                0.021                    0.000
## log_time_before_ai                0.015                    0.107
```

## log_response_writing_time	0.004	0.000
## jaccard_similarity	0.013	0.000
## embedding_similarity	0.005	0.000
## epistemic_ownership	0.054	0.021
## confidence	0.000	0.051
##	log_time_before_ai	log_response_writing_time
## user_decides	0.875	0.147
## independent_not_reliant	0.659	0.045
## instrument_vs_point_of_view	0.687	0.303
## existing_info_vs_creating_new	0.135	0.027
## domain_freq	0.537	0.795
## read_carefully	0.015	0.004
## log_ai_visible_time	0.107	0.000
## log_time_before_ai	0.000	0.008
## log_response_writing_time	0.008	0.000
## jaccard_similarity	0.027	0.000
## embedding_similarity	0.001	0.000
## epistemic_ownership	0.219	0.000
## confidence	0.036	0.104
##	jaccard_similarity	embedding_similarity
## user_decides	0.138	0.077
## independent_not_reliant	0.255	0.620
## instrument_vs_point_of_view	0.554	0.880
## existing_info_vs_creating_new	0.377	0.529
## domain_freq	0.485	0.189
## read_carefully	0.013	0.005
## log_ai_visible_time	0.000	0.000
## log_time_before_ai	0.027	0.001
## log_response_writing_time	0.000	0.000
## jaccard_similarity	0.000	0.000
## embedding_similarity	0.000	0.000
## epistemic_ownership	0.000	0.000
## confidence	0.178	0.046
##	epistemic_ownership	confidence
## user_decides	0.007	0.952
## independent_not_reliant	0.661	0.671
## instrument_vs_point_of_view	0.940	0.183
## existing_info_vs_creating_new	0.103	0.003
## domain_freq	0.412	0.002
## read_carefully	0.054	0.000
## log_ai_visible_time	0.021	0.051
## log_time_before_ai	0.219	0.036
## log_response_writing_time	0.000	0.104
## jaccard_similarity	0.000	0.178
## embedding_similarity	0.000	0.046
## epistemic_ownership	0.000	0.017
## confidence	0.017	0.000

```
cor_heatmap <- spearman_correlations

heat_colors <- colorRampPalette(
  c("#2b6cb0", "white", "#c53030")
)(101)
```

```

par(mar = c(10, 10, 4, 2))

image(
  x = seq_len(ncol(cor_heatmap)),
  y = seq_len(nrow(cor_heatmap)),
  z = t(cor_heatmap[nrow(cor_heatmap):1, ]),
  col = heat_colors,
  zlim = c(-1, 1),
  axes = FALSE,
  xlab = "",
  ylab = "",
  main = "Spearman correlation matrix"
)

axis(
  1,
  at = seq_len(ncol(cor_heatmap)),
  labels = colnames(cor_heatmap),
  las = 2,
  cex.axis = 0.7
)


axis(
  2,
  at = seq_len(nrow(cor_heatmap)),
  labels = rev(rownames(cor_heatmap)),
  las = 2,
  cex.axis = 0.7
)

for (i in seq_len(nrow(cor_heatmap))) {
  for (j in seq_len(ncol(cor_heatmap))) {
    value <- cor_heatmap[nrow(cor_heatmap) - i + 1, j]
    text(
      x = j,
      y = i,
      labels = sprintf("%.2f", value),
      cex = 0.65
    )
  }
}

box()

```


Spearman correlation matrix



```
par(mar = c(5, 4, 4, 2) + 0.1)
```

The strongest pattern is that Jaccard and embedding similarity are highly correlated, $\rho = .94$, $p < .001$, so the lexical and semantic uptake measures capture very similar information among AI users.

Both uptake measures are strongly negatively related to epistemic ownership among AI users: Jaccard similarity, $\rho = -.54$, $p < .001$; embedding similarity, $\rho = -.54$, $p < .001$. In plain terms, participants whose final responses were more AI-like tended to report less own-thinking ownership and more perceived AI influence.

AI visible time and response writing time are positively correlated, $\rho = .63$, $p < .001$, and both are negatively related to AI similarity. AI visible time was negatively related to Jaccard similarity, $\rho = -.36$, $p < .001$, and embedding similarity, $\rho = -.37$, $p < .001$. Response writing time showed an even stronger negative association with Jaccard similarity, $\rho = -.59$, $p < .001$, and embedding similarity, $\rho = -.58$, $p < .001$. This supports the idea that fast full-copy users and slower selective/editing users are behaviorally different groups.

Reading carefully was positively related to confidence, $\rho = .34$, $p < .001$, but only marginally related to epistemic ownership, $\rho = .15$, $p = .054$. Reading carefully was also negatively related to Jaccard similarity,

$\rho = -.20$, $p = .013$, and embedding similarity, $\rho = -.22$, $p = .005$, again suggesting that more attentive engagement may reflect selective use rather than direct copying.

Instrument vs. point-of-view was moderately related to existing info vs. creating new, $\rho = .42$, $p < .001$, suggesting these two mental-model beliefs tend to cluster together.

Existing info vs. creating new was not significantly related to epistemic ownership among AI users, $\rho = .13$, $p = .103$, but it was positively related to confidence, $\rho = .24$, $p = .003$. This suggests that seeing AI as helping create something new may be more connected to confidence than to ownership once people have already used AI.

User decides was positively related to epistemic ownership, $\rho = .21$, $p = .007$. People who believed AI should defer to their judgment tended to report more own-thinking ownership.

Independent/not reliant was not significantly related to epistemic ownership, $\rho = .04$, $p = .661$, reinforcing that this predictor is empirically distinct from **user_decides** and does not consistently predict ownership.